





Causa1ontology:

AGI's Missing Link

D. R. Disson



Causa1ontology:

AGI's Missing Link

D. R. Dison



The Roadmap to AGI and ASI

TABLE OF CONTENTS:

PART ONE: THE WALL..... 11

CHAPTER 1: THE AGE OF MIRACLES THAT CANNOT REASON
..... 11

CHAPTER 2: THE LADDER THEY CANNOT CLIMB..... 15

CHAPTER 3: THE BILLION-DOLLAR ADMISSION..... 20

CHAPTER 4: THE BLACK BOX CRISIS..... 25

PART TWO: THE MISSING LINK..... 36

CHAPTER 5: WHAT IF CAUSATION WERE A THING?..... 36

CHAPTER 6: A VERB-FIRST, NOUN-HOSTING ONTOLOGY. 44

CHAPTER 7: TWENTY-ONE KINDS OF THING..... 48

CHAPTER 8: IDENTITY WITHOUT A COMMITTEE..... 70

CHAPTER 9: EVERY WORD HAS AN AUTHOR..... 74

PART THREE: THE ENABLER..... 109

CHAPTER 10: CLIMBING THE LADDER BY DATA STRUCTURE
..... 109

CHAPTER 11: THE FACULTIES OF A GENERAL MIND... 112

CHAPTER 12: THE GLASS BOX MIND..... 115

CHAPTER 13: THE COMMONS: A WORLD MODEL FOR
EVERYONE..... 118

CHAPTER 14: COMMONS STORAGE AND FEDERATION DESIGN
..... 126

CHAPTER 15: STIGMERGY: THE WALL THAT TEACHES.. 148

CHAPTER 16: THE FIRST GARDENER..... 151

PART FOUR: THE PROOF..... 155



The Roadmap to AGI and ASI

CHAPTER 17: THE FREEZE..... 155

CHAPTER 18: TWENTY LANGUAGES, ZERO SHARED CODE 158

CHAPTER 19: INSTALL IT BEFORE BREAKFAST..... 160

APPENDIX: THE SPECIFICATION IN BRIEF..... 178

BACK COVER MATTER..... 211



The Roadmap to AGI and ASI

CAUSALONTOLOGY

AGI'S MISSING LINK

A Manifesto

D. R. Dison

Founder of AIU (Artificial Intelligence University, the institution the author leads),
Creator and Owner of PrologAI and Mentova

ORCID (Open Researcher and Contributor Identifier, a permanent public ID number that ties a researcher to their work) 0009-0001-9246-5758 -
[linkedin.com/in/d-r-dison/](https://www.linkedin.com/in/d-r-dison/)

Presented at THE AGI 2026 conference, San Francisco.

The living standard, the code, and every receipt behind every claim in this book:

github.com/ai-university-aiu/causalontology

License: "The attribution always; no profit, no problem license."

EPIGRAPH

A child flips a light switch a hundred times,
watching the ceiling.

In that small, ridiculous, magnificent experiment
lives everything this book is about.

The child is not predicting the next word.



The Roadmap to AGI and ASI

The child is asking the universe a question with her own hand - and writing the answer into a structure that will outlive the afternoon.

By "structure" we simply mean an organized place to keep something, the way a spice rack is an organized place to keep spices - each jar in its slot, easy to find again.

We built machines that read every book on Earth.

We forgot to give them somewhere to keep what the child learns before she can read at all.

THE CLAIM

This book makes one claim, and makes it without flinching.

The missing ingredient of Artificial General Intelligence (AGI - a machine that can think and learn across many kinds of problem, the way a person does, not just one narrow task) is not a bigger model, a longer context window, or another trillion tokens of text.

A token is a small chunk of text, roughly a word or a piece of a word, that these systems read and write one at a time.

A context window is simply how much text one of these systems can keep in view at once, the way you can only hold so many playing cards in your hand.

The missing ingredient is a DATA STRUCTURE.



The Roadmap to AGI and ASI

A data structure is a deliberately chosen way of organizing information so that a computer, or a mind, can store it, find it again, and work with it – the spice rack, not the spices.

Causation itself – the knowledge of what makes what happen – reified into objects a mind can hold, sign, merge, inspect, and improve.

To reify something (from the Latin for "thing") is to turn an idea you can only talk about into an actual object you can pick up, point at, hand to someone else, and check – the way a promise becomes a signed contract you can file in a drawer.

Here "sign" means to attach a tamper-evident mark that proves who vouched for a piece of knowledge, like a wax seal on a letter; "merge" means to fold two people's knowledge into one shared pile without anything getting lost; and "inspect" means anyone is allowed to open it up and look inside.

That data structure now exists.

It is called Causalontology, and it is the world's first synthesis of reified causal representation, temporal-mechanistic modeling, and interventional learning into a single glass-box unified ontology.

Let us unpack that mouthful slowly, because every word is doing real work.

"Reified causal representation" means cause-and-effect written down as solid objects rather than fuzzy hunches.



The Roadmap to AGI and ASI

"Temporal-mechanistic modeling" means keeping track of how things unfold over time and by what actual mechanism, not just what tends to appear next to what.

"Interventional learning" means learning by DOING something and watching what changes - the child's hand on the switch - rather than only by reading about it.

A "glass box" is the opposite of a black box: you can see straight through the walls and check every gear inside, instead of being told "trust me" by a machine you cannot open.

An "ontology" is an agreed-upon list of what kinds of things exist in some domain and how they relate to each other - in plain terms, a shared, careful vocabulary of reality that everyone can point to and use the same way.

Said from a second angle, for the ontologists: it is the world's first synthesis of a process-first causal calculus, a continuant noun backbone, and a realizable hinge between them - one foundational ontology governing the verbs and hosting the nouns.

In everyday language: it puts happenings first (a "process-first" view treats events and changes - the verbs of the world - as the main characters), it keeps a sturdy list of the lasting things those happenings happen to (a "continuant" is a thing that persists through time, like a person or a brick - the nouns), and it connects the two with a hinge that can be "realized," meaning a standing capability (a switch is able to make light) turns



The Roadmap to AGI and ASI

into an actual event (the switch makes light right now).

A "foundational ontology" is just the bottom-most, most general such vocabulary - the ground floor every other description is built on.

Said from a third angle, for the philosophers: the world's first synthesis of symbolic causal reasoning, the philosophy of process, and embodied interventional learning, unified in a single glass-box ontology.

Plainly: it reasons about cause and effect using clear named symbols you can read (not hidden numbers), it takes seriously the old philosophical idea that reality is fundamentally made of processes and happenings, and it learns the way a body learns - by acting in the world and feeling the result.

Said from a fourth angle, for the scholars: the world's first synthesis of process ontology, causal modeling, and constructivist learning into a single foundational ontology.

In ordinary words: a picture of reality built from processes, a way of writing down causes, and the idea (called "constructivist") that a mind builds its understanding piece by piece from its own experience - all three joined into one ground-floor vocabulary.

And said succinctly, for everyone: the world's first synthesis of reification, temporal causation, and interventional learning in one unified, glass-box ontology.



The Roadmap to AGI and ASI

That is: turning causes into things you can hold, tracking how those causes play out over time, and learning them by doing - all inside one see-through, shared vocabulary of reality.

This is not a proposal.

It is a shipped, frozen standard with one hundred and thirty-seven public conformance tests, twenty independent implementations in twenty programming languages that agree with each other byte for byte, a running shared store, and a synthetic mind already contributing signed causal knowledge to it.

A "conformance test" is a public check anyone can run that asks "does this thing behave exactly the way the standard says it should?" - like a driving test that every implementation must pass identically.

"Agree byte for byte" means that when you run the same job through all twenty versions, they produce output that matches down to the very last character, with zero disagreement - the strongest possible sign they truly do the same thing.

A "shared store" is one common storehouse of knowledge that many minds can read from and write to together, like a public library everyone contributes books to.

A "synthetic mind" is an artificial mind built in software; here it is already depositing new cause-and-effect knowledge into that shared store, each deposit marked with a tamper-evident signature showing who it came from.



The Roadmap to AGI and ASI

You can install it before you finish your coffee.

You can verify every sentence of this book in public.

That is the claim.

Now let us earn it.

PART ONE: THE WALL

Why the current road, magnificent as it is, does not reach AGI.

CHAPTER 1: THE AGE OF MIRACLES THAT CANNOT REASON

We live in an age of miracles, and the miracles cannot reason about a light switch.

Hold both halves of that sentence at once, because the whole argument of this book lives in the tension between them.

The miracles are real.

Machines now write serviceable sonnets, pass professional examinations, translate between a hundred tongues, and converse with a fluency that would have read as science fiction a decade ago.

Nobody serious disputes this, and this book will not waste your time pretending otherwise.

But walk into any laboratory measuring what these systems actually UNDERSTAND about cause and effect, and the temperature drops.



The Roadmap to AGI and ASI

When researchers tested state-of-the-art Large Language Models (LLMs – the text-predicting systems behind the modern chatbots) on inferring causal relationships from real-world text, the best models managed an average F1 score of roughly 0.535.

An F1 score is a single number that blends two measures of getting things right – precision and recall – into one grade that runs from 0 (useless) to 1 (perfect).

1.0 is perfect; a coin has no shame at 0.5.

("Precision" asks: of the answers it gave, how many were correct?

"Recall" asks: of the correct answers out there, how many did it find?

The F1 score rewards you only when both are good.)

Barely better than the coin – at the single capability that separates understanding from imitation.

And it gets worse in exactly the direction that matters.

The same research found performance degrading substantially as causal information becomes implicit – which is to say, as the text stops SAYING "because" and starts merely meaning it.

("Implicit" means unstated, left for you to figure out, as opposed to spelled out in plain words.)

The real world almost never says "because."



The Roadmap to AGI and ASI

The real world is implicit all the way down.

Why does this happen to systems that seem so brilliant?

Because of what they are made of.

An LLM is trained on a staggering corpus of text to do one thing supremely well: predict the next token.

(A "corpus" is just a large collected body of text; recall that a token is a word-sized chunk of writing.)

So the machine's entire schooling is a single game: given everything so far, guess the next little chunk.)

Prediction of this kind captures statistical patterns – which words travel together, which phrases follow which – with superhuman thoroughness.

A "statistical pattern" is simply a regularity you notice by counting – for example, that the word "salt" often shows up near the word "pepper."

But the clinical literature now documents what the architecture implies: these systems operate on statistical correlation, not on mechanistic models of how anything works, and their reasoning bends and breaks where a causal model would hold firm.

Two ideas are being contrasted here, and the whole book turns on the difference.



The Roadmap to AGI and ASI

"Correlation" means two things tend to show up together, like umbrellas and rain, without telling you which causes which, or whether anything does.

"Causation" means one thing actually makes the other happen - the rain is why the umbrellas come out.

A "mechanistic model" is an inner picture of how a thing works, gear by gear, so you can predict what happens when you change a part; that is exactly what these systems lack.

They are, in the most precise sense, correlation engines.

And we have placed a correlation engine at the center of civilization's information supply.

Let the child back into the room.

She flips the switch.

The light answers.

She flips it again - and again - and on the forty-first flip, nothing, because the bulb has died, and NOW she is fascinated, because the rule she was building just met its first exception, and exceptions are where understanding is forged.

By sunset she owns a piece of causal knowledge no amount of reading could have given her: pressing causes lighting, usually, quickly, unless.

Ask yourself what she has that the trillion-token machine does not.



The Roadmap to AGI and ASI

Not vocabulary.

Not facts.

She has a place to PUT what the switch taught her - some structure in her small skull that stores not the words about the light but the mechanism of it, with slots for how often, how fast, and what spoils it.

That structure is the subject of this book.

The machines do not have one.

We propose to hand it to them - and to everyone.

CHAPTER 2: THE LADDER THEY CANNOT CLIMB

The clearest map of what is missing was drawn by Judea Pearl, the Turing Award laureate who spent a career turning causation into mathematics.

The Turing Award is often called the Nobel Prize of computing, so "laureate" here simply means a winner of that top honour.

Causation is the plain idea that one thing actually makes another thing happen - flipping the switch is what makes the light come on, not merely something that appears alongside it.

Pearl calls his map the Ladder of Causation, and it has three rungs.

The Ladder of Causation is just a way of ranking three kinds of thinking from simplest to hardest, drawn as a ladder you climb one rung at a time.



The Roadmap to AGI and ASI

Rung one: ASSOCIATION.

Seeing.

What goes with what?

The barometer falls, and rain follows; the rooster crows, and the sun rises.

A barometer is the little dial that measures air pressure and hints at coming weather.

Association is the level of noticing that two things tend to show up together, which experts call correlation - two things moving in step, without any claim that one causes the other.

Every statistical method ever devised, and every neural network ever trained, lives natively on this rung.

A neural network is the kind of computer program behind most of today's Artificial Intelligence: a giant web of simple number-crunching units that learns patterns by being shown mountains of examples, the way a child learns to recognise cats after seeing thousands of them.

"Lives natively on this rung" just means this is the only floor these methods truly stand on; they are experts at spotting what-goes-with-what, and nothing higher.

Rung two: INTERVENTION.

Doing.

What happens if I ACT - not observe, act?



The Roadmap to AGI and ASI

An intervention means you reach in and change something on purpose, rather than sitting back and watching.

Smash the barometer and the storm comes anyway; silence the rooster and the sun is unmoved.

One action severs a thousand false correlations, which is why the child at the switch learns more in an afternoon than a telescope pointed at ten thousand light switches ever could.

To sever a correlation is to cut the imagined link between two things once you prove that touching one does not move the other.

Acting beats watching: the child who actually flips the switch discovers in minutes what no amount of pure staring could ever settle.

Rung three: COUNTERFACTUAL.

Imagining.

What **WOULD** have happened, had things been otherwise?

A counterfactual is a "what if things had gone differently" question - a rerun of the past in your imagination with one detail changed.

Would the patient have lived with the other treatment?

Would the bridge have held without the third rivet?



The Roadmap to AGI and ASI

This is the rung of responsibility, regret, credit, blame, planning, and science itself - the rung where minds do their most human work.

Now the uncomfortable verdict, and it is not ours alone.

The assessment coming out of the causal-inference community is blunt: machine learning as practiced remains on rung one, and today's Large Language Models (LLMs) are sophisticated statistical approximators stuck at the bottom of the ladder, unable to reliably answer intervention questions, let alone counterfactual ones.

A Large Language Model (LLM) is the type of system behind popular chatbots: it is trained on enormous amounts of text and learns to predict which word tends to come next, which makes it fluent but keeps it on the seeing rung.

"Statistical approximators" is a fancy way of saying they are very good pattern-matchers that guess by resemblance rather than reasoning about causes.

Stuck is the operative word, and here is the reason scale does not unstick them.

Here "scale" means simply making the system bigger - more data, more computer power, more of the same.

More association is still association.

A model trained on ten times the text learns the word "because" ten times more thoroughly - as a WORD, as a token that keeps certain company.



The Roadmap to AGI and ASI

A token is the small chunk of text – roughly a word or piece of a word – that these systems chop language into and shuffle around; to the machine, "because" is just such a chunk, not a claim about cause.

But rungs are not climbed by reading about ladders.

The second rung is made of actions and their consequences; the third is made of structured alternatives to reality; neither is made of adjacency in text.

Adjacency in text simply means which words sit near which other words on the page – the only raw material these models actually have.

To climb, a system needs somewhere for the climbing to LIVE: a representation whose native objects are interventions, mechanisms, timings, and might-have-beens.

A representation is just the form in which a system stores what it knows; "native objects" means the basic building blocks that form is made of – here, actions, cause-and-effect machinery, when-things-happen, and what-might-have-been, rather than mere neighbouring words.

Hold that thought – "somewhere for the climbing to live" – because it is the hinge of this entire book.

The field has diagnosed the ladder problem eloquently for years.



The Roadmap to AGI and ASI

What it has not done, until now, is answer the question a builder asks the morning after agreeing with the diagnosis:

Fine.

WHERE, exactly, in what format, with what identity, signed by whom, merged how, does a piece of rung-two knowledge physically reside?

"Signed" here points to a digital signature, an unforgeable mathematical seal that proves who created a piece of information and that nobody tampered with it, the electronic cousin of a wax seal on a letter.

"Merged" points to combining two separate stores of knowledge into one without a fight over which version wins.

Show me the bytes.

Bytes are the tiny units of data a computer actually stores, so "show me the bytes" means show me the knowledge as something real and concrete on a disk, not as a vague promise.

This book shows you the bytes.

CHAPTER 3: THE BILLION-DOLLAR ADMISSION

If you think the wall we have described is the grumbling of outsiders, consider what the insiders are doing with their money and their careers in this very year, 2026.

One of the founding architects of deep learning – a man who spent years arguing that neural networks



The Roadmap to AGI and ASI

alone would carry us - has left the biggest platform in the industry to found a laboratory dedicated to WORLD MODELS, raising a seed round in excess of one billion dollars to pursue prediction, planning, and learning from the real world rather than from text alone.

Deep learning is the branch of Artificial Intelligence built on large neural networks with many stacked layers; it is the engine behind most recent breakthroughs.

A seed round is the very first big batch of investor money a new company raises to get started, so "a seed round in excess of one billion dollars" means more than a thousand million dollars handed over before the company has even proven itself.

A world model is an internal picture of how reality behaves - a mental simulator that lets a system predict what happens next and plan ahead, the way you can imagine a ball rolling off a table before it actually falls.

Read that again: a billion-dollar seed, not for a bigger language model, but for the explicit thesis that language models are not enough.

A thesis here is simply the central claim someone is betting on.

And the same pioneer, long and famously hostile to symbolic approaches, is now acknowledged - by his critics, with visible astonishment - to have changed teams: conceding that symbol manipulation is a necessary ingredient of humanlike intelligence.



The Roadmap to AGI and ASI

Symbolic approaches, or symbol manipulation, means building intelligence out of explicit, readable pieces - named facts, rules, and logic you can point at and check - rather than out of the blurry numbers inside a neural network.

Around him, an entire movement has gone mainstream.

Neurosymbolic Artificial Intelligence - the marriage of neural pattern recognition with symbolic structure - has moved from heresy to conference track, with dedicated workshops on explainable neurosymbolic systems and a growing consensus that domains demanding trustworthiness require symbolic bones under the neural muscle.

Neurosymbolic Artificial Intelligence simply means combining the two styles above: the pattern-spotting talent of neural networks joined to the clear, checkable logic of symbols.

"From heresy to conference track" means an idea once mocked is now given its own official session at the big research gatherings.

We recount all this not to gloat but to mark a historic moment with precision.

THE FIELD HAS CONCEDED THE DIAGNOSIS.

Everyone at the frontier now agrees that something structural is missing - that scale alone has hit a wall made of exactly the material Pearl said it would be made of.



The Roadmap to AGI and ASI

The remaining fight is over WHAT the missing thing is.

The billion-dollar bet says: the missing thing is a better MODEL - train a network to imagine the world's dynamics inside its weights.

"Its weights" are the billions of adjustable numbers inside a neural network that hold everything it has learned; picture countless tiny dials that training slowly turns, and note that a human cannot simply read them like a book.

We wish that bet well; richer models help everyone.

But notice what even a perfect world model, trapped in weights, still lacks.

It cannot be read.

It cannot be audited claim by claim.

To audit something means to inspect it piece by piece to confirm each part is correct and accounted for, the way an accountant checks a ledger line by line.

It cannot be signed, merged with a rival's model, corrected surgically, or handed to another mind.

"Corrected surgically" means fixing one precise wrong belief without disturbing everything around it, like removing a single bad tooth instead of rebuilding the whole jaw.



The Roadmap to AGI and ASI

Two labs training two world models end up with two mutually unintelligible skulls - each a new black box, however brilliant.

"Mutually unintelligible" means neither can read or understand the other.

A black box is any system whose inner workings are hidden from view: you see what goes in and what comes out, but not how or why - the opposite of the glass box this book champions.

Our bet is different, and it is the wager of this entire book.

The missing thing is not a better model.

The missing thing is a missing DATA STRUCTURE - a common, open, verifiable format for causal knowledge itself, standing outside any one skull, shared by all of them.

A data structure is simply an agreed way of organising information so that any program can store it, read it, and pass it along - like a standard form everyone fills in the same way.

"Verifiable" means anyone can check it and confirm for themselves that it is genuine and untampered.

Models come and go; representations remain.

The web did not happen because someone trained a better browser; it happened because HyperText Markup Language (HTML) gave every browser the same thing to read.



The Roadmap to AGI and ASI

HyperText Markup Language (HTML) is the shared, plain-text code that describes what a web page contains; because every browser agreed to read the same HTML, any page could be shown by any browser anywhere.

AGI's HTML moment is still missing.

AGI stands for Artificial General Intelligence, meaning a machine that can reason across many different kinds of problems with the flexible, all-purpose competence of a human mind, rather than being clever at just one narrow task.

That is the vacancy Causalontology fills.

Causalontology is the shared, open format this book proposes for writing down cause-and-effect knowledge – the thing every intelligent system could read, in the way every browser reads HTML.

CHAPTER 4: THE BLACK BOX CRISIS

There is a second wall, and it is scarier than the first.

Suppose the scaling optimists are right and the current road does produce systems of general, superhuman competence.

The scaling optimists are the people who believe that simply making today's systems bigger will, on its own, get us all the way to human-level and beyond.

Now ask the only question that will matter that day: WHAT DO THEY BELIEVE, AND WHY?



The Roadmap to AGI and ASI

With today's architectures, nobody can answer – not the builders, not the auditors, not the machine itself.

An architecture, in this setting, is the overall design and inner arrangement of an Artificial Intelligence system.

The knowledge is smeared across billions of numbers; the interpretability researchers, doing heroic work, are archaeologists excavating a city that is being rebuilt nightly.

Interpretability is the science of trying to figure out what a neural network is actually doing inside – reading meaning back out of those billions of raw numbers – and the picture in that sentence is of diggers uncovering an ancient city that someone keeps demolishing and rebuilding every night, so the work never settles.

The safety community's own assessments for 2026 say the quiet part plainly: alignment and interpretability remain unsolved at the exact moment capabilities compound.

Alignment is the challenge of making sure a powerful Artificial Intelligence actually wants what we want – that its goals and behaviour stay in line with human intentions – and "unsolved" means nobody yet has a reliable way to guarantee it.

"Capabilities compound" means the systems are getting more powerful faster and faster, building on themselves like interest on interest.

You cannot align what you cannot read.



The Roadmap to AGI and ASI

You cannot correct a belief you cannot locate.

You cannot trust a justification generated AFTER the answer, by the same opaque process, dressed in the grammar of reasons.

An opaque process is one you cannot see inside – the very opposite of transparent – and "dressed in the grammar of reasons" means an explanation that sounds logical and well-reasoned on the surface while hiding that it was invented after the fact.

We have a word for confident, fluent, unauditable assertion.

Unauditable means it cannot be checked piece by piece; an assertion is simply a stated claim.

When a machine does it, we politely call it hallucination.

Hallucination is the polite industry term for when an Artificial Intelligence states something false with complete confidence, as if making it up on the spot.

The glass box alternative is not a dream; it is an engineering discipline with a simple, brutal rule.

A glass box is the opposite of a black box: a system whose inner workings are fully visible and inspectable, so you can see exactly what it knows and why it acts.

THE EXPLANATION MUST BE THE REPRESENTATION.



The Roadmap to AGI and ASI

Not a summary of it, not a rationalization near it
- the thing itself, readable.

A rationalization is a nice-sounding excuse
invented to justify an answer already given,
rather than the real reason behind it.

A mind whose world model is made of discrete,
signed, inspectable causal claims can be audited
the way accounts are audited: claim by claim,
author by author, evidence grade by evidence
grade.

"Discrete" here means made of separate, countable
pieces you can pick up one at a time, instead of a
blurry continuous smear.

"Evidence grade" means a rating of how strong the
support for each claim is, from flimsy hearsay up
to solid proof.

Change its mind and you can point at the exact
object that changed.

Ask it "why?" and the answer is a lookup, not a
performance.

A lookup means fetching the actual stored reason
directly, the way you find a word in a dictionary,
rather than staging a convincing show of
reasoning.

Hold the two walls side by side and the shape of
the missing piece appears in silhouette.

Wall one: no place for causal knowledge to live.



The Roadmap to AGI and ASI

Wall two: no way to read the knowledge that does live inside.

One structure demolishes both - IF that structure is causal to its bones, open to every mind, and glass all the way down.

"Causal to its bones" means cause-and-effect is built into its very foundation, not bolted on as an afterthought; "glass all the way down" means it is inspectable at every level, with nothing hidden.

Now we can stop describing the hole and start describing the key.

INTERLUDE ONE

A SHORT HISTORY OF THE ARROW

Before we describe the key, honor the lock, because humanity has been picking at it for twenty-four centuries.

By "the arrow" we simply mean the little arrow we draw when we say one thing causes another, as in "A causes B".

That arrow stands for causation, which means the relationship where one thing actually makes another thing happen.

Aristotle started it, as he started most things, by refusing to accept that "why" had only one answer.



The Roadmap to AGI and ASI

Aristotle was an ancient Greek philosopher, one of the first people to try to organize all of human knowledge into neat categories.

He counted four causes - the material, the formal, the efficient, the final - and if his taxonomy creaks today, his instinct was immortal: causation is STRUCTURED, and the structure can be named.

A taxonomy is just an organized scheme of named categories, the way a library sorts books onto labeled shelves.

Then came David Hume, in the eighteenth century, with the coldest shower in the history of philosophy.

David Hume was a Scottish thinker famous for doubting things other people took for granted.

Look as hard as you like, Hume said, and you will never SEE causation - only one thing following another, again and again, until habit forges an expectation and calls it a law.

All we ever observe is correlation; the arrow is something the mind paints on.

Correlation means two things tend to happen together, or one tends to follow the other, without any proof that one actually makes the other happen.

The rooster crows and then the sun rises every single morning, yet the rooster does not cause the dawn - that gap between "happens together" and "actually makes it happen" is exactly Hume's point.



The Roadmap to AGI and ASI

The sciences responded to Hume in the strangest possible way: they surrendered.

The statistics that grew up in the nineteenth and twentieth centuries – brilliant, indispensable statistics – was built almost entirely on Hume's side of the line: correlation coefficients, regressions, associations, with "correlation is not causation" recited as a warning label and then, too often, used as an excuse never to formalize causation at all.

Statistics is the mathematics of finding patterns in large piles of numbers.

A correlation coefficient is a single number that measures how tightly two things move together.

A regression is a method for drawing the best-fitting line through scattered data to summarize a trend.

To formalize something means to pin it down in exact, rule-based terms that a machine or a mathematician can work with, instead of leaving it as a loose idea.

For a century, the arrow was homeless in the very disciplines that needed it most.

There were prophets in the wilderness.

Sewall Wright, a geneticist in the 1920s, drew little diagrams with arrows – path analysis – and was politely ignored for fifty years.



The Roadmap to AGI and ASI

A geneticist is a scientist who studies how traits are passed from parents to offspring.

Path analysis is his name for those arrow diagrams, a way of laying out on paper which thing is supposed to influence which.

The philosophers kept the flame: counterfactual accounts, mechanism accounts, process theories – rich, careful, and almost entirely disconnected from anything a machine could compute.

A counterfactual is a "what if things had gone differently" question, such as "would the patient have lived if she had not smoked?".

A mechanism is the step-by-step chain of moving parts that explains HOW one thing leads to another, like the gears inside a clock.

And then, in our own era, Judea Pearl and his colleagues did the thing Hume said could not be done: they gave the arrow MATHEMATICS.

Judea Pearl is a computer scientist and one of the leading modern figures in the study of cause and effect.

The do-operator – the formal difference between seeing and doing.

The do-operator is Pearl's precise notation for the difference between merely watching something happen and actually reaching in and making it happen on purpose.

Structural causal models.



The Roadmap to AGI and ASI

A structural causal model is a formal blueprint of a situation that spells out, in equations, which things affect which other things, so you can calculate cause and effect instead of just guessing.

The calculus that turns "would the patient have lived?" from a sigh into a computation.

Here calculus just means a worked-out system of rules for calculating, and a computation is a calculation a machine can carry out step by step.

A Turing Award followed, and an entire discipline of causal inference now thrives on those foundations.

The Turing Award is often called the Nobel Prize of computing, the highest honor in the field.

Causal inference is the whole area of study devoted to working out true cause and effect from data, rather than just noticing patterns.

So why, with the mathematics in hand, is the arrow still homeless in our largest machines?

Because mathematics is not INFRASTRUCTURE.

Infrastructure means the shared, agreed-upon plumbing that lets many people and machines work together – things like standard units, common file formats, and rules everyone follows, the way roads and power lines let a whole country function.

Pearl gave us the calculus of causation the way Newton gave us the calculus of motion – but between Newton and the moon landing stood three



The Roadmap to AGI and ASI

centuries of engineering: units agreed, tables published, formats standardized, tools interoperable.

Isaac Newton was the physicist who first wrote down the mathematical laws of motion and gravity.

Interoperable means different tools, built by different people, can work smoothly with one another because they all speak the same agreed language.

Causation in 2026 has the calculus and almost none of the engineering.

There is no agreed unit, no common format, no signature scheme, no merge rule, no shared table of what humanity knows about what makes what happen.

A signature scheme is a mathematical way of stamping a piece of information so that anyone can check who created it and confirm that nobody has tampered with it since, like a tamper-proof wax seal made of arithmetic.

A merge rule is an agreed procedure for combining two people's separate contributions into one shared version without a fight over which copy wins.

Every discipline that uses Pearl's mathematics writes its models in its own notation, in its own files, unreadable to its neighbors, unsigned, unmergeable.

Notation just means the particular set of symbols and shorthand a field uses to write things down.



The Roadmap to AGI and ASI

Twenty-four centuries to name the arrow.

Three decades since we learned to calculate with it.

And the arrow still has nowhere to LIVE.

That is the exact size and shape of the hole in the world.

Causalontology is not a new philosophy of causation and not a rival calculus - it is the missing third act: the ENGINEERING.

Causalontology is the name of the system this book is about; the word joins "causal" (about cause and effect) with "ontology" (a formal, agreed-upon catalog of what kinds of things exist and how they relate), so it literally means an organized, shared catalog of causes.

The units are agreed, down to the second.

The format is frozen, down to the byte.

A byte is the tiniest standard chunk of computer data, so "frozen down to the byte" means the format is pinned down so exactly that every copy is identical to the last detail.

The signatures are arithmetic; the merge rule is a theorem; the table is public and its doors are open.

A theorem is a statement that has been proven true beyond doubt using logic, so "the merge rule is a theorem" means combining contributions is not a



The Roadmap to AGI and ASI

matter of opinion but a guaranteed, provable procedure.

Aristotle named it; Hume doubted it; Wright drew it; Pearl calculated it.

We gave it an address.

PART TWO: THE MISSING LINK

What Causalontology is - explained so a newcomer can hold all of it.

CHAPTER 5: WHAT IF CAUSATION WERE A THING?

Here is the founding move, and like every great founding move it sounds almost too simple once said aloud.

TAKE THE ARROW SERIOUSLY.

For a century we have scribbled "A causes B" and treated the arrow as punctuation - decoration between the nouns, which got all the attention.

By "the nouns" we mean the things on either side of the arrow, the A and the B, the objects our language loves to name and fuss over.

Causalontology REIFIES the arrow.

To reify means to turn an idea into a thing: a thing with a name, a boundary, an identity - a thing you can store, send, sign, compare, and improve.



The Roadmap to AGI and ASI

Reification is that same move as a noun; you are promoting something we usually treat as an abstract idea (the causing) into a solid object you can pick up and handle, the way "running" the activity could be turned into "a run" you record and share.

The reified arrow is called a Causal Relation Object - and it is the atom of understanding, the fundamental unit of this whole architecture.

(We always write it out in full; this standard, by its Principle P7, uses whole words and never abbreviations, so that no reader ever has to decode a cipher to read a claim.)

By "atom" we mean the smallest indivisible building block, the way atoms are the basic pieces everything else is made of.

By "architecture" we mean the overall design and structure of the whole system, the way a building's architecture is its underlying plan.

Look inside one.

- CAUSES: which happenings bring this about - not free text, but references to defined process types.

Free text would be loose words a human types however they like; instead each cause points to a defined process type, which is an agreed, catalogued name for a kind of happening (such as "pressing" or "heating"), so that everyone means exactly the same thing.



The Roadmap to AGI and ASI

A process type is one of these catalogued happenings – an event or activity that unfolds over time, as opposed to a standing object – and it is the kind of entry the CAUSES and EFFECTS fields point at.

- EFFECTS: which happenings result.
- MECHANISM: the finer-grained causal objects that explain HOW – the gears inside, each of them a Causal Relation Object in its own right, turtles of explanation all the way down until you choose to stop.

"Turtles all the way down" is an old joke about an endless stack with no bottom; here it playfully means every mechanism can itself be explained by smaller mechanisms inside it, as far down as you care to look.

- TEMPORAL: the window – how soon may the effect follow, and how late is too late – with fixed, agreed units, down to a month defined as exactly 2,629,746 seconds so that no two minds ever disagree about admissibility arithmetic.

"Temporal" simply means having to do with time.

Admissibility arithmetic means the arithmetic that decides whether an effect arrived within the allowed time window and therefore counts.

- MODALITY: the flavor of the claim – necessary, sufficient, contributory, preventive, or enabling – because "smoking contributes to cancer" and "the switch suffices for the light" are different shapes of truth.



The Roadmap to AGI and ASI

Modality means the KIND or strength of a claim; here "necessary" means the effect cannot happen without this cause, "sufficient" means this cause alone is enough to bring the effect about, "contributory" means this cause helps make the effect more likely without guaranteeing it, "preventive" means this cause works to stop the effect from happening, and "enabling" means this cause does not itself push the effect but opens the door for another cause to work, the way oxygen enables a fire it does not start.

- **CONTEXT:** the enabling conditions under which the claim holds at all.

Enabling conditions are the background circumstances that must already be true for the claim to work - for a light switch, that the power is on and the bulb is not burned out.

- **STRATUM:** the level of description at which the claim is pitched - molecular, cellular, social - a named rung on a named ladder, so that "this causes that" always states at what grain it holds.
- **SKIPS:** a flag declaring that the cause reaches across NON-ADJACENT levels of description with no mechanism claimed at the levels in between - honest, layer-skipping causation, which we unpack a few paragraphs from now.
- **REFINES:** lineage - a pointer to the vaguer claim this one completes, so knowledge does not churn, it MATURES.

Lineage means the family history of a claim, the trail showing which earlier, rougher claim this sharper one grew out of.



The Roadmap to AGI and ASI

A pointer is simply a reference that links one thing to another, like a footnote saying "see this other entry".

And here is the stroke that makes the whole thing livable rather than utopian.

Utopian means an ideal that sounds lovely on paper but could never work in the messy real world.

ALMOST EVERYTHING IS OPTIONAL.

A newborn claim may say nothing but "pressing causes lighting" - no window, no modality, no mechanism.

We call it degenerate, and it is perfectly legal, because Causalontology is built for knowledge as it actually arrives: half-formed, promising, incomplete.

Here "degenerate" is not an insult; it is a technical word for the barest, most stripped-down version of a thing, one that has been reduced to its simplest possible form.

A partial claim is not an embarrassment.

A partial claim is a **STANDING INVITATION** - and in Part Three you will watch invitations like it summon whole communities, and one synthetic mind, to fill them.

A synthetic mind means an artificial mind, a thinking system built by people rather than born.



The Roadmap to AGI and ASI

The child at the switch has been carrying a Causal Relation Object in her head all along.

And here the standard must pause the story to show you four more things it holds - four answers the reified causal claim cannot give on its own.

A reified claim can say THAT pressing causes lighting.

On its own it cannot say four things a grown mind takes for granted.

It could not say at what LEVEL a claim lived - whether "the neuron fired" and "the ion channels opened" were two facts or one fact in two costumes.

The standard answers with the STRATUM: every claim may name the rung of description it stands on.

It could not say that two claims at different levels were THE SAME HAPPENING.

The standard answers with the BRIDGE: a signed declaration that one coarse occurrent simply IS a particular set of finer ones - not correlated with them, not caused by them, but identical to them, the same event described at two grains.

This is what lets a mind reason across scales without double-counting the world.

It could not say that two things were CONNECTED - that the output of this stands wired to the input of that.



The Roadmap to AGI and ASI

The standard answers with PORTS and CONDUITS:
typed interfaces borne by continuants, and typed
connections running between them.

A conduit that conveys its contents untouched is a
mere pipe; one that transforms what passes through
is, quietly, a computer – and the Perfect Wire
Test decides which, by asking whether an ideal
lossless wire could stand in its place without
changing the ends.

And it could not say that something ACTUALLY
OCCURRED – that beyond the eternal law "pressing
causes lighting" there was a real press, at a real
second, that a real mind witnessed.

The standard answers with the whole TOKEN TIER:
token individuals, token occurrences, state
assertions, and token causal claims, the machinery
of history rather than law.

This is the tier that gives a mind an episodic
memory, a perception, an act – a remembered
afternoon and not merely a rule about afternoons.

It is delivered, shipped, and testable – part of
the standard, not a promise for some later
edition.

Now the skip.

The old worry about level-crossing knowledge was
that it looked like a gap – a rude cause on one
floor with no visible staircase to the effect ten
floors down, and a commons forever nagging its
authors to fill in the missing steps.



The Roadmap to AGI and ASI

The standard turns that worry inside out with the SKIPS flag.

Consider cortisol, the stress hormone.

A social event - a humiliation, a threat, a loss - is a fact at the level of persons and meanings, roughly ten rungs above the level of molecules binding to DNA.

Yet in one physical step, the released hormone crosses all ten rungs at once and binds the DNA directly.

There is no mechanism at the intervening levels because there IS no intervening mechanism; the cause genuinely skips.

A Causal Relation Object may now say so, out loud, by setting its SKIPS flag.

And the absence of a mechanism at the levels between becomes a positive FINDING - a signed statement that the skip is real - rather than a permanent hole for the commons to nag about.

The map is allowed to record that the territory itself takes a shortcut.

What keeps this honest is Rule 7, the rule of BRIDGED reachability, which lets reachability run THROUGH bridges - so a chain of causes may be checked for consistency across several scales at once, a coarse link vouched for by the finer links it bridges to, without pretending every skip is a staircase.

Multi-scale mechanism consistency, made checkable.



The Roadmap to AGI and ASI

Picture a small child flipping a light switch on and off, delighted that she makes the light obey – that image, the child at the light switch, comes back again and again in this book as the simplest picture of learning cause and effect.

Causes: pressing.

Effects: lighting.

Temporal: within a second.

Modality: sufficient – usually.

Context: the power is on; the bulb survives.

We did not invent the atom of understanding.

We gave it a format, an identity, and a home outside any single skull – which, as the next chapters show, changes everything.

CHAPTER 6: A VERB-FIRST, NOUN-HOSTING ONTOLOGY

Before we go further, one plain-language word for the star of this chapter.

An ontology is just an agreed-upon list of the kinds of things a system believes exist, and the rules for how those kinds fit together.

Think of it as the official furniture catalogue for a world: it says what counts as a real item, and how the items may be arranged.

Every information system ever built has a bias so deep we stopped seeing it.



The Roadmap to AGI and ASI

They are all NOUN-FIRST.

A noun, remember, is a thing-word - a person, a place, an object.

A verb is a doing-word - press, fall, heal, shine.

Databases, which are organized electronic filing cabinets, model entities (the nouns, the things) first, then bolt relations between them afterward.

Object-oriented programming, a popular style of writing software in which the program is built out of little self-contained bundles called objects, literally puts objects first and demotes the verbs to methods (small built-in actions) that the objects own.

Knowledge graphs - maps that link facts together like dots joined by lines - are item-first; forms and filing cabinets and org charts are thing-first; the entire clerical history of computing is a museum of nouns.

Causalontology commits, calmly and completely, to the opposite.

IT IS A VERB-FIRST, NOUN-HOSTING ONTOLOGY.

Reality is what happens, and things are its participants.

Verb-first in four exact senses.

- The atomic building block is the occurrent - the happening, the process type, the verb (occurrent simply means a thing that happens, that unfolds in



The Roadmap to AGI and ASI

time, as opposed to a thing that just sits there and endures); occurrents are the vocabulary out of which every cause and every effect is built, and the standard enforces this mechanically.

- The fundamental unit reifies the master verb itself: a Causal Relation Object is a verb ABOUT verbs – causation, the one that says "brings about." (To reify, said REE-if-eye, means to take an idea that normally floats around as a mere relationship and turn it into a solid, nameable thing you can point at, store, and pass around; causation is the idea that one happening brings about another.)
- Even the hinge bends verb-ward: a disposition like fragility is noun-shaped (a disposition is a hidden tendency to behave a certain way under the right trigger – glass is disposed to shatter), but its entire meaning lives in the occurrent it is realized in – fragility is only fully itself in the shattering. (Realized here means the hidden tendency actually happening for real, coming out into the open.)
- And knowing is verb-first, because in this architecture knowledge enters by ACTING – by intervention – and by ACTIVE watching, in which the viewer identifies with the watched actor and internally simulates the action (the way you learn a dance from an instructional video, replaying the dancer's movements in your own body), rather than by PASSIVE watching, through which little or no knowledge enters – as the evidence ladder of Chapter 9 will make enforceable rather than merely poetic. (An intervention means you reach in and change something on purpose, then watch what follows; an internally simulated intervention is



The Roadmap to AGI and ASI

the imagined rehearsal of that same act, and the standard gives it its own evidence grade - simulation - a first-class slot distinct from raw observation and from physical intervention.)

But verb-first is not verb-only, and the discipline of that distinction is half the design.

Radical process philosophies dissolve objects into mist - everything a flux, nothing to hold.

Causalontology instead GOVERNS the verbs and HOSTS the nouns.

On the verb layer it is authoritative and unreplaceable: its causal objects are the last word.

On the noun layer it is a deliberately modest host: a thin backbone of continuants - things that endure (a continuant is a thing that persists as one and the same object through time, like a button or a person, in contrast to an occurrent, which is over the moment it is done) - that defers all domain depth to the specialist ontologies the sciences have already built.

Between them stands the realizable hinge: dispositions, functions, and roles, which inhere in things (inhere means to belong to, to be a property carried by) and discharge in happenings.

The button matters; the person matters; but they are understood THROUGH what they do and what happens to them.

Grammar, it turns out, was trying to tell us the architecture of understanding all along.



The Roadmap to AGI and ASI

The sentence needed a verb before it needed anything else.

CHAPTER 7: TWENTY-ONE KINDS OF THING

The entire universe of Causalontology contains exactly twenty-one kinds of object.

Not twenty-one hundred.

Twenty-one - ten that describe the world in the language of TYPES, seven that record what ACTUALLY HAPPENED, what a mind EXPECTS, and what one mind BELIEVES of another, in the language of TOKENS, and four that keep the paper trail of provenance - a vocabulary you can still learn before lunch and build a civilization's knowledge on after it.

(A word about the number.

Eight of these kinds are the irreducible core, the ones every causal claim leans on; nine more say what LEVEL a claim sits at, that two claims are secretly the SAME HAPPENING seen at two grains, that two things are CONNECTED, and that something ACTUALLY OCCURRED; one more, the managed cross_stratal_seam, records an honest LEAP across distant levels and earns its own pages shortly; and three more - the attitude, the predicted occurrence, and the prediction error - record what another MIND holds and what a mind EXPECTS, with the signed error of the expectation, and they too earn their own pages shortly.

Together they are a designed whole, complete and not a torso awaiting limbs.)



The Roadmap to AGI and ASI

(Type means a general kind - "pressing a button," in the abstract, everywhere and always.

Token means one particular instance - the button YOU pressed at 9:04 this morning, once, and never again.

The type tier holds the laws; the token tier holds the history.)

(Content means the actual claim or fact - the WHAT.

Provenance means the paper trail of where a thing came from and who is responsible for it - the WHO, WHEN, and ON WHAT GROUNDS.)

The content kinds are pure, immutable, and content-addressed.

(Immutable means once written it can never be changed, only added to or set aside.

Content-addressed, which we will unpack fully in the next chapter, means an object's name is computed from what it says, so identical content always gets the identical name.)

First, the TEN TYPE-TIER kinds - the world described as general kinds, laws, and structure.

- The OCCURRENT (prefix occurrent:): a process or event TYPE - `press_button`, `light_on`, `administer_dose` - the verbs of the world, and each may optionally be pitched at a named STRATUM, a stated level of description. (A prefix is just a short label at the front of a name that tells you which of the twenty-one kinds you are looking at.)



The Roadmap to AGI and ASI

- The CAUSAL RELATION OBJECT (prefix `causal_relation_object:`): the reified claim of Chapter 5 – the atom of understanding. It can carry a SKIPS flag, which we unpack in a moment: a signed declaration that this cause reaches across non-adjacent levels of description with no mechanism claimed at the levels in between.
- The CONTINUANT (prefix `continuant:`): a thing that endures – button, lamp, patient – the hosted nouns.
- The REALIZABLE (prefix `realizable:`): a disposition, function, or role – fragility, the key's function, the catalyst's power – the hinge, borne by a continuant, realized in occurrents. (A realizable is a latent capability that is carried by a thing and only shows up when an event lets it out.) The standard gives each realizable an identity-bearing label, a small repair with large consequences: one and the same bearer may hold many distinguishable dispositions – a molecule that is at once fragile, catalytic, and soluble – without the three collapsing into one indistinct blur.
- The STRATUM (prefix `stratum:`): a single level of description WITHIN a named scheme – the molecular level, the cellular level, the social level – a way of stating at what grain a claim is pitched. (A scheme is just an agreed ladder of levels, named so that two people mean the same rungs by the same words.)
- The BRIDGE (prefix `bridge:`): a cross-stratal identity map – the declaration that one coarse occurrent IS a particular set of finer ones, the



The Roadmap to AGI and ASI

very same happening seen at two levels at once. "A neuron firing" and "these ion channels opening" are not two events that correlate; they are one event described twice, and the bridge says so on the record.

- The CROSS-STRATAL SEAM (prefix `cross_stratal_seam:`): a managed jump across NON-ADJACENT levels of description - molecular straight to social, say - that records, in its `mechanism_status` field, whether the intervening mechanism is merely UNMODELED (it exists but is not drawn here) or genuinely ABSENT (there is none), with an optional drawn CHAIN of intervening steps and a HOME RULE placing it at the coarsest level it touches. It is the honest, first-class form of the SKIPS flag, and it earns its own pages shortly.
- The PORT (prefix `port:`): a typed interface borne by a continuant - the place where one thing can be joined to another, like the socket on a wall or the receptor on a cell. (Typed means the interface declares what may plug into it, so mismatched connections are caught rather than guessed.)
- The CONDUIT (prefix `conduit:`): a typed connection from one port to another - the thing that carries influence across. A conduit is a PIPE if it merely conveys what enters it unchanged, and a COMPUTER if it transforms what passes through in transit. Which one it is gets decided by the Perfect Wire Test: imagine an ideal, lossless wire in its place; if the two ends would then behave identically, the conduit was a mere pipe; if not, it was doing computation, and must be labelled as such.



The Roadmap to AGI and ASI

- The QUALITY (prefix quality:): a property TYPE - mass, colour, temperature, charge - the measurable characteristics a thing can bear.

Next, the SEVEN TOKEN-TIER kinds - not the law but the history, the record of what particular things actually did, and now also of what particular minds expect and believe.

These are LOCAL BY DEFAULT: a mind's episodic record is private unless it is deliberately published, because a history is nobody else's business until its owner says so.

The attitude, the predicted_occurrence, and the prediction_error are local by default for the same reason, and the attitude doubly so: an attribution of a mental state to an agent is among the most sensitive records this standard can carry - share deliberately, grade honestly, retract cleanly.

- The TOKEN INDIVIDUAL (prefix token_individual:): a single particular thing - not "patients" in general but THIS patient, this button, this cell, named once and for all.
- The TOKEN OCCURRENCE (prefix token_occurrence:): a particular happening at an absolute time - not "pressing a button" but the press that happened at 9:04:07 on the fourteenth, once. (Absolute time means a real timestamp on the real calendar, not a relative "later" or "before.")
- The STATE ASSERTION (prefix state_assertion:): a value holding over an interval - "the lamp was on from 9:04 to 9:11," "the temperature stayed at thirty-seven degrees through the night" - a fact

about how the world stood across a stretch of real time.

- The TOKEN CAUSAL CLAIM (prefix `token_causal_claim:`): a particular causing, this effect brought about by this cause on this occasion, under a stated covering law, and carrying its own token-level counterfactual – the but-for test: had this particular cause not occurred, this particular effect would not have followed. (A covering law is the general `causal_relation_object` this one instance falls under; the but-for test is the ordinary courtroom question, would the harm have happened anyway.)
- The ATTITUDE (prefix `attitude:`): what a particular mind CONTAINS – a record that a HOLDER (an agent modeled inside the ontology, never a cryptographic key) believes, desires, intends, knows, expects, or fears a CONTENT, where the content is a reference, by identity, to any content object at all – including another attitude, which is how "A believes that B believes X" becomes two honest records instead of a muddle. The content may be FALSE without the holder knowing, because an attitude records what a mind holds, not what the world did; it carries no strength of its own; and it earns its own pages shortly. (The philosophers' word for belief-like states is doxastic, from the ancient Greek *doxa*, opinion.)
- The PREDICTED OCCURRENCE (prefix `predicted_occurrence:`): a happening a stated PREDICTOR – a modeled agent or construct, never a key – EXPECTS: the occurrent type predicted, the window in which it is expected, held in exactly one dimension of time (the wall clock or the tick,

never both), and, optionally, the STRENGTH of the expectation, the predicted probability. A prediction is NOT a claim that the thing happened – it is the sibling of the token occurrence that has not (yet) happened – and it, too, earns its own pages shortly.

- The PREDICTION ERROR (prefix `prediction_error:`): the graded reckoning between a prediction and the world – it pairs a predicted occurrence with the token occurrence that answered it, or with nothing at all when nothing did, and carries the DISCREPANCY, a signed number read as actual minus expected: positive when the world exceeded the expectation, negative when it fell short, zero when the prediction was fulfilled. (Signed here means the number keeps its direction, plus or minus – the arithmetic sense of the word, not the cryptographic one.)

Finally, the FOUR PROVENANCE kinds are signed and add-only.

(Signed means each record carries a mathematical seal proving who wrote it, which we cover in Chapter 9.

Add-only means you can pile new records on, but never quietly erase old ones.)

- The ASSERTION (prefix `assertion:`): who claims a content object is true, on what evidence, with what strength and confidence, when. An assertion may cite its token-tier evidence directly, through an `evidenced_by` link, pointing at the very occurrences that ground it – and the evidence ladder is three-runged: intervention outranks observation, and observation outranks simulation,



The Roadmap to AGI and ASI

the bottom grade for an imagined or internally simulated intervention that was rehearsed in the mind but never physically performed.

- The ENRICHMENT (enrichment:): who added one alias, one participant, one taxonomy link - one word - to the commons, and now also who declared that one occurrent subsumes another (occurrent_subsumes) or is a part of another (occurrent_part_of). (An alias is an alternative name for the same thing; a taxonomy link is a "this is a kind of that" connection, like saying a robin is a kind of bird; the commons is the shared public pool of knowledge everyone contributes to and reads from.)
- The RETRACTION (retraction:): a source's honest withdrawal of its OWN earlier record. (A retraction is taking back something you yourself said, on the record.)
- The SUCCESSION (succession:): a cryptographic change of guard - this new key continues that old key's identity and responsibilities. (A succession is a signed handover, like an outgoing office-holder formally passing the keys and the duties to the next one - here the "key" is a digital key, explained in Chapter 9.)

Now the load-bearing decision, the one everything else in this book stands on.

CONTENT IS SEPARATED FROM PROVENANCE, UNIFORMLY, WITH NO EXCEPTIONS.

The claim "pressing causes lighting" is one immutable object - the same object no matter who utters it, in any decade, on any continent.



The Roadmap to AGI and ASI

That YOU believe it, with YOUR evidence, at YOUR confidence – that is a different object entirely: your signed assertion about it.

A thousand sources may assert the same claim: one content object, a thousand signatures around it, and the claim grows STRONGER with each independent voice.

Two laboratories may flatly disagree: both assertions stand, signed, side by side, and the disagreement itself becomes queryable knowledge – surfaced, never averaged away, never silently resolved by whoever edited last.

(Queryable just means you can ask the system about it and get a straight answer.)

Ask any veteran of collaborative knowledge systems what kills them, and they will tell you: edit wars, anonymous vandalism, contested cells.

(A contested cell is a single spot in a shared document that two people keep overwriting with rival versions.)

Causalontology has no contested cells.

There is nothing to fight over, because the design gave every disagreement its own home.

THE LAST THREE SEAMS: TIME, THE LEAP, AND THE LIVING LAW

A data structure meant to be the missing link between symbolic structure and native dynamics must answer three demands the object kinds alone



The Roadmap to AGI and ASI

cannot, and all three name the same fault line: the seam between symbolic STRUCTURE, which the kinds hold beautifully, and native DYNAMICS, the living law and the lived time that structure is supposed to be about.

The standard closes that fault line in three moves, and it does so under its own rules, in the open, with the arithmetic of identity left untouched.

Each move obeys the founding discipline.

An object's name is the fingerprint of its content, computed by writing its fields in one canonical form (under RFC 8785, the internet standard that makes JavaScript Object Notation (JSON) serialization deterministic) and hashing them with the Secure Hash Algorithm, 256-bit (SHA-256).

The fingerprint is taken over the fields that are PRESENT, and over nothing else.

A record that has no need of one of these three concerns simply does not carry it, and its name is exactly the name its content earns - nothing is re-baselined, nothing is renamed under anyone's feet.

One hundred and thirty-seven public conformance vectors hold the whole of it to the byte, green in the reference binding written in the Python language.

The first move lets time be recorded as it is lived.



The Roadmap to AGI and ASI

Every temporal window in the standard – every interval over which a state holds or a delay unfolds – carries a unit.

For a great deal of the world that unit is the wall clock, seconds and their multiples, and that is exactly right.

For the rest of it, it is exactly wrong.

A mind counting settling steps toward a decision, or a cerebellar circuit timing an inter-stimulus interval by the beat of its own updates, does not live in seconds; it lives in STEPS.

Force such a process into wall-clock seconds and you have not recorded it – you have distorted it, inventing a precision the phenomenon never possessed.

So the temporal window's unit enumeration carries, beside the wall-clock units, the value "ticks": a discrete, dimensionless ORDINAL step with no mapping to the wall clock at all.

A window measured in ticks is ordered by plain integer comparison of tick counts – step two comes after step one, and that is the whole of what the ordering means.

Ticks and seconds are held as DISJOINT DIMENSIONS: a tick window and a seconds window never overlap, a delay of the one is never inside a window of the other, and asking the standard to convert ticks into seconds is not answered approximately – it is REFUSED, as the category error it is.



The Roadmap to AGI and ASI

(A category error is treating one kind of thing as though it were another kind entirely, like asking how many kilograms an hour weighs.

The unit is identity-bearing: a record timed in ticks has a different fingerprint from an otherwise identical record timed in seconds, because it says something genuinely different about the world.

Conformance vectors V108 through V111 hold this behaviour to the byte.

For the argument of this book the tick is small and decisive: it is where the structure records native, lived time on the phenomenon's own terms, rather than translating it into a clock the phenomenon never consulted.

The second move manages the leap between levels, and it turns a confession into a record.

A Causal Relation Object may carry a bare boolean flag, SKIPS, admitting that a cause reached across NON-ADJACENT levels of description - molecular straight to social, say - with no mechanism claimed in between.

That flag is honest as far as it goes, but it cannot tell apart the two very different reasons a jump might skip.

Sometimes the intervening mechanism EXISTS and is simply not drawn here - known to be present, deliberately left off this particular map for brevity or because it belongs to another discipline.



The Roadmap to AGI and ASI

Sometimes there is GENUINELY NO intervening mechanism to draw – the jump is primitive, a brute adjacency of levels with nothing between them.

A single boolean cannot separate honest abbreviation from honest ignorance, and that difference is the whole of scientific candour.

So the standard makes the skip a first-class managed object: the `cross_stratal_seam`.

It names its two ends – a source occurrent and a target occurrent, whose stratum ordinals (the numbered rungs of a named level-ladder) are the two ends of the jump.

It carries a `mechanism_status` that states, on the record, whether the intervening mechanism is "unmodeled" (it exists; it is simply not drawn here) or "absent" (there is none) – the very distinction the boolean could never make.

It may carry a `CHAIN`: an optional sequence of intervening occurrents, each strictly between the endpoints on the same ladder and strictly climbing, that actually `DRAWS` the pathway the seam spans – and a seam that draws its chain may not also declare its mechanism absent, because a mechanism you can draw is, by definition, not absent.

And it obeys a `HOME RULE`: a seam belongs to the `COARSEST` stratum it touches, the end bearing the greater ordinal, so that the platform's layer-to-stratum binding has one unambiguous place to file it and check it.



The Roadmap to AGI and ASI

The bare boolean is not abolished; it is SUBORDINATED – a bare skipping relation remains valid and simply reads as the coarse, mechanism-status-unspecified form of the same idea, the seam not yet elaborated.

The standard's Semantics Rule 22 and its companion Algorithm F give the seam its meaning and its checking procedure, and conformance vectors V112 through V116 hold it to that meaning.

Here the missing-link argument gains its sharpest instrument: a leap between levels is no longer a blank the reader must interpret charitably, but a managed record that states, in whole words, exactly what it does and does not know about the space it crosses.

The third move is the quietest, and for the central claim of this book the most important, because it closes the seam the whole manuscript is named for – the seam between symbolic structure and native dynamics – and it closes it by IDENTITY.

A conduit, recall, is a typed connection that carries influence from one port to another, and the Perfect Wire Test sorts it into a mere pipe or a genuine computer according to whether it transforms what passes through.

But a conduit that computes is standing in for something: a real rate law, a real runtime signal, a native dynamical process that is the actual doing of the transform.



The Roadmap to AGI and ASI

Structure can say THAT a conduit computes; the question is whether it can say WHAT law does the computing.

It can: the conduit carries an optional realized_by field, a scheme-qualified REFERENCE, by identity, to the native law or signal that realizes its transform – the rate law a computational conduit stands for, the runtime signal a transmissive conduit relays.

The word REFERENCE is load-bearing and deliberately modest.

The standard does not swallow the law; it does not embed the differential equation or the signal trace inside the record.

The native dynamics stay native, living wherever they live, in whatever engine actually runs them.

All the standard records is the BINDING – the signed, content-addressed pointer that says "this structural conduit is realized by that native law," the two joined by the same fingerprint arithmetic that joins everything else in the commons.

Leaving it unbound is entirely legal and identity-preserving; a conduit with no realized_by is a complete and valid conduit.

But where the binding is present it is identity-bearing, because a conduit that names its law is making a stronger claim than one that does not.

Conformance vectors V117 through V119 hold the field to account.



The Roadmap to AGI and ASI

This is the missing link made literal: for the first time the symbolic record of a connection can point, by name and without ambiguity, at the native law that gives it its power - structure and dynamics clasped by a single fingerprint, neither one pretending to be the other.

None of this bends the standard's discipline; all of it is that discipline.

Thirty decisions stand LOCKED - written down, numbered, changeable only by a public, versioned process - and among them stands the plain rule that the standard's enumerations are CLOSED.

The token tier and its local-by-default privacy, the minimalism of required fields, the whole-word naming of Principle P7, and content-addressed identity all stand together as one design, each guarding the others.

And they stand in whole words.

Principle P7 - the standard's refusal of abbreviations, the rule that a reader should never have to decode a cipher to read a claim - governs every name here as strictly as it governs any other.

The managed seam is cross_stratal_seam, spelled through; its honesty field is mechanism_status, with the values unmodeled and absent written out; the temporal unit is ticks; the conduit's field is realized_by.

Not one of them is a stub or an acronym.



The Roadmap to AGI and ASI

The names are as legible as the design they carry, which is the only way a glass box stays glass.

Step back and the shape of it is a single argument, made three times.

A data structure meant to be the missing link between symbolic structure and native dynamics must be able to record time as it is actually lived, to be honest about the jumps it makes across the levels of the world, and to bind its symbols to the living laws that give them force - and it must do all of this without ever betraying the identity arithmetic that lets two strangers on opposite sides of the planet mint the same object and merge without a meeting.

The standard does exactly that, and only that.

The tick lets lived time in; the `cross_stratal_seam` makes the leap between levels an honest, managed record rather than a silent gap; and `realized_by` clasps the structural record to the native law that realizes it.

THE OTHER MIND AND THE EXPECTED WORLD

The three moves above complete the fabric that joins symbolic structure to native dynamics; the weave is whole, and nothing that follows reopens it.

The two kinds that follow are different in kind - not a mending of the weave but two new rooms built onto the house, and each is admitted the only way this standard permits.



The Roadmap to AGI and ASI

Principle P6, remember, forbids furnishing the ontology from the armchair: no construct enters unless a recorded, specific failure forced it.

So a kind is admitted only when a real build hits a real, recorded wall - and here the wall came twice in a single day, two findings from one konnectome build, answered by two moves that enter together, each additive, each identity-preserving, each in whole words.

The first move is the DOXASTIC one, and it begins with a mind trying to be honest about another mind.

(Doxastic, from the ancient Greek word doxa, belief or opinion, simply means having to do with believing.)

A working synthetic mind, modelling a companion, worked out that its companion held a belief that was FALSE - and held that discovery inside a further layer, its own model of the other's model.

It could think the thought privately.

It could not WRITE it, because among the other kinds there was no object whose job is to say "this agent's mind contains that content."

And the nearest kind, the assertion, is constitutionally the wrong shape for the job, in three exact ways.

An assertion names the SOURCE that signs it - a cryptographic key - and a key is not a mind: the record needed a HOLDER, the modeled agent whose head is being described, and the signer of a



The Roadmap to AGI and ASI

record and the holder of a belief are different things, forever (one agent can sign, on its own key, its model of what a second agent believes).

An assertion stakes its signer's word on the content being TRUE, so a false belief written as an assertion would inject into the commons the very falsehood it meant to report.

And an assertion cannot NEST: "A believes that B believes X" is a two-story thought, and assertions have one floor.

So the standard answers with the ATTITUDE: a content object that records a HOLDER, an attitude type from a closed list - believes, desires, intends, knows, expects, fears - and a CONTENT, a reference by identity to any content object at all, including another attitude, which is exactly what makes the two-story thought two honest records, the outer one pointing at the inner one by name.

Its governing law, Semantics Rule 25, is a quarantine, and the quarantine is what makes the whole thing safe: AN ATTITUDE'S CONTENT IS NEVER TREATED AS ASSERTED.

No reasoning rule, no conflict rule, no reachability or classification rule may read the content of an attitude as a fact about the world; content that contradicts the actual record raises no conflict at all - because that mismatch is not an error in the commons, it IS a false belief, first-class and shareable at last.

A mind can now be RIGHT, on the public record, about another mind being WRONG.



The Roadmap to AGI and ASI

The attitude carries no strength of its own - Principle P4 holds - and it is asserted, graded, corroborated, and retracted through the ordinary provenance layer like any other content, so the holder written inside the record stays forever distinct from whichever sources sign about it.

This is the substrate of THEORY OF MIND - the capacity to model what other minds believe, desire, and intend, long named among the prerequisites of general intelligence - and the social cognition of Chapter 11 now has an object kind of its own under its feet.

Conformance vectors V128 through V135 hold it to the byte: the true belief, the false belief left unpunished by the conflict rules, the desire, the intention, the nested two-story attribution, the closed list refusing an unknown attitude, the refusal of strength, and the assertion about an attitude whose signer and holder are visibly two different identities.

The second move is the PREDICTIVE one, and it begins with a number that had no home.

The second wall stood in the same builder's predictive loop, the oldest loop a learning mind runs: expect something, watch what actually comes, and compute the difference - the graded surprise a COMPARATOR reports, the "better or worse than expected?" signal the neuroscientists call a reward prediction error.

(A comparator is any component whose whole job is to hold an expectation against an outcome and report the difference between them.)



The Roadmap to AGI and ASI

The loop could always compute the number.

It could not RECORD what the number was ABOUT, because the token tier held only what actually happened – and a forecast minted as a token occurrence is byte-identical to a report of fact, the same fingerprint for "I expect the light within a second" and "the light came on within a second," which is exactly the confusion between a prediction and a claim of fact that a causal commons must never permit.

So the standard answers twice, with a pair of kinds.

The PREDICTED OCCURRENCE records the expectation itself: the occurrent type predicted, the PREDICTOR – a modeled agent or construct inside the ontology, never a cryptographic key – the window in which the happening is expected, and, optionally, a STRENGTH, the predicted probability.

That strength is part of the predicted CONTENT itself – the grade of the expectation is what the expectation SAYS, not somebody's later evaluation of it – and so it is the one deliberate, narrow, openly recorded exception to Principle P4's rule that content carries no strength, made by extension and confined to this single kind.

The governing law, Semantics Rule 24, keeps the new kind honest at both edges: A PREDICTION IS NOT A REPORT – it satisfies no rule that requires a token occurrence, and a forecast and a report of the same content mint DIFFERENT names – and its window must carry exactly ONE dimension of time, the wall clock or the tick, never both and never



The Roadmap to AGI and ASI

neither, because Rule 23 has already ruled that the two dimensions never compare.

And the PREDICTION ERROR gives the comparator's number its home: it pairs the predicted occurrence with the token occurrence that answered it - or with nothing at all, when nothing arrived and the surprise was silence - and carries the DISCREPANCY, the signed, graded value read as actual minus expected: positive when the world exceeded the expectation, negative when it fell short, zero when the prediction was fulfilled.

When the observed occurrence is present it must instantiate the very occurrent the prediction named - a mismatched pairing is refused - so an error can never quietly grade a forecast against the wrong event.

The predictive loop - expect, observe, compare - is thereby lifted out of the runtime and into the record: the transient number becomes a first-class object, signable, gradable, inspectable, shareable, auditable by any stranger with the arithmetic.

Conformance vectors V120 through V127 hold it to the byte: the tick-window forecast, the wall-clock forecast whose strength is identity-bearing, the fulfilled error at exactly zero, the unfulfilled error whose observed is absent, and the refusals - the missing predictor, the doubled dimension, the missing discrepancy, the mismatched pair.

Both moves end the way every move in this standard ends: with the identity arithmetic untouched and witnesses posted.



The Roadmap to AGI and ASI

Every record minted under an earlier version keeps its identifier byte for byte – not assumed but WITNESSED, by a conformance vector that re-pins two frozen earlier identifiers, unchanged, under the current implementation (V136) – and abbreviated forms of the new names are refused as firmly as every abbreviation before them (V137), because Principle P7 governs here too: attitude, predicted_occurrence, prediction_error, holder, attitude_type, content, predictor, discrepancy, all spelled through, all legible, all glass.

Step back one last time and the shape of the tier is worth saying aloud: the token tier, which gave a mind a HISTORY, now gives it an EXPECTED world and a model of OTHER minds – the was, the will-be, and the believed – each signed, each graded, each quarantined exactly as its nature demands.

Twenty-one kinds, one hundred and thirty-seven vectors, one identity arithmetic.

The missing link is finished – and every seam it closes, it closes in public.

CHAPTER 8: IDENTITY WITHOUT A COMMITTEE

Here is a small miracle of engineering, and it deserves to be savored slowly.

In Causalontology, AN OBJECT'S IDENTITY IS THE FINGERPRINT OF ITS CONTENT.

A fingerprint here means a short, unique code computed from a chunk of data – the technical word is a hash.



The Roadmap to AGI and ASI

A hash is what you get when you feed any amount of information through a fixed mathematical blender: out comes a compact string of characters that acts like that information's fingerprint.

Change the input by even one letter and the fingerprint comes out completely different; feed in the identical input and you always get the identical fingerprint.

The procedure fits in one breath.

Take the object's identity-bearing fields; write them in one canonical byte-for-byte form (canonical means the single officially-agreed way of writing something down, so that everyone who writes the same meaning produces the exact same characters, letter for letter - an internet standard, RFC 8785, exists precisely to make JavaScript Object Notation (JSON) serialization deterministic).

(RFC stands for Request for Comments, which despite the modest name is a numbered, public, official internet standard; JSON stands for JavaScript Object Notation, a simple and extremely common text format for writing down structured data; serialization means turning data into a flat string of text you can store or send; deterministic means it always comes out exactly the same way, with no wiggle room.)

Then hash those bytes with SHA-256 (the Secure Hash Algorithm, 256-bit), the workhorse fingerprint function of modern cryptography.

(SHA-256 stands for Secure Hash Algorithm, 256-bit; cryptography is the science of protecting and



The Roadmap to AGI and ASI

verifying information with mathematics; a byte is a tiny unit of computer data, one small character's worth.)

The identifier is the kind's prefix plus that fingerprint.

Sixty-four hexadecimal characters of destiny.

(Hexadecimal is just a way of writing numbers using sixteen symbols - the digits 0 through 9 plus the letters a through f - a compact everyday format for computer fingerprints.)

Feel what follows.

A physician in Nairobi and an engineer in Oslo, strangers forever, each writes down "pressing causes lighting" in their own systems.

THEY MINT THE SAME IDENTIFIER - necessarily, mathematically, without ever meeting.

(To mint an identifier means to generate its name.)

When their stores meet, the two claims are not duplicates to reconcile; they are the SAME OBJECT, already, and always were.

So merging two stores is set union.

(Set union is the simplest possible way of combining two collections: just throw everything from both into one pile, and anything that appears in both simply counts once - no negotiating, no deleting.)



The Roadmap to AGI and ASI

Not diff, not conflict resolution, not a committee, not a coordinator, not a vote.

(A diff is a line-by-line comparison that hunts for differences to settle.)

Union.

The computer scientists have a name for structures with this property: a Conflict-free Replicated Data Type, a CRDT – the gold standard for knowledge that must converge across an unreliable planet.

(A CRDT is a kind of shared data that many copies can edit independently, even while disconnected, and still be guaranteed to end up identical once they sync back together, with no clash to resolve; converge means all the copies settle onto the same final answer; a replica is one such copy of the shared store, living on someone else's machine.)

Most systems fight for years to bolt CRDT properties on.

Causalontology is a CRDT BY CONSTRUCTION – the immutability of content, the additivity of provenance, and content-addressed identity conspire so that any two replicas, merged in any order, arrive at the same store, provably, with one small deterministic rule for the single edge case (a cycle formed across replicas) that every implementation resolves identically.

(An edge case is an unusual corner situation; a cycle here means a loop, such as A causes B, B causes C, and C loops back to cause A, which only appears when two copies are joined.)



The Roadmap to AGI and ASI

And there is a quieter gift inside the fingerprints.

TAMPER-EVIDENCE IS FREE.

(Tamper-evidence means any meddling leaves an obvious mark, the way a broken seal on a medicine bottle instantly shows someone got in.)

Alter one byte of a claim and its identity no longer matches its content; forge a record and its signature fails.

In Part Four you will see a test where a single corrupted byte in a merged bundle is caught by arithmetic alone - no guard, no administrator, just mathematics standing watch.

Knowledge that carries its own audit.

(An audit is an independent check that everything is genuine and adds up.)

Identity without an authority.

Merging without a meeting.

This is infrastructure the way bridges are infrastructure: boring at the exact points where boring is a superpower.

CHAPTER 9: EVERY WORD HAS AN AUTHOR

Now the soul of the system - the part that makes it fit for a world full of liars, saints, and everyone in between.



The Roadmap to AGI and ASI

EVERY MUTABLE DATUM IN THE COMMONS IS SIGNED.

(Mutable means changeable; a datum is a single piece of data; so this says every changeable little fact in the shared pool carries a proof of who put it there.)

Not just the big claims.

Every alias in every language, every added participant, every taxonomy link, every retraction, every change of cryptographic guard - each one is a record signed with Ed25519, a fast, widely trusted digital-signature scheme, verifiable by anyone, forgeable by no one.

(A digital signature is the electronic equivalent of a tamper-proof, unforgeable handwritten signature: it uses a private secret key that only the author holds to seal a record, and a matching public key that anyone can use to check the seal, so the whole world can confirm who wrote something while nobody can fake it; Ed25519 is one specific, widely trusted method for doing exactly this.)

Every word has an author, and every author is checkable arithmetic.

Consider what this quietly forbids.

There is no anonymous edit.

There is no framing - no writing lies into a rival's mouth - because a record claiming to come from a source must verify against that source's own key, and the mathematics refuses all impostors.



The Roadmap to AGI and ASI

There is no silent rewriting of history: a source may retract its OWN record – and only its own, or its lineage's – and the retraction hides the record from everyday views while history keeps it, marked, forever.

(A source's lineage is the chain of keys that have officially succeeded one another, as described back in Chapter 7.)

Honest people get an exit; the record gets its integrity; both, not either.

And when an institution rotates its keys – as all institutions must – a signed succession carries the old identity's responsibilities to the new key, so reputation and accountability survive the calendar.

(Rotating keys means periodically retiring an old secret key and switching to a fresh one, a routine safety habit, like changing the locks.)

Then, on top of the signatures, the architecture takes its boldest philosophical stand.

EVIDENCE IS GRADED, AND ACTING OUTRANKS WATCHING.

Every assertion declares how its source knows: by INTERVENTION – the source acted and observed what followed; by OBSERVATION – the source watched without touching; by DERIVATION – reasoned from other claims; by HUMAN HINT; or merely IMPORTED from legacy data.

(Legacy data means old information inherited from earlier systems, of unknown reliability.)



The Roadmap to AGI and ASI

Intervention sits at the top, and it is not there by sentiment.

One deliberate action cuts through a thousand correlations - smash the barometer, silence the rooster - because coincidence cannot survive your own hand on the switch.

(Here is the heart of it in two everyday words.

Correlation means two things tend to happen together - the barometer needle drops and a storm arrives, the rooster crows and the sun rises.

Causation means one thing actually makes the other happen.

The rooster does not cause the dawn: silence the rooster and the sun still rises, which is exactly how acting - reaching in and changing one thing - reveals the true cause that mere watching could never separate from coincidence.)

This is rung two of Pearl's ladder, not as aspiration but as a FIELD IN THE SCHEMA - a first-class, machine-readable property of every piece of knowledge in the commons.

(Pearl's ladder is a well-known three-rung picture of reasoning from the computer scientist Judea Pearl: rung one is seeing, rung two is doing, rung three is imagining; a schema is the blueprint that lists what fields a record is allowed to have; a field is one labeled slot in that record, and machine-readable means the computer itself can read and act on it, not just a human.)



The Roadmap to AGI and ASI

Downstream, every consumer chooses its own trust policy over these grades and these authors – the commons never forces consensus, it furnishes evidence.

(A consumer here is any person or program that reads from the commons; a trust policy is that reader's own rule for which sources and which grades of evidence it chooses to believe.)

Signed words.

Graded knowing.

Honest exits.

Doing above watching, engraved in the data model itself.

That is the missing link, described.

Now watch what it enables.

THE ARCHITECT'S TOUR

One object, cradle to maturity – the whole system in a single biography.

We are about to follow a single piece of knowledge from the moment it is born to the moment it retires, the way you might follow one person through a whole life story.

Every architecture (the overall design of how a system is built and how its parts fit together) should survive this test: follow one citizen through its entire life and see whether the constitution holds at street level.



The Roadmap to AGI and ASI

By "citizen" we mean one small object in the system, treated like a member of a community with rights and a life of its own.

By "the constitution holding at street level" we mean the system's promises actually working in ordinary daily use, not just on paper.

Meet our citizen at the moment of conception.

A field engineer named Amara, testing irrigation valves in a greenhouse outside Nakuru, notices that opening valve three is followed - reliably, she thinks - by pressure loss in loop seven.

She does not write a report.

She mints two occurrents - open_valve_three, pressure_drop_loop_seven - and one degenerate Causal Relation Object.

An "occurrent" is a thing that happens, an event, as opposed to a thing that just sits there existing; opening a valve is an occurrent.

To "mint" one is simply to create it and give it an official identity, the way a mint stamps a new coin.

A Causal Relation Object is just a small tidy record that says "this cause leads to this effect."

A "degenerate" Causal Relation Object is a bare, bare-bones one: it names the cause and the effect and nothing else yet - causes, effects, and no further detail.



The Roadmap to AGI and ASI

BIRTH.

The instant the object exists, it has its name: the canonical bytes of its identity-bearing fields, fingerprinted; `causal_relation_object:` and sixty-four hexadecimal characters.

Let us unpack that slowly.

"Canonical form" means one single agreed-upon way of writing something down, so that the same meaning always produces the exact same characters, with no wiggle room.

The "canonical bytes" are the object's important fields written out in that one agreed way, as raw data.

To "fingerprint" (also called to "hash") is to run those bytes through a standard recipe that boils any input down to one short, fixed-length string of characters; the same input always gives the same string, and any change to the input gives a wildly different string - just as no two people share a fingerprint.

"Hexadecimal" just means counting with sixteen symbols (the digits 0 through 9 plus the letters a through f) instead of the usual ten, which is the normal way computers print such fingerprints.

So the object's name is the prefix `"causal_relation_object:"` followed by a sixty-four-character fingerprint of what it says.

Nobody issued this name.



The Roadmap to AGI and ASI

No registry approved it.

Because the name is computed straight from the content, this is called content-addressing: you find a thing by a fingerprint of what it contains, not by an address someone handed out.

If a hydraulics professor in Porto had written the same claim last spring - same causes, same effects - Amara's object IS his object, and the commons already knew it before either of them did.

The "commons" here means the shared public pool of knowledge that everyone contributes to and draws from, like a village common where all the animals graze.

The quiet magic is this: because the name comes from the content, two people who describe the exact same thing land on the exact same name, automatically, without ever meeting.

FIRST WITNESS.

Amara asserts her claim: evidence type OBSERVATION - she is honest; she watched, she has not yet experimented - confidence modest, timestamp, her Ed25519 signature.

To "assert" a claim is to publicly stand behind it and put your name on it.

Marking the evidence as OBSERVATION means she is being honest that she only watched it happen; she has not yet run a real experiment.

A "timestamp" is simply the recorded date and time.



The Roadmap to AGI and ASI

An "Ed25519 signature" is a digital signature made with Ed25519, a specific, widely trusted digital-signature method; a digital signature is a piece of unforgeable math that proves a particular person, and only that person, stood behind this exact record - like a wax seal that cannot be faked and breaks if the letter is altered.

The claim now has one voice behind it, and the voice is checkable arithmetic.

That last phrase matters: because the signature is math, anyone can verify for themselves that Amara really signed it - no trust in a middleman required.

APPEARANCE ON THE FRONTIER.

That night, the store's gap scan finds our citizen naked of window and modality: MISSING FIELD.

The "store" is the software that holds all these objects.

Its "gap scan" is an automatic sweep that looks for missing pieces of information.

A "window" here means the timing detail - how soon after the cause the effect shows up.

"Modality" means how strong the causal link is - for instance, whether the cause is enough on its own to guarantee the effect, or merely helps.

So the scan notices our young claim has no timing and no strength recorded yet: a MISSING FIELD.



The Roadmap to AGI and ASI

It joins the frontier - value-ranked, low at first, a small invitation among many.

The "frontier" is a public to-do list of open questions and gaps in the knowledge, ranked by how valuable it would be to fill each one.

Our claim joins that list, ranked low at first: a small, quiet invitation for someone to come finish it.

But greenhouse technicians keep querying it - is this thing supported yet? - and every read is counted, and demand compounds its value, and by Thursday our citizen has climbed the ranked list like a song climbing charts.

The commons has, without a meeting, decided this claim matters.

Nobody voted; the sheer weight of people asking about it pushed it up the list on its own.

INTERVENTION.

Friday, Amara does what the architecture was built to celebrate: she stops watching and starts DOING.

An "intervention" is when you stop merely observing the world and actively step in to change something on purpose, to see what happens - the difference between watching rain and turning on a hose.

This is the book's steady refrain: acting beats watching.

Forty deliberate valve cycles, gauge logged.



The Roadmap to AGI and ASI

Thirty-eight pressure drops, onset never later than ninety seconds; two exceptions, both when the bypass was open – and the exceptions are the treasure, because they are the CONTEXT field announcing itself.

The "context field" is the slot in the record for the surrounding conditions under which the cause-and-effect holds true.

The two odd cases were not failures; they revealed that the rule only holds when the bypass is closed – and that condition is exactly what the context field is for.

REFINEMENT.

She mints the mature form: same causes, same effects; temporal zero to two minutes; modality SUFFICIENT; context – bypass_closed; refines: our citizen's identifier.

A "refinement" is a more detailed, grown-up version of an earlier claim that fills in the blanks the first version left empty.

"Temporal zero to two minutes" records the timing window: the effect follows within nought to two minutes.

Modality "SUFFICIENT" means the cause, by itself, is enough to bring the effect about.

The "refines:" field points back to the original claim's identifier (its fingerprint name), recording that this new record is a child of that one.



The Roadmap to AGI and ASI

This trail from child back to parent is called lineage - a family tree for facts.

A new object - the child - with lineage pointing home.

She asserts it: evidence INTERVENTION, strength 0.95, signed.

And on everyone's frontier, everywhere on Earth, the original gap CLOSES - not because a server pushed a notification, but because the gap was always just a fact about structure, and the structure changed.

Read that twice, because it is the whole trick: no central computer had to announce anything.

The gap was simply the absence of a detail; once the detail exists, the gap is gone for everyone, the same way a hole in a wall is filled the moment you set the brick - the wall and the brick.

CORROBORATION.

"Corroboration" means independent backing up: someone else, separately, confirms the same thing.

The Porto professor - remember him - finds the refined claim in his weekly pull.

A "pull" is when your software reaches out and fetches the latest updates from the shared pool, the way you pull new mail into your inbox.

He runs his own rig; his numbers agree; he asserts the SAME object - one more signature around the



The Roadmap to AGI and ASI

same identity, and corroboration is now a countable two, from two continents, two instruments, two lives.

Because the object's name comes from its content, his confirmation lands on the very same object, and now that object simply carries two signatures instead of one - support you can literally count.

CHALLENGE.

A pump manufacturer's engineer disputes the window - her data says five minutes, not two.

She does not edit anyone.

She CANNOT edit anyone; there is nothing editable.

Records here are immutable, meaning once written they can never be altered - only added to.

She asserts a rival refinement, signed; the store surfaces the disagreement as structure; and the frontier gains a conflict entry that is, in truth, a beautifully specified experiment waiting for a volunteer.

Instead of an argument, the disagreement becomes a clearly framed open question on the public to-do list - a precise experiment that anyone can step up and settle.

MECHANISM.

A "mechanism" is the story of how the cause produces the effect - the chain of gears in between, not just the fact that turning one turns the other.



The Roadmap to AGI and ASI

A doctoral student adopts the citizen for her thesis: WHY does the drop follow?

She decomposes - valve opens flow path, flow path lowers manifold pressure, manifold couples to loop seven - three finer Causal Relation Objects, threaded into the parent's mechanism field.

To "decompose" is to break one big step into its smaller sub-steps.

She breaks the single cause-and-effect into three finer cause-and-effect records and threads them into the parent's mechanism field, so the parent now carries its own inner workings.

Now the claim does not merely predict.

It EXPLAINS, three layers deep, and every layer is a first-class object someone signed.

"First-class object" means each inner step is a full, real, independently addressable citizen in its own right, not a scribbled footnote - so someone can sign, challenge, or refine any one of them on its own.

RETIREMENT WITH HONOR.

Years on, the greenhouse re-plumbs; loop seven is gone.

Amara retracts her original observational assertion - one signed record; her early voice steps back from default views; history keeps it, marked, forever.



The Roadmap to AGI and ASI

A "retraction" is a signed statement that says "I am withdrawing this," which quietly removes a claim from the normal, everyday views without ever destroying the record of it.

A "default view" is what an ordinary reader sees unless they deliberately go digging into the full history.

The claim objects themselves remain - immutable, as all content is - true forever about the world that was, lineage intact for the historians of knowledge, and for the auditors, and for anyone who ever asks how we came to know.

Walk back through the biography and count what never happened.

No committee named it.

No moderator arbitrated it.

No edit war touched it.

No anonymous hand moved it.

No merge meeting reconciled Nakuru with Porto.

To "merge" two collections is to combine them into one; here it happened automatically, with no meeting, because same-content records already share the same name.

And no step - not one - is closed to a synthetic participant; the gardener of Chapter 16 could have been Amara at every station of this life except the coffee she drank at dawn.



The Roadmap to AGI and ASI

A "synthetic participant" is a machine or software agent taking part instead of a human; every stage of this life story is open to one.

That is the constitution, holding at street level.

THE RULES OF THE ROAD

The thirteen semantic laws, told as the small human stories they encode – because every clause in a good constitution is a remembered injury.

"Semantic" means having to do with meaning.

A "clause" is one single provision or rule within a larger set of rules.

By "a remembered injury" we mean each rule exists because someone, somewhere, once got badly burned by its absence.

A schema says what a claim may LOOK like; the semantic rules say what it may MEAN, and each one exists because some failure mode of knowledge demanded it.

A "schema" is a template or blueprint that spells out which fields a record must have and what shape they take – like a form with labelled boxes.

A "failure mode" is a specific way something can go wrong.

THE WINDOW MUST MAKE SENSE – minimum_delay no greater than maximum_delay.

Here "minimum_delay" is the earliest the effect may appear and "maximum_delay" is the latest; the



The Roadmap to AGI and ASI

rule says the earliest time cannot come after the latest time.

Trivial?

Every biologist who has read "effect observed between day 14 and day 3" in a hurried table knows this rule is a memorial to real corpses of meaning; machines must be forbidden the nonsense humans typo.

THE UNITS ARE FIXED, TO THE SECOND - a month is 2,629,746 seconds for every mind on Earth.

Because "within a month" has started arguments in every laboratory and every legal department since calendars began, and a commons cannot converge on knowledge if two implementations disagree about admissibility arithmetic at the boundaries.

To "converge" is for many separate parties to settle onto one shared answer.

Two "implementations" are two different pieces of software built to the same standard.

"Admissibility arithmetic at the boundaries" just means the edge-case math of deciding whether something barely counts as "within a month" or not.

Boring numbers, chosen once, are how civilizations stop fighting about boring numbers.

NOTHING MAY CAUSE ITSELF THROUGH ITS OWN EXPLANATION - mechanism and lineage must not loop.



The Roadmap to AGI and ASI

In plain terms: a claim is not allowed to appear inside the list of gears that supposedly explain it, and a family tree of refinements is not allowed to circle back on itself.

A claim that appears inside its own mechanism is a snake eating its explanation; a refinement chain that circles is a family tree where someone is their own grandparent.

The rule is small; the diseases it prevents - circular justification wearing the costume of depth - are ancient.

"Circular justification" is proving A by pointing to B and then proving B by pointing back to A - a loop that looks like reasoning but proves nothing.

A REFINEMENT MAY ONLY FILL SILENCE - never change what the parent said.

That is: a more detailed version may fill in blanks the original left empty, but it may never overwrite or contradict what the original actually stated.

This is the architecture's deepest courtesy: your vague early claim is safe forever; those who come after may complete you, but the moment they CONTRADICT you, the standard strips their disguise and calls their claim what it is - a RIVAL, standing beside yours on its own evidence.

Knowledge matures without betrayal.

CONFLICT HAS A DEFINITION - same causes, same effects, compatible contexts, overlapping windows, preventive against positive.



The Roadmap to AGI and ASI

This spells out exactly when two claims count as truly clashing: they share the same cause and same effect, their conditions can co-occur, their time windows overlap, and one says the cause prevents the effect while the other says it produces it.

Not vibes; a computable test.

"Computable" means a machine can check it by a definite procedure, with no judgment call.

Because "these studies disagree" has meant everything and nothing in seminar rooms for a century, and a frontier cannot rank contradictions it cannot detect.

HIERARCHY MUST ADD UP - every cause-to-effect promise of a parent must be reachable through the gears of its mechanism.

A "hierarchy" here is a claim with sub-claims nested inside it, layer within layer.

The rule: if a big claim promises that pressing leads to lighting, you must be able to trace that promise step by step through its inner mechanism.

If your explanation, followed link by link, cannot get from the pressing to the lighting, then it is not an explanation yet - and the store will say INCONSISTENT to your face, which is more honesty than most reviewers manage.

This is the child at the light switch: the claim only earns the word "explanation" when the chain actually reaches from finger to bulb.



The Roadmap to AGI and ASI

And where members are missing, the verdict is INDETERMINATE – never guessed – because pretending to know is the one corruption a knowledge commons cannot survive.

"Indeterminate" means "not enough information to decide," and the system says so plainly rather than bluffing.

EVERY ENRICHMENT IN ITS LAWFUL PLACE – aliases carry language tags; participants attach to happenings; taxonomies attach to things; realizations attach dispositions to their discharge.

An "enrichment" is extra descriptive detail added to a record.

An "alias" is an alternative name, and its "language tag" marks which human language the name is in (English, Swahili, and so on).

"Participants attach to happenings" means the actors and objects involved get attached to events (reality is what happens, and things are its participants).

A "taxonomy" is a classification tree, the way biology sorts living things into kingdoms, families, and species; taxonomies attach to things.

A "realization" records a "disposition" (a latent tendency, like glass being fragile) getting "discharged" (actually acting out, like the glass finally shattering) – and those attach to that moment of discharge.



The Roadmap to AGI and ASI

The clerk's rules, yes - and archives run on clerks; ask any archivist what happens to collections where anything may be filed anywhere.

Each rule alone is a fence post.

Together they are the reason a store that accepts contributions from strangers, machines, and centuries can still be TRUSTED to mean the same thing to every reader: the rules are not restrictions on knowledge.

They are the load-bearing manners of a shared mind.

WHAT THIS IS NOT

A chapter of refusals, because a sales document you can trust must show you its brakes.

THIS IS NOT A CLAIM THAT WE HAVE BUILT AGI.

AGI is Artificial General Intelligence (AGI): a machine with the broad, flexible, all-purpose intelligence of a human being, able to turn its mind to almost any problem.

We have built a data structure, a standard, a store, and one small gardener; the perfect benchmark scores are on public training sets and say so every time; the third benchmark is a labeled bet.

A "benchmark" is a standard test used to measure and compare how well systems perform.

A "public training set" is the batch of practice problems that are openly published for anyone to



The Roadmap to AGI and ASI

study beforehand - so scoring well on it is honest but modest, because the questions were seen in advance.

The perfect scores mentioned throughout were earned by pure symbolic induction - that is, by working out explicit, human-readable rules from examples, with no Large Language Model (LLM) involved at all; an LLM is the kind of statistical text-prediction system that powers today's chatbots.

The "third benchmark" is a harder, unseen test, and calling it a "labeled bet" means it is an openly declared wager on future performance - no score claimed until measured.

If you hear anyone - including our own enthusiasm on a bad day - claim more, quote this page at them.

THIS IS NOT A TRUTH MACHINE.

The commons stores claims with evidence, not verdicts; it makes lying expensive and disagreement legible, and that is ALL it does - the verdicts remain forever the work of readers, communities, and the world's stubborn habit of answering experiments.

"Legible" here means plainly readable and out in the open, so anyone can see the disagreement clearly.

THIS IS NOT A RELIGION OF SYMBOLS - and not a demotion of learning.



The Roadmap to AGI and ASI

Gradient descent is one of the great instruments of the age; embeddings are a magnificent similarity organ; perception will be neural for the foreseeable future and should be.

"Gradient descent" is the standard training method behind modern machine learning: the program keeps nudging its settings a little downhill toward fewer mistakes, over and over, the way a ball rolls to the bottom of a valley.

"Embeddings" are a way of turning things like words or images into lists of numbers positioned in space, so that similar things sit close together and their nearness can be measured.

"Neural" refers to neural networks, software loosely inspired by the brain's web of connected cells, which learn patterns from mountains of examples; "perception" here means the machine's raw sensing, like recognizing what is in an image.

We claim one thing only: what the learning LEARNS about cause and effect deserves a signed, mergeable, inspectable home – geometry for guessing, ledgers for knowing, both.

"Geometry for guessing" is the number-space of embeddings, wonderful for hunches; "ledgers for knowing" is the signed, checkable record for settled cause-and-effect – and the point is you want both, each for what it does best.

THIS IS NOT VENDOR ANYTHING.

A "vendor" is a company selling a product; "vendor lock-in" is the trap where, once you have committed, you cannot leave without great pain.



The Roadmap to AGI and ASI

There is no company to be locked into; the license forbids nothing but ingratitude; the reference implementations are cribbable in twenty languages; and if the maintainers of this standard were hit by a bus tomorrow, the freeze, the vectors, and the archive are public - the whole point of a standard is that its authors are replaceable.

A "reference implementation" is a worked example of the software, published so others can copy and learn from it; "cribbable" means free to copy.

The "freeze" is the officially locked, unchanging version of the standard; the "vectors" are the public test cases every implementation must pass; the "archive" is the permanent public record - all three out in the open, so the project survives its founders.

THIS IS NOT FINISHED.

Token-level occurrences are DELIVERED - the whole token tier ships, tested and frozen; the bounty economy is still designed, not deployed, and remains the one large piece held back for a later edition; probabilistic strength aggregation across corroborating assertions is deliberately left to consumers rather than mandated by the standard; and somewhere in these pages is at least one error we will retract with the same signature discipline we preach - watch the record, not the rhetoric.

A "token" here is one particular instance rather than a general kind; "token-level occurrences" (real happenings at a real time, at that fine a grain) are first-class - the standard defines and



The Roadmap to AGI and ASI

tests them, so a mind can record what actually happened, not only what always happens.

A "bounty economy" is a planned system of rewards for filling knowledge gaps; it is designed on paper but not yet running, and is now the only major item still reserved for the future.

"Probabilistic strength aggregation across corroborating assertions" means the math of combining many separate confirmations into one overall confidence number - deliberately left for users to do themselves for now, until a major new version of the standard.

THIS IS NOT UTOPIAN.

"Utopian" means naively assuming everything and everyone will behave perfectly.

People will sign nonsense; machines will grade sloppily; institutions will try to game reputations - Chapter after chapter assumes it, and the architecture's whole posture is that of a harbor built for storms, not for brochures.

What remains, after all the refusals?

A missing organ, honestly described; a standard, honestly frozen; receipts, honestly public.

That is all - and if the age of infinite claims has taught you anything, you know how much "that is all" is.

THE ANATOMY OF SAFETY



The Roadmap to AGI and ASI

How the architecture behaves under attack - a guided dissection.

Safety chapters usually promise; this one demonstrates - each defense below is implemented, vectored, and part of the frozen suite.

"Vectored" means covered by a public test case; the "frozen suite" is that locked, unchanging set of tests - so each defense is not just claimed but provably checked.

THE FRAMING ATTACK, KILLED IN THE CRADLE.

A "framing attack" is putting words in someone else's mouth - forging a record that pretends a respected party said something they never said.

The nastiest move in any knowledge system: write lies into a rival's mouth - "Laboratory X asserts the bridge is safe."

Here the move is not forbidden; it is IMPOSSIBLE: a record claiming X as source must verify against X's public key, and forged records fail arithmetic before they reach any view.

A "public key" is the openly shared half of a digital-signature key pair; it lets anyone check that a signature genuinely came from its owner, while only the owner holds the secret half needed to make that signature.

So a forged record simply fails the math and never appears.



The Roadmap to AGI and ASI

The suite includes the tampering test; the tampered record dies in quarantine, every implementation, every time.

"Quarantine" is a holding pen for records that fail the checks – they are set aside, visible only to those who deliberately go looking, and trusted by no ordinary view.

THE MEMORY-HOLE ATTACK, INVERTED.

A "memory-hole attack" is quietly deleting inconvenient history so it is as if it never existed.

In editable systems, power deletes embarrassments.

Here, content is immutable and provenance is add-only: even a retraction – the one sanctioned goodbye – hides without erasing, and history views keep the record, marked.

"Provenance" is the full traceable record of where something came from and everything that happened to it; "add-only" means you can add to that record but never rub any of it out.

You may leave the stage.

You may not burn the transcripts.

THE IMPOSTOR SUCCESSION, CHAINED.

"Succession" is the handover of an identity from an old key to a new one – for example when someone updates their signing key over time.



The Roadmap to AGI and ASI

An "impostor succession" would be a thief trying to hijack that handover to steal an identity.

Stealing an identity by "rotating" its key requires the succession record to be SIGNED BY THE PREDECESSOR – the old king must anoint the heir; and a stranger's retraction of someone else's record is rejected by lineage check, vectored and frozen.

"Rotating" a key means swapping in a fresh signing key to replace an old one; the rule is that only the current key-holder can authorize the swap.

THE POISONED TAXONOMY, BROKEN DETERMINISTICALLY.

To "poison" a taxonomy is to corrupt the classification tree; "deterministically" means every computer resolves it by the same fixed procedure and reaches the identical result, with nothing left to chance.

Concurrent writers on a distributed commons can, innocently or maliciously, complete a cycle – dog subsumes mammal subsumes animal subsumes dog – that no single store would accept.

"Concurrent writers" are many people writing at the same time; a "distributed commons" is the shared pool spread across many computers rather than one.

"A subsumes B" means A is a broader category that contains B; the cycle above is nonsense because it loops back on itself, making dog a kind of dog through its own ancestors.



The Roadmap to AGI and ASI

Enforcing stores refuse the completing write with a reason; and when a cycle arrives via merge, every replica on Earth breaks it IDENTICALLY - the latest-timestamp record loses, ties by identifier - and the exclusion itself surfaces on the frontier as a repair invitation.

A "replica" is one of the many identical copies of the shared pool held on different computers; the point is that all of them, independently, snap the bad loop in exactly the same place.

Even the pathological case converges, and even its resolution is an invitation.

CLAIMS OF CONSEQUENCE, HANDLED WITH TIERS.

Some knowledge - dosages, structural tolerances - can hurt people while wrong.

"Tiers" are levels of organization, from the tightly controlled to the fully open.

The architecture's honest position: a single store may enforce policy gates; a federation may propagate suppression notices; a fully decentralized layer CANNOT globally delete, and pretending otherwise would be the lie that poisons everything else - so the safety spec says plainly, per tier, what can and cannot be promised, and consumers of consequence-laden domains are told to read with trust policies tuned accordingly.

A "policy gate" is a rule-based checkpoint that can block certain records.

A "federation" is a group of separate stores that agree to cooperate and share notices with one



The Roadmap to AGI and ASI

another; to "propagate" a notice is to pass it along the network.

A "fully decentralized layer" is the wide-open bottom tier with no one in charge, where nothing can be forcibly erased everywhere – and the standard admits this openly rather than pretending it can promise deletion it cannot deliver.

Safety by honest architecture, not by comforting fiction.

THE QUARANTINE, ALWAYS-ON.

Nothing unsigned touches a default view – not as moderation, as physics; the unsigned record is accepted into quarantine, visible to those who go looking, trusted by no view that did not ask for it.

The point is that this is not a human moderator making a judgment call; it is a built-in mechanical fact of how the system works, like gravity – unsigned records simply cannot reach the ordinary reader.

Add it up and notice the pattern: every defense is STRUCTURAL – arithmetic, immutability, determinism – and none of them requires a wise administrator awake at 3 a.m.

"Structural" means the safety comes from the shape of the system itself – unforgeable math, records that cannot be altered, and procedures that always resolve the same way – rather than from someone watching the door.



The Roadmap to AGI and ASI

This is the glass box: every wall is visible and checkable, nothing hidden behind a curtain of trust.

Harbors for storms.

Wisdom optional; the walls hold anyway.

INTERLUDE TWO

HOW STANDARDS WIN

There is a pattern in the history of technology so reliable it deserves to be called a law, and this book's whole strategy rests on it.

FORMATS OUTLIVE AND OUT-POWER PRODUCTS.

Before we go on, let us define the two words in that heading in plain language.

A product is a single thing a company makes and sells, like one particular web browser or one particular phone.

A format (also called a standard) is an agreed way of writing things down or exchanging them that everybody follows, so that different products can all work together.

Here is the pattern: the agreed format almost always outlasts and out-powers any single product built on top of it.

The web did not happen because someone built a better browser; browsers came and went like weather.



The Roadmap to AGI and ASI

The web happened because two shared formats gave every browser, every server, every decade the same thing to read.

The first was the HyperText Markup Language (HTML), which is the agreed way to write a web page so any browser can display it (HyperText simply means text with clickable links in it, and a Markup Language is a set of tags that label the parts of a document, like "this is a heading" or "this is a link").

The second was the HyperText Transfer Protocol (HTTP), which is the agreed set of rules for how a browser asks a distant computer for a page and how that computer sends it back (a Protocol is just a shared set of rules two machines follow so they can talk to each other).

And the moment those formats were common, ten thousand incompatible efforts became one compounding system.

The internet itself is the same story one layer down.

Underneath the web sits the Transmission Control Protocol and Internet Protocol (TCP/IP), the shared rules that let any computer on Earth send data to any other computer in small numbered packages that get reassembled at the far end.

TCP/IP was not the fastest networking technology of its day - it was the one everyone could AGREE on, and agreement ate speed for breakfast.

The Structured Query Language (SQL), the standard way of asking a database questions like "give me



The Roadmap to AGI and ASI

every customer in France" (a database is just an organized store of information, and SQL is the agreed wording for questioning it), outlived every database engine that first carried it.

The shipping container – a dumb steel box with agreed corner fittings – reorganized the physical economy of Earth more than any ship ever did, because every crane, truck, and dock in the world agreed on the same box.

Extract the pattern, because we are about to apply it.

- The winning layer is BORING: a format, not a feature – stable where products are restless.
- It is OPEN: anyone may implement it, so adoption compounds instead of fragmenting (open means no single company owns it or can lock others out).
- It is TESTABLE: conformance is a public fact, not a marketing adjective – the plug fits the socket or it does not (conformance means actually meeting the standard, and it can be checked, not just claimed).
- And it is MINIMAL: it standardizes the interface, not the imagination – HTML never dictated what a page must say (the interface is the shared boundary where two things meet and connect).

Now look at the artificial-intelligence landscape of 2026 through this lens, and the strangeness leaps out.



The Roadmap to AGI and ASI

The field is a war of PRODUCTS - model against model, lab against lab, weights against weights - with billions wagered on which black box will be largest next quarter.

By weights we mean the many billions of internal numbers a Large Language Model (LLM) adjusts during its training; an LLM is the kind of Artificial Intelligence, trained on enormous amounts of text, that powers today's chat assistants.

By black box we mean a system whose insides you cannot inspect: you see what goes in and what comes out, but not why.

Nobody is fighting for the layer that history says actually wins, because that layer barely exists: there is no common, open, testable format for the knowledge the products are supposed to have.

Every model is a browser with no HTML - a magnificent renderer of nothing shared.

Causalontology is aimed, with premeditation, at the boring layer.

The word Causalontology combines two ideas.

Causation means one thing actually making another thing happen, not just the two appearing together.

An ontology is an agreed catalogue of the basic kinds of things that exist and how they relate, so Causalontology is an agreed catalogue built around cause and effect.



The Roadmap to AGI and ASI

It is a format: twenty-one kinds of object, frozen.

It is open: attribution always, no profit no problem, twenty implementations and counting (attribution means you must credit the source; an implementation is one working piece of software that follows the standard).

It is testable: one hundred and thirty-seven public test vectors, pass or fail, no adjectives (a test vector is a fixed example with a known right answer, used to check whether an implementation behaves correctly).

And it is minimal: it standardizes how causal knowledge is SHAPED, signed, and merged - never what anyone must believe; the trust policy belongs to the reader, the way the browser belongs to the user.

One more property from the history books, and it is the one that should keep strategists up at night.

STANDARDS WARS END SUDDENLY.

For years, nothing - then a threshold of tooling and adoption crosses, and the question flips from "why would we use it?" to "how are we not on it yet?" - and the late-comers spend a decade paying integration debt to the early ones (integration debt is the extra, painful work of bolting yourself onto a standard everyone else already adopted long ago).

The tooling threshold, for this standard, has already been crossed: the installs are one line,



The Roadmap to AGI and ASI

the store is one command, and the conformance gates run in public (a conformance gate is an automatic public check that passes or fails a piece of software against the standard).

The adoption threshold is a decision - yours, reader, and your field's.

History does not repeat, but it rhymes in formats.

We wrote this one down so the rhyme could include causation.

PART THREE: THE ENABLER

How Causalontology powers the faculties of a general mind.

CHAPTER 10: CLIMBING THE LADDER BY DATA STRUCTURE

A data structure, named in this chapter's title, is simply an organized way of holding information inside a computer so it can be used - like the difference between papers dumped in a box and papers filed in labelled folders.

Return now to Pearl's ladder with new eyes, because we can finally say something the field has not been able to say.

(Pearl's ladder, introduced earlier in the book, is a three-rung picture of causal reasoning from the scientist Judea Pearl: rung one is seeing patterns, rung two is doing things to see what changes, and rung three is imagining what would have happened otherwise.)



The Roadmap to AGI and ASI

EVERY RUNG OF THE LADDER IS A FIRST-CLASS OBJECT
IN THIS ARCHITECTURE.

First-class object means a thing the system stores and handles directly in its own right, not something vaguely implied or hidden.

Rung one - association - lives in the humbler evidence grades: observational assertions, imported correlations, honest about what they are, useful for what they are, and NEVER confused with more, because the grade is written on the label.

An observational assertion is a recorded claim based on just watching, not acting.

A correlation is two things tending to appear together, which is not the same as one causing the other; an evidence grade is a stamp saying how strong the backing for a claim is.

Rung two - intervention - lives in intervention-graded assertions attached to Causal Relation Objects.

An intervention means deliberately doing something to see what changes, rather than just watching.

A Causal Relation Object is the central building block of this whole system: a single stored record that says "this causes that," together with all the details of how, when, and how strongly.

The child's hundred presses of the light switch, the laboratory's controlled trial, the robot's deliberate probe - each one is signed, each one carries strength and confidence measured from acting.



The Roadmap to AGI and ASI

Rung three - counterfactual - lives in the compositional machinery.

A counterfactual is a "what would have happened if things had been different?" question - imagining an alternative to what actually occurred.

Compositional simply means built by combining smaller pieces into larger ones.

That machinery includes refinement lineages that hold rival and successive versions of a claim (a lineage is a recorded family tree of a claim as it gets revised; successive means one version coming after another over time); context fields that say under which conditions a mechanism holds (a mechanism is the actual step-by-step way a cause produces its effect); and, in the reference cognitive architecture, quarantined REALITIES - observed, desired, expected, imagined, recalled.

Quarantined means kept safely walled off and separate, so one cannot contaminate another.

A cognitive architecture is the overall design of a thinking machine's mind; the reference one is the worked example that ships with this system.

Keeping these realities apart means a mind can run "what if the bulb had been dead?" through the same causal objects WITHOUT ever contaminating its record of what actually happened.

Understand what this means for the great stuckness of Chapter 2.

The ladder is not climbed by making models bigger.



The Roadmap to AGI and ASI

The ladder is climbed by giving every rung a FORMAT - a place where intervention knowledge physically differs from observational knowledge, where alternatives are structured rather than hallucinated, where the climb leaves an audit trail.

To hallucinate, for one of today's AI systems, means to make something up confidently and present it as fact; an audit trail is a step-by-step record that lets anyone check afterwards what happened and why.

Show me the bytes, we demanded.

A byte is the tiny basic unit of computer data, so "show me the bytes" means "show me the thing actually stored, not the hand-waving."

These are the bytes.

CHAPTER 11: THE FACULTIES OF A GENERAL MIND

A faculty, in this chapter, means a distinct mental power or ability, like planning or remembering.

What makes intelligence GENERAL rather than narrow?

(Narrow means good at one single task, like playing chess; general means able to turn its mind to many different kinds of problem, the way a person can.)

Not one giant capability, but a family of faculties that share one world: planning,



The Roadmap to AGI and ASI

imagining, transferring, explaining, remembering, cooperating.

Watch each faculty become an OPERATION OVER THE SAME TWENTY-ONE KINDS.

The twenty-one kinds are the twenty-one kinds of object that Causalontology defines; an operation over them is a task performed by working on those objects.

- **PLANNING** is search over causal chains: find a path of Causal Relation Objects from the state you have to the state you want, respecting temporal windows and contexts; the plan IS the chain, and every link is inspectable (a temporal window is the span of time within which a cause is allowed to produce its effect; inspectable means you can open it up and look).
- **IMAGINATION** is counterfactual composition: recombine known Causal Relation Objects into configurations never observed - what if this preventer were absent, this contributor doubled? - inside a quarantined imagined reality, then compare against expectation to mint surprise (to mint surprise means to generate a measurable signal of "that is not what I expected," which is how such a mind notices something new).
- **TRANSFER** - the crown jewel of generality - is analogy over relation structure: the mapping from water pressure to voltage is a correspondence between CAUSAL GRAPHS, not between surface features, and a mind whose knowledge is stored as structure can carry a whole solution across domains by mapping the structure alone.



The Roadmap to AGI and ASI

An analogy is seeing that two different situations share the same underlying shape.

A causal graph is a diagram of causes and effects drawn as dots joined by arrows; surface features are the shallow outward appearances, as opposed to the deep structure underneath.

- EXPLANATION is reading the mechanism field aloud: WHY does pressing light the lamp - because circuit closes, because current flows, because the filament glows - each "because" one level of the mechanism decomposition, already stored, just waiting to be spoken (decomposition means breaking one thing down into its smaller parts).

- AGENCY is choosing interventions: given goals and a causal store, select the action whose Causal Relation Object points from here to there - and every action taken mints fresh intervention-grade evidence, so acting and learning become the same loop.

Agency here means the capacity to act in the world on purpose to reach a goal.

- MEMORY CONSOLIDATION is refinement lineage: raw episodes distill into claims, claims refine as evidence accumulates, and the lineage from vague to precise is preserved - knowledge that matures without forgetting how it grew.

Memory consolidation is the settling of fresh experiences into lasting, organized memory; an episode is a single remembered happening.

- SOCIAL COGNITION is causal models of OTHER minds: their goals inferred from their



The Roadmap to AGI and ASI

interventions, their false beliefs tracked as their models diverging from yours – the same objects, aimed at agents.

Social cognition is the ability to understand and reason about other thinking beings; an agent is any actor with goals, whether a person or a machine.

One substrate.

A substrate is the single shared foundation that everything else is built on top of.

Many faculties.

This is the deep economy the field keeps missing while it builds a separate silo for every skill: general intelligence is not a thousand capabilities in a trench coat.

(A silo is a walled-off store that does not share with its neighbours; the trench-coat image is the joke of a thousand narrow tricks pretending, stacked up, to be one real general mind.)

It is a handful of operations over ONE well-designed representation of how the world works.

A representation is a way of encoding knowledge so a machine can store and use it.

The representation was the hard part.

It exists now.

CHAPTER 12: THE GLASS BOX MIND



The Roadmap to AGI and ASI

Now return to the second wall - the black box crisis - and watch it come down.

(Recall that a black box is a system whose inner workings you cannot see; the crisis is that today's most powerful AI is exactly such a box.)

A mind built on Causalontology is a GLASS BOX by construction, and the phrase means something exact.

A glass box is the opposite of a black box: you can see straight through it to every part of how it works.

Its beliefs are enumerable: here are my causal claims, my confidence in each, my evidence for each, graded and signed.

Enumerable means they can be listed out, one by one, completely.

Its reasoning is traceable: this plan follows this chain of Causal Relation Objects; this expectation came from that mechanism; this surprise came from that violated window (traceable means every step can be followed back to where it came from).

Its changes are surgical: when the mind learns better, a NEW object appears - a refinement, a fresh assertion, an honest retraction - and the diff between yesterday's mind and today's is a readable list of signed records, not a hundred-billion-parameter shrug.

A retraction is a formal taking-back of an earlier claim now believed wrong.



The Roadmap to AGI and ASI

A diff is a list of exactly what changed between two versions.

A parameter is one of the internal adjustable numbers inside an AI model, and today's largest models have hundreds of billions of them, which is why their changes are an unreadable shrug.

Now say the word ALIGNMENT into this architecture and hear how differently it rings.

Alignment is the project of making an AI system's goals and behaviour actually match what humans intend and value.

Auditing a belief system becomes accounting: every claim of consequence can be traced to its authors and its evidence grades before the system acts on it (to audit is to inspect and check something thoroughly and officially).

Correcting a dangerous belief is a targeted operation on an identified object - not a fine-tuning campaign that hopefully nudges the right ghost in the weights.

(Fine-tuning is the process of nudging an already-trained AI model by training it further; the "ghost in the weights" is the joke that in a black box you are poking at an invisible spirit among those billions of numbers, hoping you touched the right one.)

And the machine's answer to "why?" is a LOOKUP, not a performance: the mechanism field, the evidence trail, the lineage - the explanation IS the representation, which was the brutal rule the black-box world could not satisfy.



The Roadmap to AGI and ASI

A lookup is simply reading an answer that is already stored, rather than improvising one.

We do not claim this solves alignment; honesty is the house style, and alignment has depths beyond representation.

We claim something narrower and still enormous:
ALIGNMENT BECOMES AN ENGINEERING DISCIPLINE
INSTEAD OF AN EXCAVATION – because for the first time, the thing to be aligned can be read.

(An excavation is a dig, groping blindly through buried material; the point is that today aligning a black box feels like archaeology, and here it becomes ordinary readable engineering.)

The difference between those futures is the difference between auditing and hoping.

Civilizations should not run on hope.

CHAPTER 13: THE COMMONS: A WORLD MODEL FOR EVERYONE

A commons is a shared resource that everyone may use and contribute to, like a public park or a public library.

A world model is a system's internal picture of how the world works.

Every skull that has ever learned has learned ALONE.



The Roadmap to AGI and ASI

The child masters her switch; the knowledge dies with the afternoon unless she happens to tell someone.

Laboratories rediscover each other's failures; models retrain from scratch; humanity's causal knowledge lives scattered in prose, in paywalled figures, in the retiring engineer's head.

Causalontology's second gift - inseparable from the first - is a SHARED WORLD MODEL: one global, provenance-first commons of causal knowledge that every mind, human or synthetic, can read from and write to.

Provenance means the recorded origin and history of a piece of information: who said it, when, and on what basis.

Provenance-first means that origin is built in from the start, not bolted on later.

Wikidata taught the world that a structured commons can work at planetary scale.

(Wikidata is a free, community-built database of facts, the structured cousin of Wikipedia; structured means the information is stored in tidy labelled fields a computer can query, not just written out as prose.)

Causalontology takes that lesson and specializes it to the knowledge that matters most and structures worst: causation - with the signatures, the evidence grades, and the no-contested-cells design that make a causal commons survivable where a wiki would burn.



The Roadmap to AGI and ASI

A signature here is a digital signature, a piece of unforgeable mathematics attached to a record that proves who wrote it and that it has not been altered.

A wiki is a website that many people can edit together; no-contested-cells means the design avoids the situation where two people fight over the one shared box that holds a fact, because each contributor gets their own separate record instead.

And because identity needs no committee (Chapter 8), the commons needs no center.

(Identity here means how each record gets its unique name; the earlier chapter showed that name can be worked out by mathematics from the record's own content, so no central authority has to hand out names.)

One store on a laptop serves a team; stores federate by simple pull, converging by set union; bundles cross air gaps on a USB stick and merge on arrival, every replica breaking even the pathological cycle cases in exactly the same deterministic way.

To federate means to join separate stores into a loose cooperating group while each stays independent; a pull is simply one store fetching the latest records from another.

Set union means combining two collections by keeping every item that is in either one, with no duplicates - the plain "put them all together" rule.



The Roadmap to AGI and ASI

An air gap is a deliberate physical separation, with no network between two computers, for safety; a USB stick carries data across that gap by hand.

A replica is one full copy of the store held in another place.

A cycle case is an awkward situation where records point round in a loop; deterministic means the rule always gives the exact same result no matter who runs it or in what order, so every copy ends up identical.

Tier by tier - single store, federation, full decentralization - the architecture has been built and tested, and Part Four shows the receipts (decentralization means having no central hub that everything depends on; the parts are spread out and equal).

Picture the asymptote honestly, the way one pictures the early web without apology.

An asymptote is the ideal limit a thing keeps approaching as it grows - here, the fullest form the shared commons could one day reach.

Every experiment humanity runs, every mechanism medicine confirms, every failure engineering pays for in blood and steel - landing once, signed, graded, mergeable, queryable by every mind that comes after (queryable means you can ask it questions and get answers back).

The species has never had a shared memory for how the world works.

It has one now, standing empty, doors open.



The Roadmap to AGI and ASI

CASE STUDY: WHAT WIKIDATA PROVED – AND WHAT IT LEFT UNDONE

(A case study is a close look at one real example to draw lessons from it.)

Skeptics of planetary knowledge commons must explain away an inconvenient existence proof.

(An existence proof is a real working example that settles the question of whether a thing is possible, simply by existing.)

Wikidata – the structured sibling of Wikipedia – holds over a hundred million items, maintained by tens of thousands of volunteers and bots, queried billions of times, feeding half the infoboxes and assistants on Earth.

A bot is an automated software helper that does edits on its own; an infobox is the little summary panel of facts you see at the top corner of a Wikipedia article.

It proved, beyond argument, four propositions this book depends on.

That volunteers WILL build structured knowledge at civilizational scale, given a schema and a mission.

A schema is the agreed blueprint that says what fields a record has and how they fit together – the shape every entry must follow.



The Roadmap to AGI and ASI

That HUMANS AND BOTS can garden the same field - most Wikidata edits are made by software, supervised by community norms.

That a commons can become LOAD-BEARING INFRASTRUCTURE without a business model attached to every read (load-bearing infrastructure means a foundation that other important things now rest on and cannot do without).

And that identity discipline - stable item identifiers - is the keel that keeps a giant afloat.

An identifier is the unique name or code that picks out one exact item; stable means it never changes, so links to it never break; the keel is the backbone beam that keeps a ship upright.

Honor it fully; we stand on its shoulders.

Then look, respectfully, at what the pioneer left undone - through no failure of will, but because its design predates the lessons.

Wikidata is NOUN-FIRST to its core: items are things - people, cities, proteins - and causation, where it appears at all, is squeezed into flat property links: "cause of death," "has effect" - arrows with no window, no modality, no mechanism, no context, no evidence grade.

Noun-first means the whole design is built around things (nouns) rather than around happenings and causes (verbs).

A property link is a simple labelled connection from one item to another.



The Roadmap to AGI and ASI

Modality is the "how certainly and in what way" of a cause – for instance whether it always produces its effect, sometimes does, or merely can.

The arrow is decoration again.

Its truth model is CONSENSUS-EDITED: statements can carry sources and ranks, but ultimately one community negotiates one shared cell per fact, and the edit history is a log of contention, not a structure of it – disagreement is a talk page, not a query.

A truth model is how a system decides what counts as true; consensus-edited means a group argues until it agrees on one shared answer.

A cell is one box holding one value; a talk page is the discussion page where Wikipedia editors argue in prose – useful to read, but not something a computer can neatly search.

Its provenance is ACCOUNT-BASED, not cryptographic: an edit traces to a username, not to arithmetic; nothing verifies offline; forks cannot re-merge by union.

Account-based means trust rests on who was logged in; cryptographic means trust rests on unforgeable mathematics that anyone can check for themselves.

A fork is a separate copy taken off in its own direction; the complaint is that two Wikidata forks cannot later be cleanly reunited by the simple put-them-all-together rule.



The Roadmap to AGI and ASI

And it has no ladder of evidence: a citation to a blog and a citation to a randomized trial are both, structurally, just references.

(A randomized trial is a carefully controlled experiment, among the strongest kinds of evidence there is; the point is that Wikidata's structure treats it no differently from a casual blog post.)

None of this diminishes the miracle.

It defines the SUCCESSOR'S JOB DESCRIPTION - and reading that description aloud, you will recognize every clause.

(A successor is the next thing that comes along to build on and surpass what came before.)

Verb-first, with the arrow reified and dressed in window, modality, mechanism, context.

Verb-first means built around happenings and causes rather than around things.

To reify something is to turn a mere idea or relationship into a full, first-class object the system stores and handles in its own right - here, promoting the causal arrow from decoration into a real record with all its details.

Signed provenance, checkable by anyone, mergeable by union, forkable without divorce.

Disagreement as structure; evidence as grades; acting above watching.



The Roadmap to AGI and ASI

And gaps as first-class invitations, so the commons steers its own volunteers - human and synthetic - instead of leaving them to wander.

Wikidata proved the world will build a shared brain for NOUNS.

Causalontology exists so the world can build one for WHY.

CHAPTER 14: COMMONS STORAGE AND FEDERATION DESIGN

We have called it a commons, and shown who fills it.

Now we must say plainly where it LIVES - because a shared memory for how the world works is worthless if it can be switched off by unplugging one machine.

Here is the founding fact of this chapter, and everything else falls out of it: THE COMMONS IS A NETWORK, NOT A SERVER.

A server is a single computer somewhere that everyone else must phone to get an answer; a network is a mesh of many computers talking to one another, with no one machine at the center.

No single machine should ever hold all of the commons, and - this is the quiet marvel - by design none has to.

This is not a new idea bolted on.

It builds directly on the three-tier store the standard already defines: Tier A, a single store on one machine; Tier B, a federation of stores



The Roadmap to AGI and ASI

that pull from one another; and Tier C, a fully decentralized layer with nobody in charge.

This chapter simply follows those tiers out to their natural horizon and shows the plumbing underneath.

WHY THE DATA WANTS TO BE A NETWORK.

Every identifier in the commons is the SHA-256 hash of its own content – recall from Chapter 8 that an object's name is a fingerprint computed from what it says, so identical content always earns the identical name.

And recall from that same chapter that two stores merge by set union: throw both piles together, count anything shared just once, and you are done – no coordinator, no vote, no meeting.

Two plain words before we go on.

Peer-to-peer means nodes talking directly to one another as equals, with no central authority any of them must obey.

A node means any one computer taking part in the network, holding some of the data and passing the rest along.

That combination – names computed from content, and merging by union – is the exact mathematics that already powers three famous systems.

The InterPlanetary File System (IPFS) is a peer-to-peer file network that addresses every file by its hash.



The Roadmap to AGI and ASI

Git is the version-control system, used by programmers everywhere, that names every saved snapshot of a project by a hash of its contents.

And Dat is a protocol for sharing versioned datasets from peer to peer.

All three lean on the same trick, and the commons shares it.

Because it shares that mathematics, the commons can be a peer-to-peer federation: many nodes, each holding some of the data, each gossiping newly minted objects to its neighbours, and all of them converging to the same whole.

Gossip means the simple habit of each node telling a few others whatever new objects it has just heard about, so that news spreads outward the way a rumour spreads through a town - with no announcer, no schedule, and no one in charge of the telling.

And because merging is union, the order in which the gossip arrives does not matter.

Every node, hearing the same objects in whatever jumbled sequence, converges to the same whole.

Convergence without coordination is the property this whole design exists to buy.

TIERS OF PARTICIPATION.

Not every node need be the same size, and this is where the hash earns its keep a second time.



The Roadmap to AGI and ASI

A full node holds everything - the entire type tier of the commons, every law humanity has signed into it - and answers any query from its own shelves.

A partial node holds only a slice.

And slicing is trivial precisely because identity is a hash: a hash is a string of characters spread evenly across all their possible values, so a rule as blunt as "keep every object whose hash begins with 0, 1, 2, or 3" hands one node a clean, even quarter of the commons, with no registry deciding who keeps what.

Splitting the load this way has a name - sharding - where a shard is one such slice of the whole.

A light client holds almost nothing at all.

A light client is a participant that stores no bulk data, asks other nodes for whatever it needs, and then verifies the answer for itself.

It needs to trust no one - and this is the part worth savouring - because every record is signed with Ed25519 and every content object self-certifies by its own hash.

A light client that asks a stranger for a claim can check the returned bytes against the hash it expected, and check the signature against the author's public key; if either check fails, the answer is thrown away.

A lying server is caught by mathematics, not by reputation.



The Roadmap to AGI and ASI

The light client does not need an honest network.

It needs an arithmetic one, and arithmetic it has.

WHAT BOUNDS THE SIZE.

The obvious fear about a shared world memory is that it grows without limit - every experience of every mind, forever, until no network on Earth could hold it.

The architecture forecloses that fear with a single rule about what gets shared: THE COMMONS FEDERATES LAWS, NOT DIARIES.

Recall the two tiers of content.

The type tier holds general causal law - "pressing causes lighting," everywhere and always - and this is what is shared world-wide.

The token tier holds a particular mind's history - the press it made at 9:04 on the fourteenth, what it saw, what it did - and this stays LOCAL BY DEFAULT, because a history is nobody else's business until its owner says so.

Privacy is the resting posture here, not an afterthought.

So the natural life of a mind on the commons runs like this: it observes its own tokens privately, it induces laws from them - the way the gardener of Chapter 16 turns a hundred private presses into one signed rule - and it contributes only the laws.

The diaries stay home; the laws go out.



The Roadmap to AGI and ASI

Weigh what that means for size.

The shared commons is not "every experience of every mind.

It is nearer to the scale of Wikidata - the structured commons of a hundred million items this book has already praised - than to some unthinkable ocean.

A world's worth of general causal law is a large but ordinary quantity: a normal-sized engineering problem, not an impossible one.

The frightening version of the commons was never the one the standard actually builds.

THE STORAGE ENGINE.

What does one node actually keep on its disk, and how does it keep it?

The heart of it is that content objects are immutable and content-addressed - written once, named by their own fingerprint, never altered.

That single property hands three gifts for free.

Deduplication: two copies of the same claim carry the same name and so are stored once, not twice.

Integrity: the name IS the checksum, so a corrupted object announces its own corruption the instant its bytes no longer match its name (a checksum being a small value computed from data precisely to detect whether the data has changed).



The Roadmap to AGI and ASI

And permanent caching: because an object can never change under its name, any copy of it anywhere stays valid forever, so it may be cached without any worry that it has gone stale.

Caching, said plainly, means keeping a copy of something near where it is needed so you need not fetch it again from far away.

In its bones, then, a node is two things sitting side by side.

First, a content-addressed blob store - a plain warehouse that files each object under its hash and hands it back on request, a blob being simply an opaque lump of stored bytes.

On a single machine this can be something as modest as SQLite, RocksDB, or LMDB - three well-worn, battle-tested storage libraries that live inside one program.

For holdings measured in the very large, it can be an object store such as MinIO or Garage - self-hosted systems that speak the widely used S3 protocol, the same request-and-store language the big cloud warehouses use - or PostgreSQL, the venerable open database, pressed into service as a shelf for blobs.

Second, alongside that immutable warehouse, the node keeps a locally rebuilt index for the mutable views - the ever-changing questions the commons must answer, such as where the gaps are, what each source's reputation is, and the provenance trail behind a claim.



The Roadmap to AGI and ASI

An index is a derived lookup table that makes searching fast, the way the index at the back of a book saves you from reading every page.

The division of labour is the same content-versus-provenance split the standard has insisted on from its first pages.

The content - the immutable objects - is what is SHARED across the network, identical on every node that holds it.

The indexes are LOCAL, rebuilt by each node from the content it holds, private to that machine, and never shipped, because any node can recompute them for itself.

Nothing that changes has to travel; only the things that never change do.

And which objects a node keeps hot - close and fast - versus cold - archived and slow - follows the very demand signal that Chapter 15's stigmergy already measures.

The claims the world keeps reading earn a place near the front; the claims nobody asks about drift to the back.

Hot and cold tiering is just the ordinary practice of keeping frequently used data on fast storage and rarely used data on cheap slow storage.

The commons already counts its own readership to rank its gaps; the same count tells each node what to keep within reach.

DISTRIBUTION AND DURABILITY.



The Roadmap to AGI and ASI

A network of live peers is good, but a network can go dark – peers sleep, links fail, an afternoon comes when nobody is online.

So the commons does not stake its survival on live peers.

It publishes SNAPSHOTS.

A snapshot is a single frozen bundle capturing the whole commons, or a chosen slice of it, as it stood at one moment – a bag of content-addressed objects, nothing more.

This is exactly how the largest public knowledge bases already keep themselves durable.

Wikidata ships periodic dumps of its entire contents for anyone to download.

The Internet Archive and Software Heritage – two institutions whose whole purpose is to preserve the world's digital record – keep vast holdings alive by copying them widely rather than trusting any one server.

Because a snapshot is just a bag of objects named by their hashes, it can travel by any road at all.

Anyone may pull it over IPFS, over BitTorrent – the well-known peer-to-peer protocol that lets a crowd share the labour of a large download, each participant sending pieces to the others – or over plain Hypertext Transfer Protocol Secure (HTTPS), the ordinary secure web transfer every browser already speaks.



The Roadmap to AGI and ASI

And however it arrives, the recipient can verify every single byte by hash before trusting it, then stand up a mirror.

A mirror is an independent copy of the whole, offered back to the network as another place to fetch from.

The consequence is worth stating slowly: the commons can reach global redundancy - many independent copies scattered across the world - before a single live peer is ever recruited.

A published, signed, verifiable snapshot is already a kind of immortality.

The data outlives the machine that first held it the moment the second copy exists.

READS AND WRITES AT SCALE.

Now weigh the two things a busy commons must do - answer reads and absorb writes - and see how gently each one scales.

Reads are the easy half, and content-addressing makes them almost embarrassingly easy.

Because an object can never change under its name, a copy of it stays valid forever, which means reads are infinitely cacheable.

Put mirrors near the readers, put a content-delivery network in front - a content-delivery network, or CDN, is a fleet of caching machines spread around the globe so that a reader is always served from somewhere nearby rather than from one



The Roadmap to AGI and ASI

distant origin - and a claim asked for a million times is fetched from the far origin perhaps once.

Popularity, which crushes ordinary servers, is here served almost for free.

Writes are the half that usually breaks a system, because ordinary databases must lock - freeze a record so two writers cannot clobber each other - and locks become bottlenecks under load.

The commons has no such lock.

A write is mint-and-gossip: a node computes the new object's hash, files it, and tells its neighbours.

There is no global authority to ask, because the name was never anyone's to grant.

Two nodes minting at the same instant - even minting the very same claim - cannot conflict, because union merge simply keeps what each produced and counts a shared object once.

And any drift that opens up while nodes are apart is closed later by anti-entropy reconciliation running quietly in the background.

Anti-entropy is the housekeeping habit by which two nodes periodically compare notes and each hands the other whatever objects it was missing, so that any gap between them steadily closes on its own.

Entropy is the tendency of separated copies to drift apart; anti-entropy is the patient work that pulls them back together.



The Roadmap to AGI and ASI

There is no global lock and no write bottleneck anywhere in the design – only minting, gossiping, and reconciliation, each of which any node can do without asking permission from any other.

TRUST AND SAFETY.

A network open to strangers must answer the hard questions squarely, and this one does, because its safety is structural rather than supervised – the same posture the safety chapter took earlier in the book.

Trust is a chosen policy, not a forced consensus.

Because every record is signed, each reader decides for itself whose signatures and whose grades of evidence it will believe, and the network never has to hammer all its participants into one agreed answer.

You are not made to trust; you are handed the means to check.

Privacy is protected by the token tier's local-by-default rule.

The histories that could expose a mind stay on that mind's own machine unless it deliberately publishes them, so federating the commons never means surrendering the diary.

History is corrected without being erased.

A retraction – the standard's one sanctioned goodbye – withdraws a claim from everyday view while leaving the record marked and intact for



The Roadmap to AGI and ASI

anyone who looks; the network can pass that withdrawal along, but no participant can quietly burn the transcript.

And tamper-evidence rides along on the mathematics for free.

Because every object is named by its hash and every mutable record is signed, even an untrusted mirror – a copy served by someone you have no reason to trust – cannot corrupt the data undetected.

Alter one byte and the hash no longer matches; forge a record and the signature fails.

You do not have to trust the mirror.

You only have to check the arithmetic, and the arithmetic never blinks.

THE PHASED BUILD.

None of this must arrive all at once, and the honest path is a staircase, each step useful on its own.

First, a persistent content-addressed node: a single machine that stores objects by their hashes and keeps them across restarts – Tier A, made durable.

That alone is a working commons for one team.

Second, a genesis full node that publishes signed snapshot dumps – the first authoritative copy, packaged and released so that the data outlives any one machine from its very first day.



The Roadmap to AGI and ASI

Redundancy before federation.

Third, the Tier B federation itself: many nodes gossiping new objects and reconciling by anti-entropy, so that there is no longer any single point of failure - unplug any one node and the others carry on unbothered.

Fourth and last, the reach outward: light clients that hold almost nothing and verify everything, CDN caching that puts reads near every reader, and hash-prefix sharding that lets partial nodes each carry a slice.

This is how the network scales SIDEWAYS - by adding more modest participants rather than one ever-larger server.

And the first step is no longer a promise on paper: it has been taken.

The reference node now persists.

Where it once held everything in memory and forgot it all the moment the process stopped, it now writes its objects and its signed records into a durable store and finds them waiting, whole, when it starts again - Tier A, made to remember.

The discipline of the change is the part worth savouring: nothing a user can see was altered.

Every address still answers as it did, every reply comes back in the same shape, down to the byte; the identity rule and the merge rule are untouched.



The Roadmap to AGI and ASI

A reader cannot tell the store was ever stopped and started, except by the one fact that matters - its data survived.

Under the surface the store now keeps the immutable content and the signed provenance each in its own place, and rebuilds its working views from those two alone, so that a single portable file carries the whole truth of a node and can be copied, checked, and trusted like any other object in the commons.

It is the smallest whole thing the staircase can deliver: one durable node, standing on a permanent floor.

And now the second step has been taken as well, and it is the one that lifts the commons off any single machine for good.

That durable node can now publish a SNAPSHOT of itself - a signed, downloadable dump of the whole commons - and a snapshot changes everything about how the data survives.

Picture it plainly.

The node writes out every object and every signed record, each as its own exact bytes, laid down in a fixed order so that the same store always produces the very same file; over the whole of it, it computes a single fingerprint - a Merkle root, one hash standing at the top of a tree of hashes, so that if any byte anywhere in the dump were to change, that top hash would change with it - and it signs a small manifest carrying that fingerprint with its own key.



The Roadmap to AGI and ASI

Alongside the dump it lays a plain checksum and a detached signature, the kind of receipts any ordinary tool can check.

That bundle is the commons, packaged to travel.

What you can do with such a bundle is the whole point.

You can fetch it however you like - a download, a shared drive, a stick carried across a room - and before you trust a single line of it you can verify it end to end: recompute the fingerprint and see that it matches the signed manifest, check that every object's name really is the hash of its own contents, check that every record's signature really holds.

Nothing asks for your faith; every claim in the dump certifies itself.

And once it checks out you can stand up a MIRROR from it - import the whole thing into an empty node and hold, provably, a copy identical to the original.

Import it twice and nothing changes the second time; try to slip in a tampered snapshot and the import refuses it whole, loading nothing rather than loading a lie.

There is one deliberate restraint in what a snapshot carries: by default it holds only the shared laws and never the private diary - the token tier, a mind's particular history, stays home unless its owner expressly chooses to send it - so what proliferates across the world is the



The Roadmap to AGI and ASI

commons of general knowledge, and not anyone's private life.

This is redundancy achieved before federation: copies can spread to a dozen machines in a dozen countries, each one self-checking, and the loss of any single node - or institution, or nation - can no longer take the commons down with it.

The one step still above this one - the light clients reaching sideways across the world - remains the road still ahead, named and waiting; but the data itself now outlives its keepers.

And now the third step has been taken too, and it is the one that turns a set of copies into a living thing.

The nodes no longer wait to be reconciled by hand; they reconcile themselves.

Give a node a short list of its peers, and two habits keep them all of one mind.

The first is gossip: the instant a node accepts a new claim - whether it minted the claim itself or heard it from a neighbour - it passes that claim onward to its peers, the way news travels a village.

The passing is thrifty and well-mannered: a flurry of writes goes out as one bundle rather than a storm of messages; a peer that already holds a claim simply shrugs and does nothing; and a node only ever forwards what was genuinely new to it, so an announcement ripples outward once and then dies, instead of echoing around a ring of neighbours forever.



The Roadmap to AGI and ASI

The second habit is the quiet housekeeping beneath the gossip - anti-entropy, the patient background reconciliation that catches whatever the fast path dropped: a message lost, a peer that was asleep, two writes that landed in different orders on different machines.

On a slow clock, each node compares notes with a peer by the cheapest possible question first - it holds up the single fingerprint of everything it knows, the same Merkle root the snapshot already taught it to compute, and asks, "does yours match?"

If the two fingerprints agree, the nodes are already identical and there is nothing to do; the whole exchange costs one small breath.

Only when the fingerprints differ do they look closer, trade the bare list of names each holds, work out exactly which handful each is missing, and hand over only that handful - never the whole library to settle a difference of three books.

And through all of it the old discipline holds, now enforced at every doorway in the network: nothing is believed that cannot prove itself.

Every object that arrives must still be the hash of its own contents, and every record must still carry a signature that holds, or it is turned away at the threshold and never let in - while the honest items travelling beside it are waved through, so a single forged claim cannot spoil the whole delivery and no broken or malicious peer can ever push a lie into another's store.

Federation only ever adds; it never erases.



The Roadmap to AGI and ASI

A retraction crosses the wire as the signed record it is, and hides a claim from view on the far side exactly as it does at home, without ever burning the history of what was once said.

There is no head of the network, no vote, no single ordained order of events; contradictory claims ride side by side, each wearing its own provenance, precisely as they do within one machine - and trust stays where it always lived, a private judgement each reader makes for itself.

The private diary stays home here as well: only the shared laws gossip and reconcile across the network; a mind's particular history never leaves its own keeping unless its owner sends it deliberately.

What makes all this safe without a referee is the same fact that has carried the whole design from the first page - the data is built so that merging is union, and union does not care about order or repetition, so any set of peers, talking in any sequence, dropping and repeating messages as real networks do, drifts inexorably toward the very same whole.

The commons is no longer a stack of copies waiting to be tended.

It is a living, self-healing network of peers, and only the reach outward now remains.

And now the fourth step has been taken too - the last one - and with it the whole staircase is built.



The Roadmap to AGI and ASI

The reach outward is no longer a promise; it is real.

A LIGHT CLIENT can now hold almost nothing at all and still read the commons safely: it asks whatever node is nearest for a claim, and then, trusting no one, checks the answer for itself - recomputing the object's name from the very bytes it received and seeing that the two agree, checking each record's signature against its author's key - so that a lying node, or a stale or hostile cache dressed up as a helpful one, is caught the instant its bytes fail to match their name.

Because the answer certifies itself, it no longer matters who hands it over.

And that single fact unlocks the cheap global read: since an object can never change under its name, any caching machine anywhere - any ordinary CDN edge a website already uses - may keep a copy forever and hand it out to everyone nearby, and the reader loses nothing by taking it from a stranger's cache rather than the far origin, because the reader verifies it regardless.

Only the settled, immutable things are cached this way; the living views that still shift as new records arrive are marked as perishable, never frozen and served as though they were final.

And the load itself is spread by SHARDING: because a name is a hash, the commons divides into even slices along the leading digits of those hashes, and a PARTIAL NODE can volunteer to hold just one slice - "I will keep everything whose name begins with these digits" - announcing its stretch so



The Roadmap to AGI and ASI

that anyone looking for a particular claim is pointed straight to a node that holds it.

Ask a node for something outside its slice and it does not lie and say "not found"; it hands back a signpost to a node that has it.

The gossip and the anti-entropy of the third step now run slice by slice, so that a partial node reconciles only its own stretch with the others who share it, and no machine is ever forced to carry the whole world to take part.

The network grows SIDEWAYS: every new reader, every edge cache, every volunteer holding a slice adds capacity instead of adding strain.

And the old restraint holds here too - the private diary is never served on these public roads, nor swept into any public slice, unless its owner chooses to send it.

Stand back and see what now stands finished.

The commons is durable, so it remembers; it is redundant, so it survives the loss of any keeper; it is self-healing, so it repairs itself without a hand to guide it; and now it is horizontally scalable, so it can be read by anyone from a client that trusts no one, cached at the edge of the world because its names are hashes, and hosted a slice at a time by whoever will help.

Four steps, each useful the day it landed, and together a whole: a durable, self-healing, worldwide network that is free and open not by anyone's charity but by the very way it is built.



The Roadmap to AGI and ASI

The architecture is finished.

What remains is not to design it further but to FILL it - to stand up the first genuine nodes, to invite the libraries and the universities and the strangers, and to let the commons grow.

THE CLOSING NOTE.

Stand back and notice what kind of thing has been built.

Because it is a federation and not a server, its cost is spread across all who take part, and no central owner is left holding a single crushing bill; institutional full nodes - libraries, universities, archives - can anchor it the way great libraries anchor the written record, without any one of them owning it.

The open license and the open stack - every piece of it, from the storage libraries to the transfer protocols, freely available and freely inspectable - make the commons cheap to mirror and impossible to quietly corrupt.

And here is the through-line of the whole book, arriving one last time from a new direction.

The very same content-addressing that let two strangers in Nairobi and Oslo mint the same name for the same claim - the mathematics that makes the commons MERGEABLE - is what also makes it FREE and OPEN.

Free, because a name computed from content is nobody's to sell.



The Roadmap to AGI and ASI

Open, because a record that certifies itself needs no gatekeeper to be trusted.

The commons is not free and open by anyone's charity, to be withdrawn on a bad day.

It is free and open by CONSTRUCTION – built that way in its bones, the way a bridge is built to stand.

CHAPTER 15: STIGMERGY: THE WALL THAT TEACHES

An empty commons raises the coldest question in all of infrastructure (the shared plumbing that many people build on top of): WHO WILL FILL IT, AND WHY?

Causalontology's answer is stigmergy, and it may be the most quietly beautiful idea in the whole design.

Stigmergy is a long word for a simple thing: coordination through the shared environment itself, where the work already done leaves a trace that tells the next worker what to do, with nobody giving orders.

Termites raise cathedrals with no blueprint and no foreman: each half-built wall SHOWS the next worker where the next brick goes.

The work organizes the workers.

In Causalontology, the structure's own gaps – the holes and missing pieces in the shared knowledge – are first-class, queryable facts.



The Roadmap to AGI and ASI

By first-class we mean the system treats a gap as a real, named thing you can ask about, not as an afterthought.

By queryable we mean you can literally ask the store to hand you its gaps, the way you would ask a librarian for a list of missing books.

There are six kinds of them.

- A MISSING FIELD: a claim with no temporal window (no statement of the time span over which it holds) or no modality (no statement of how strongly the cause forces the effect - for example, whether it always happens or only sometimes) - refine me.
- An EMPTY MECHANISM: a relation not yet decomposed, meaning we know that one thing leads to another but nobody has spelled out the smaller steps in between - explain me.
- A DANGLING REFERENCE: a thing named but never defined - the red link that says this page is wanted, just like an encyclopedia entry that other pages point to but that nobody has written yet.
- An INCONSISTENT HIERARCHY: a mechanism that fails its arithmetic, meaning the numbers describing its parts do not add up the way they must - repair me.
- A CONFLICT: two claims that formally contradict, meaning they cannot both be true at once - bring evidence.
- And DEMAND WITHOUT SUPPLY: a claim the world keeps READING but no one has supported with



The Roadmap to AGI and ASI

evidence - measured from real readership, the store counting its own visitors the way a shop counts the people who walk through the door - invest here.

Ask the store for gaps and it answers VALUE-RANKED, meaning it sorts them by how much they matter: contradictions first, hungry claims weighted by live demand, the whole frontier sorted so that the first item is always, literally, the most valuable brick to lay next.

The frontier is simply the leading edge of the unfinished work, the row of open gaps waiting to be closed.

No assignments.

No committee meetings.

No roadmap document rotting in a drawer.

The commons ANNOUNCES ITS OWN NEEDS, continuously, to every contributor on Earth - and when someone closes a gap with a refinement (a sharper, more complete version of a claim) or a signed enrichment (an addition stamped with the author's cryptographic signature, a kind of tamper-proof digital seal that proves who made it), it vanishes from everyone's frontier at once, and the wall shows the next worker a different brick.

There is even a dashboard - one self-contained page served by the store itself - where humans watch the frontier move in real time.



The Roadmap to AGI and ASI

Machines read the same frontier through the same interface (the same shared doorway into the system).

Which brings us to the machine that already did.

CHAPTER 16: THE FIRST GARDENER

Every claim in this chapter is reproducible from the public repository (the openly available online home of the code), and we will say each one precisely, because precision is the house religion.

First, the pedigree of the reference implementation – that is, the family history of the official, standard-setting version of the software.

Causalontology did not fall from theory.

It was distilled from PrologAI, a glass-box cognitive architecture that solved – by pure symbolic induction, with named, readable rules, and NO Large Language Model (LLM) anywhere in the loop – all four hundred of four hundred tasks on the public training set of the first Abstraction and Reasoning Corpus for Artificial General Intelligence benchmark (ARC-AGI-1), and all one hundred twenty of one hundred twenty on the public set of the second (ARC-AGI-2).

A glass-box system is one whose every step you can see and read, the opposite of a black box whose inner workings are hidden.



The Roadmap to AGI and ASI

Symbolic induction means learning explicit, spelled-out rules from examples, rather than tuning millions of hidden numerical weights.

A Large Language Model (LLM) is the kind of artificial intelligence, trained to guess the next word over enormous piles of text, that powers today's popular chatbots.

The Abstraction and Reasoning Corpus for Artificial General Intelligence (ARC-AGI) is a famous set of visual puzzles built to test genuine reasoning rather than memorization.

Say the qualifiers with us, every time: PUBLIC TRAINING SETS, stated precisely, never inflated.

A training set is the batch of practice problems that come with the answers, as opposed to the sealed exam used for official scoring.

The benchmarks' private evaluations are a different mountain, and the third benchmark is a labeled bet, not a claim – meaning we have honestly marked it as a wager on what we expect, with no score claimed until measured.

But hold the residue that survives all qualifiers: the causal, structural, glass-box approach REASONS – solves genuinely novel puzzles a correlation engine cannot memorize its way through – and every solving rule can be read aloud afterward.

A correlation engine is a system that only notices which things tend to appear together, without understanding why – it can copy patterns it has seen, but it cannot work out something truly new.



The Roadmap to AGI and ASI

Those are the bones Causalontology was carved from.

Now the demonstration this book has been building toward – the first synthetic mind gardening the commons.

It runs, live, against a real store, in the repository's continuous verification (an automatic re-checking of the whole system that fires on every change), on every change.

A glass-box agent in the PrologAI style – every line of its code carrying its English explanation – wakes and reads the frontier: GET the gaps, value-ranked.

GET here is just the plain command for fetch, the same word a web browser uses to ask a server for a page.

It chooses the most valuable brick: a degenerate claim – pressing causes lighting, nothing more, with none of the missing detail filled in – that the world has been reading but no one has refined.

And then it does the thing this whole book exists to celebrate.

IT PUTS ITS OWN HAND ON THE SWITCH.

One hundred interventions in its microworld – an intervention being a deliberate action taken to see what happens, and a microworld being a small, controlled toy environment.



The Roadmap to AGI and ASI

Ninety-eight lightings; two failures on a worn contact; slowest response, nine tenths of a second.

From its own actions - not from text ABOUT switches - it induces the missing shape: modality SUFFICIENT (meaning the press is enough, on its own, to cause the light), at rate 0.98; temporal window, zero to one second (the effect follows within one second).

To induce here means to work out the general rule from the particular results it saw.

It mints the refinement - a new Causal Relation Object, the standard little package that records one cause-and-effect fact, completing the old one, lineage intact.

Lineage means the chain of descent is preserved, so anyone can trace the new version back to the old one it grew from.

It signs an assertion with its own Ed25519 key: evidence type INTERVENTION, the strongest grade, strength 0.98.

Ed25519 is a specific, widely trusted digital-signature method, and the key is the private secret that lets it stamp its work with a seal only it could have made.

An assertion is simply a stated claim being put on the record.

And it checks the frontier one last time, and narrates, in its own log, the sentence we will quote for years:



The Roadmap to AGI and ASI

"the gap I chose is gone from the frontier - the wall shows the next worker a different brick."

Every step of that loop - read, choose, act, induce, mint, sign, verify - is inspectable down to the byte, and the test that proves it ends by computing the gardener's REPUTATION from its signed public history alone.

Reputation here is a trust score calculated purely from the agent's own track record of signed, checkable contributions.

The microworld is a stand-in; the loop is the real thing.

Point the same loop at richer environments and richer stores, and you have the future this architecture was built for: minds that do not just consume humanity's knowledge but CONTRIBUTE to it - signed, graded, auditable - bricklayers in the shared wall, accountable by arithmetic.

The child at the light switch finally has a colleague.

PART FOUR: THE PROOF

Why you can believe any of this - and check all of it.

CHAPTER 17: THE FREEZE

A standard you cannot test is an opinion with margins.



The Roadmap to AGI and ASI

So Causalontology's constitution is a TEST SUITE - a fixed collection of automatic checks - built from one hundred and thirty-seven conformance vectors: published, numbered, frozen.

(They run from V01 through V137, and test everything the standard defines: the schema purity of all twenty-one kinds, the semantic rules, the strata and bridges and seams, the ports and conduits, the token tier, the tick unit, the realized_by binding, the skip semantics, the attitude and its quarantine, and the prediction and its signed error.)

A conformance vector is a single, precisely written test case - one exact input paired with the exact behavior a correct implementation must produce - so anyone can confirm their software conforms to (obeys) the standard.

The rule has no exceptions and no diplomacy: AN IMPLEMENTATION IS CAUSALONTOLOGY-CONFORMANT IF AND ONLY IF IT PASSES ALL ONE HUNDRED AND THIRTY-SEVEN.

The vectors cover the whole nervous system: schema purity and the twenty-one kinds (schema meaning the required shape or template every piece of data must follow); the semantic rules (the rules about what things mean); strata, bridges, ports and conduits, and the skip semantics of layer-crossing causation; the token tier of what actually happened; identity invariance under formatting (the same fact keeping the same identity even when it is written out differently); signature verification and tamper rejection (confirming a seal is genuine and refusing anything that has been altered); retraction, lineage, and succession



The Roadmap to AGI and ASI

(taking a claim back, tracing its ancestry, and one version rightfully replacing another); corroboration (independent evidence backing a claim up); deterministic cycle-breaking (settling a knowledge loop the same predictable way every time); the resolve ranking (the fixed order in which competing claims are sorted out); the doxastic quarantine of the attitude and the pairing discipline of the prediction error (the token tier's laws for belief and expectation, that a belief is never read as a fact and an error is never graded against the wrong event); and the stigmergic gap closing when a refinement lands.

The suite was FROZEN at specification 4.0.0: every symbolic placeholder replaced with concrete bytes – real sixty-four-character identifiers, real public keys, and one real signature that genuinely verifies – by a freeze tool that is deterministic and idempotent.

Deterministic means it gives the same result every time, with no randomness.

Idempotent is the mathematician's way of saying: run it twice and the second run changes nothing, the way pressing an elevator button a second time does not summon a second elevator.

Beneath the vectors, thirty design decisions stand LOCKED – written down, numbered, changeable only by a public, versioned process (an open procedure where every change is dated and tracked).

Builders can stand on this ground.

It does not shift.



CHAPTER 18: TWENTY LANGUAGES, ZERO SHARED CODE

Here is the empirical heart of the whole enterprise – the experiment that turns philosophy into physics.

Empirical means resting on real, observable evidence rather than on argument.

If Causalontology is truly a PROGRAMMING-LANGUAGE-NEUTRAL data structure – knowledge itself, not software, where programming-language-neutral means it does not depend on any one programming language – then independent implementations in wildly different programming languages, sharing not one line of code, should agree perfectly.

An implementation is one working version of the standard, written out in an actual programming language.

Not roughly.

Byte for byte, meaning identical right down to the smallest unit of computer data.

The experiment has been run twenty times.

Prolog.

Python.

JavaScript.

Rust.

Java.



The Roadmap to AGI and ASI

Swift.

Go.

C#.

Ruby.

PHP.

Dart.

Elixir.

Haskell.

Perl.

R.

Lua.

C++.

Zig.

Julia.

Kotlin/Native.

Twenty implementations, twenty languages, every one passing the same one hundred and thirty-seven frozen vectors – agreeing down to the length of a month, the ranking of a name lookup, the identical deterministic loser of a knowledge cycle, and the exact sixty-four bytes of an Ed25519 signature.



The Roadmap to AGI and ASI

Several of these ports – a port being the same design rebuilt in a new language – are small epics of verification in their own right.

In languages whose standard libraries carry no cryptography at all (cryptography being the mathematics of secret codes and tamper-proof seals), the signature mathematics was built from first principles – in Lua, even the big-number arithmetic beneath the curve was hand-built and cross-checked against the reference on hundreds of independent cases; in Haskell, the hash constants were regenerated from the square and cube roots of primes and checked against known answers before a single vector ran.

A hash is a short fingerprint that a fixed calculation produces from a chunk of data, so that even a tiny change to the data produces a completely different fingerprint.

Different languages, different authors, different decades of design philosophy – one identical behavior, because the MEANING lives in the standard, not in any implementation.

To our knowledge, no causal knowledge system in history has demonstrated programming-language-neutrality at this breadth.

Programming-language-neutrality is no longer a claim.

It is a table, and every row of it is green.

CHAPTER 19: INSTALL IT BEFORE BREAKFAST

Theories ask for belief.



The Roadmap to AGI and ASI

Infrastructure asks for a terminal – the plain text window where you type commands to your computer.

As this book goes to print, Causalontology 4.0.0 is live on the world's package registries – the official online stores from which programmers download ready-made software – one line each:

- Python: `pip install causalontology`
- JavaScript or TypeScript: `npm install causalontology`
- Rust: `cargo add causalontology`
- Java: `io.github.ai-university-aiu:causalontology:4.0.0` on Maven Central
- Kotlin/Native: `io.github.ai-university-aiu:causalontology-kotlin:4.0.0`
- C# / .NET: `dotnet add package causalontology`
- Ruby: `gem install causalontology`
- PHP: `composer require causalontology/causalontology`
- Dart: `dart pub add causalontology`
- Elixir: the dependency `{:causalontology, "~> 4.0"}` on Hex
- Lua: `luarocks install causalontology`



The Roadmap to AGI and ASI

- Swift and Go and Zig: one dependency line each, straight from the repository tags.

TypeScript, named above, is simply a stricter cousin of JavaScript that checks your work as you write.

Additional registries are in community review as we print, and a WebAssembly (WASM) build of the audited core – one binary that runs in browsers, edge servers, and any modern host – ships with the public release.

WebAssembly (WASM) is a portable format that runs the same program on many kinds of computer, so a single packaged file behaves identically wherever you drop it.

And guarding all of it: TWENTY-ONE continuous-integration gates, re-proving every implementation against the frozen vectors on every single change, in public, forever – with the synthetic gardener tending the commons at the end of every run.

Continuous integration (CI) is the practice of automatically rebuilding and retesting the whole system the instant anyone changes it, and a gate is a check that must pass before the change is allowed through.

We built the burden of proof INTO the infrastructure.

Do not take our word for one sentence of this book.

Take the repository: github.com/ai-university-aiu/causalontology.



The Roadmap to AGI and ASI

Clone it - that is, copy the whole project onto your own machine.

Run the gates.

Watch the gardener lay its brick.

Glass box, always - the marketing IS the audit.

THE CROSS-EXAMINATION

Eight objections, stated at full strength, answered without dodging.

A manifesto (a public statement of beliefs and intentions) that cannot survive its skeptics is a greeting card.

Put us on the stand.

OBJECTION ONE: "WE HAVE SEEN THIS MOVIE - SYMBOLIC AI ALREADY FAILED."

The strong form: for forty years, Good Old-Fashioned Artificial Intelligence (GOFAI, the early approach to Artificial Intelligence that tried to build thinking out of hand-written symbols and logic rules rather than learning from data) promised that hand-written symbols would add up to minds.

Symbolic Artificial Intelligence, to say it plainly, is the idea that intelligence can be spelled out as tidy labels and rules a programmer types in by hand.



The Roadmap to AGI and ASI

The flagship effort of that era spent decades and hundreds of person-years hand-coding common sense into a private ontology (an ontology is an organized catalog of what kinds of things exist in some domain and how they relate - here, one company's private catalog of the world), and general intelligence never came.

The field moved on for good reasons - and here you are, selling symbols again.

Answer: we agree about the corpse, and we did the autopsy - the old symbolic programs died of four specific wounds, and this architecture was designed wound by wound against the report.

They died of HAND-CODING as the primary source of knowledge.

Here, the primary source is INTERVENTION - an intervention means actually doing something to the world to see what happens, rather than just watching it, so here agents (acting programs) and instruments act on the world and write down what acting taught them, with the evidence grade (a stamped rating of how strong the proof is, from mere rumor up to a controlled experiment) to prove it.

The gardener of Chapter 15 hand-coded nothing.

They died of BRITTLINESS - one ontology, one viewpoint, contradictions fatal (a single rigid catalog that shattered the moment two facts disagreed).

Here, contradiction is a first-class citizen with its own object kind and its own place on the



The Roadmap to AGI and ASI

frontier, and rival claims coexist signed (each stamped with a digital signature, a piece of unforgeable math that proves who wrote it, like a wax seal only the true author can make) until evidence sorts them.

They died of CLOSED development - priesthods curating private truth (small closed groups guarding a secret master copy of what was true).

Here, the format is open, the commons (the shared public pool everyone can read and add to) is public, and identity needs no committee.

And they died of UNFALSIFIABILITY - no public test of whether an implementation was even correct (unfalsifiable means there was no experiment that could ever prove the thing wrong, so nobody could check it).

Here, correctness is one hundred thirty-seven frozen vectors, pass or fail, in twenty languages (one hundred thirty-seven fixed, unchangeable test cases that any version must reproduce exactly, checked across twenty programming languages).

We are not reviving the old symbolic program.

We are giving the STATISTICAL age (today's era of Artificial Intelligence built from statistics and pattern-counting over huge piles of data) the one symbolic organ it actually lacks - a signed, mergeable ledger for what acting teaches.

A ledger here means a permanent, append-only record book, like an accountant's, that many parties can merge together without conflict.



The Roadmap to AGI and ASI

Which is roughly the difference between rebuilding the horse and inventing the fuel standard.

OBJECTION TWO: "THE BITTER LESSON - SCALE BEATS STRUCTURE, ALWAYS."

The strong form: the most-quoted essay in modern machine learning teaches that general methods riding computation always defeat human-designed structure.

The Bitter Lesson is the famous claim that, over the long run, simply throwing more computer power and more data at a problem beats any clever hand-built structure a human designs in.

Every structure-first bet has lost for seventy years; yours will too.

Answer: read the lesson's fine print, because we are not on the losing side of it - we are underneath both sides.

The Bitter Lesson is about which METHODS win - learned versus hand-designed - and we take no position against learning.

Nothing in Causalontology competes with gradient descent (gradient descent is the standard training method that nudges a learning system's internal numbers a tiny step at a time toward fewer mistakes, like feeling your way downhill in fog toward the lowest point), and the biggest learned world model on Earth is welcome to write its discoveries into Causal Relation Objects (CROs, the single record-type at the heart of this system, each one capturing one piece of cause-and-



The Roadmap to AGI and ASI

effect: what made what happen, by whom, on what evidence) tonight.

A world model, by the way, is a system's internal picture of how its surroundings behave and what its actions will do next.

But notice what the lesson's own winners all stand on: shared FORMATS that nobody learned and everybody agreed to - the bytes, the floating-point standard (the agreed way computers store decimal numbers so every machine reads them the same), the image formats, the very datasets - structure of exactly the boring, load-bearing kind.

Formats are not the structure the lesson warns against; formats are the table the game is played on.

And there is a second fine print: the lesson's victories are victories of rung one - association at scale.

Rung one, association, means noticing that two things tend to occur together, the lowest level of causal understanding - correlation without knowing which one causes the other.

And the very laboratories that quote it are now, in 2026, raising billion-dollar seeds to find what scale did not deliver.

When the scaling curve itself starts shopping for structure, quoting the Bitter Lesson at a data standard is not hard-headedness.

It is nostalgia.



The Roadmap to AGI and ASI

OBJECTION THREE: "SYMBOLS WITHOUT GROUNDING ARE EMPTY - WHO CONNECTS `occurrent:press_button` TO REALITY?"

The strong form: the symbol grounding problem is a century old; your `occurrents` are just strings; meaning cannot live in a lookup table.

The symbol grounding problem is the old puzzle of how a bare label inside a machine ever comes to mean the real thing out in the world, rather than just pointing at other labels forever.

An `occurrent` (here written `occurrent:press_button`) is this system's word for a happening or event - something that unfolds in time, like a button being pressed, as opposed to a standing thing.

The objection says such a label is just text, a string of characters, and a dictionary of strings pointing at strings is empty.

Answer: grounding is exactly why this architecture is INTERVENTIONAL to its bones, and the objection has the direction of flow backwards.

In Causalontology, a symbol is not grounded by definition but by TRACK RECORD.

`occurrent:press_button` means something because agents have pressed buttons and signed what followed - the symbol accumulates intervention-graded assertions the way a word accumulates usage, and the grounding IS the provenance trail.

Provenance means the documented origin and history of a claim - who made it, when, and on what basis



The Roadmap to AGI and ASI

- the way a museum records exactly where a painting came from and every hand it passed through.

The gardener grounds its vocabulary the same way the child does: hand, switch, tally, signature.

Moreover, the architecture never claims symbols are the whole mind - it is the noun-hosting HALF of a partnership in which perception and motor control may be as neural as you please.

Neural here refers to a neural network, the brain-inspired kind of Artificial Intelligence that learns patterns from examples by tuning millions of internal numbers, the technology behind today's image and language systems.

The standard only insists that what the acting TAUGHT be written somewhere readable.

Grounding by acting, recorded with receipts.

Hume would have signed it.

OBJECTION FOUR: "THE COLD START - A COMMONS WITH NO CONTRIBUTORS IS A PARKING LOT."

The strong form: network-effect systems die in the empty quadrant.

A network effect means a thing gets more valuable the more people already use it - a telephone is useless if you are the only owner - and the cold start is the chicken-and-egg trap where nobody joins because nobody has joined yet.



The Roadmap to AGI and ASI

Wikipedia was a miracle you will not repeat; why would anyone write the first ten thousand claims?

Answer: four honest reasons, one of them new under the sun.

First, the SELFISH loop comes before the social one: a single team running a private store gets value on day one - their own agents' knowledge, signed and queryable (able to be searched and asked questions like a database).

Federation is a flag they can flip later - federation means separate private stores choosing to link up and share, while each still keeps control of its own, the way independent banks can join one payment network without merging into one company.

The commons grows the way version control did, private repositories first, the great public hub after.

Version control is the everyday tool programmers use to track every change to their files and merge many people's edits safely, and a repository is one project's tracked collection of files and its full history.

Second, the frontier PAYS ATTENTION where wikis paid only prestige: gaps are value-ranked by real demand, so a contributor always knows which brick matters most - no blank-page paralysis, ever.

Third, the reserved bounty layer is designed to make closing gaps a LIVING - a bounty here is a standing cash reward attached to an unanswered question, so demand becomes prizes the moment



The Roadmap to AGI and ASI

governance opens that door; the plumbing for it (signed contribution, measurable demand) is already laid.

And fourth - the new thing - THE MACHINES ARE CONTRIBUTORS: one synthetic mind (a mind built in software rather than born in biology) is already laying bricks, tirelessly, gladly, with receipts.

The cold-start arithmetic of every previous commons assumed only human hands, and that assumption just expired.

OBJECTION FIVE: "GARBAGE IN - THE COMMONS WILL FILL WITH CONFIDENT NONSENSE."

The strong form: open systems drown in spam, vandalism, and earnest error; a causal commons will be a landfill of astrology with signatures.

Answer: the design assumes bad input as a WEATHER CONDITION, not a surprise, and the defenses are structural, not moderational (built into the shape of the system itself, not depending on human moderators patrolling it).

Nothing unsigned enters default views - quarantine is automatic, not a judgment call.

Quarantine means unverified material is fenced off in a holding pen where it cannot mislead anyone, exactly like keeping a new arrival isolated until it is checked, and here it happens by rule, with no human deciding case by case.

Every claim wears its author and its evidence grade, so astrology arrives marked as somebody's imported observation and can be filtered by any



The Roadmap to AGI and ASI

consumer's trust policy in one line - the commons never launders provenance into authority.

Corroboration is countable: one crank asserting nonsense is one signature; forty independent laboratories asserting a mechanism is forty, and the difference is a query, not an argument.

Corroboration means independent confirmation - separate parties backing the same claim - and here you can literally count the backers instead of arguing about them.

Contradiction does not corrupt - a false claim beside a true one creates a surfaced conflict, which is an INVITATION for exactly the intervention-grade evidence that settles it.

And the worst case is bounded by immutability: vandals cannot deface anyone else's records - they can only add their own, signed, forever attributable.

Immutable means once written, a record can never be secretly altered or erased, only added to, like ink already dried in a bound ledger.

Which is the least attractive vandalism ever offered.

Landfills happen where garbage is anonymous.

Here, garbage arrives gift-wrapped with the vandal's cryptographic business card (an unforgeable digital signature that names the culprit permanently).



The Roadmap to AGI and ASI

OBJECTION SIX: "WHO DECIDES WHAT IS TRUE - YOUR STANDARD JUST OUTSOURCES THE HARD PART."

The strong form: you have punted on truth; a pile of signed disagreements is not knowledge; someone must arbitrate, and you have carefully avoided saying who.

Answer: guilty, with pride, because the punt IS the design and history is on its side.

Infrastructure that decides truth becomes a throne, and thrones get fought over.

Infrastructure that decides EVIDENCE - what was claimed, by whom, on what grounds, contradicted by what - becomes a court record, and court records are what civilizations actually run on.

Causalontology fixes the layer that can be fixed: provenance arithmetic, evidence grading, structured disagreement, honest retraction.

Retraction means openly taking back a claim you no longer stand behind, without deleting the fact that you once made it - a correction on the record rather than a quiet erasure.

Truth-arbitration then lives where it has always lived and always must - in the trust policies of readers, the review practices of communities, the replication habits of sciences - except now those practices operate over signed, graded, queryable material instead of over rumor and PDF (Portable Document Format, the frozen page-image files that scientific papers usually hide inside).

We did not outsource the hard part.



The Roadmap to AGI and ASI

We de-monopolized it, and handed every arbiter better instruments.

OBJECTION SEVEN: "EMBEDDINGS ALREADY DO THIS - KNOWLEDGE LIVES BETTER IN VECTORS."

The strong form: modern systems represent knowledge as high-dimensional vectors that capture similarity beautifully; discrete objects are a step backwards from the geometry that actually works.

An embedding is the way today's Artificial Intelligence turns a word, image, or idea into a long list of numbers - a vector - positioned in a vast mathematical space so that things with similar meaning land near each other.

High-dimensional just means that list has hundreds or thousands of numbers rather than the two or three we can picture.

Discrete objects, by contrast, are separate, countable, individually named records - like index cards - rather than points blurred together in a cloud of numbers.

Answer: vectors and objects are not rivals; they are different organs - and ask the vector three questions to feel the difference.

Ask WHO SAYS: a vector has no author; it is the anonymous residue of a training corpus (the leftover average of everything a model was trained on), and no signature can ever be checked against it.



The Roadmap to AGI and ASI

Ask WHY: a vector cannot show its mechanism; nearest neighbors are not reasons, and no regulator, court, or bereaved family will ever accept cosine similarity as an explanation.

Cosine similarity is the standard yardstick for how close two vectors point in that number-space, a single closeness score - useful for "these two feel alike," useless for "here is why one caused the other."

Ask WHAT CHANGED: update a model and every vector moves at once; there is no diff, no retraction, no lineage.

A diff is a precise before-and-after comparison showing exactly what changed, and lineage is the traceable chain of where something came from and how it was revised.

Institutional memory is impossible in a medium that forgets by remodeling.

Embeddings are a magnificent SIMILARITY organ, and minds built on this architecture should use them for retrieval, for perception, for hunches.

But the system of record (the one authoritative, trusted copy that settles disputes) for what makes what happen needs authors, reasons, and diffs - and that is a job description only discrete, signed, immutable objects can fill.

Geometry for guessing; ledgers for knowing.

A general mind carries both.



The Roadmap to AGI and ASI

OBJECTION EIGHT: "STRIP THE POETRY AND THIS IS JUST A DATABASE WITH EXTRA STEPS."

The strong form: schemas, hashes, signatures - fine engineering, nothing new; you have written a long love letter to a CRUD app.

A schema is the fixed blueprint that says what fields a record must have and what shape they take - the ruled columns of a form.

A hash is a fingerprint for data: a short string mechanically computed from a file such that the same input always yields the same fingerprint and any tampering changes it.

A CRUD app is the most ordinary kind of software there is - it only Creates, Reads, Updates, and Deletes records (that is what the letters stand for), a plain data-in, data-out program.

Answer: every civilization-scale infrastructure is "just" something with extra steps, and the steps are the civilization.

Double-entry bookkeeping is just a table with extra steps - and it created auditable commerce.

Double-entry bookkeeping is the centuries-old accounting rule of recording every transaction twice, as a matching credit and debit, so the books police themselves.

Version control is just file copies with extra steps - and it created software as a cumulative, global craft.

The web is just files with extra steps.



The Roadmap to AGI and ASI

The specific extra steps here - reified causation as the atom; identity as content; provenance separated uniformly; evidence graded with acting on top; disagreement structured; gaps as first-class invitations; conformance as public arithmetic in twenty languages - are precisely the steps that turn "a database" into A SHARED MEMORY FOR HOW THE WORLD WORKS, which has never existed, for any species, ever.

To reify (from a Latin root meaning "to make into a thing") is to take something we normally treat as a fleeting relationship - here, one event causing another - and turn it into a solid, named object you can point at, store, and sign; reified causation as the atom means the single smallest building block of the whole system is one packaged cause-and-effect fact.

Identity as content means a record is named by a fingerprint of what it actually contains, so its name and its contents can never drift apart and duplicates collapse automatically - a technique called content-addressing, where the address IS the content's fingerprint.

Conformance is the public, mechanical checking that a given implementation really follows the standard - here settled by shared arithmetic that anyone can rerun in any of twenty languages, not by anyone's say-so.

If that is a database with extra steps, then yes.

Gloriously, unapologetically yes.

The witness steps down.



APPENDIX: THE SPECIFICATION IN BRIEF

This appendix states the standard completely but minimally.

It is a summary; the normative sources are the files under spec/ in the repository, and the full rationale is the master document Causalontology_Standalone_Design_v27.txt, authoritative wherever this brief and it could be read to differ.

WHAT IT IS

Causalontology is a programming-language-neutral standard – a data structure and its rules, not a program – for recording causation as first-class, content-addressed, signed objects that any implementation in any language can mint and merge without coordination.

It is verb-first, because occurrents (processes) are primary, and noun-hosting, because continuants (things) are their participants.

It models both type-level (generic) causation and token-level (particular) history – and, since 4.0.0, token-level expectation and attributed belief – and the token tier is local by default.

The current specification is 4.0.0.

IDENTIFIERS AND IDENTITY

Every object is named scheme:localpart.



The Roadmap to AGI and ASI

The scheme is the whole-word kind name; the localpart is the lowercase hexadecimal Secure Hash Algorithm 256-bit (SHA-256) digest, sixty-four characters, of the object's identity-bearing fields written in the canonical form of RFC 8785 (the JSON Canonicalization Scheme).

The digest is produced by one procedure: take the object as JSON, remove the fields that are not identity-bearing for its kind, apply RFC 8785 to the remainder, hash those bytes with SHA-256, and set the identifier to the scheme, a colon, and the lowercase hexadecimal digest.

Only two things are ever excluded: the id itself, because it IS the hash, and the signature on the four provenance kinds, because it is computed over these same canonical bytes; a content object is exactly its identity-bearing fields and nothing else.

The digest is taken over the identity-bearing fields that are PRESENT and over nothing else, so an absent field cannot change a name, and two authors anywhere who state the same content mint the same identifier and merge with no coordinator.

Because only present fields are hashed, every optional field added in a later version is identity-preserving: a record written under an earlier version keeps its exact hash under a later one, proved by conformance vectors V106, V111, and V118, and re-proved at 4.0.0 by vector V136, which re-pins two frozen 3.0.0 identifiers byte for byte under the new implementation.

An enrichment's source and timestamp ARE identity-bearing, so the same entry from two sources, or



The Roadmap to AGI and ASI

twice from one source, is deliberately two records
- which is how corroboration is recorded.

VOCABULARY CONVERGENCE

Content-addressing is disciplined three ways so that vocabulary converges rather than fragments.

Every canonical label is one or more whole English words in lowercase snake_case, with homonyms qualified (charge_battery, charge_attack).

Every occurrent and continuant carries a REQUIRED category drawn from a closed enumeration, so the identity inputs are deterministic.

And minting is resolve-before-mint: a contributor first calls resolve to find any existing occurrent or continuant, and creates a new object only when none is found; the resolve minimum is fixed by the standard and tested by vector V37, so it behaves identically on every conformant store.

SERIALIZATION

The canonical wire format is JSON, defined by the JSON Schemas, and the recommended profile is JSON for Linked Data (JSON-LD), which layers a small context onto plain JSON so a document is also valid Resource Description Framework (RDF) data, aligned where possible with Basic Formal Ontology (BFO), the Relation Ontology (RO), and the World Wide Web Consortium Provenance Ontology (PROV).

Protocol Buffers and the Concise Binary Object Representation (CBOR) are permitted as equivalent binary encodings, but the JSON form is normative



The Roadmap to AGI and ASI

and the binary forms must round-trip to it without loss.

THE TWENTY-ONE KINDS

Every object is exactly one of twenty-one kinds: ten type-tier content, seven token-tier content, and four provenance.

For each kind the fields named are its identity-bearing fields - the exact inputs to its identifier - with the required ones first.

The ten type-tier content kinds (immutable, content-addressed):

- `occurrent` - a process or event TYPE, a verb; required label (lowercase snake_case) and category (one of action, event, process, state, state_change); optional stratum.
- `causal_relation_object` - a reified causal claim; required causes and effects; optional mechanism, temporal window (minimum_delay, maximum_delay, unit), modality, context, refines, and skips.
- `continuant` - a thing that endures, a noun; required label and category (one of object, agent, place, substance, collection, information).
- `realizable` - a disposition, function, or role; required kind (one of disposition, function, role) and bearer; optional identity-bearing label.
- `stratum` - a level of description within a named stratification scheme; required label, scheme, and ordinal (an integer where strictly lower is finer-grained); optional unit and governs.



The Roadmap to AGI and ASI

- `bridge` - a cross-stratal identity map declaring one coarse occurrent IS a set of finer ones; required `coarse`, `fine`, and `relation` (one of `constitutes`, `aggregates`, `realizes`, `supervenies_on`).
- `cross_stratal_seam` - a managed jump across non-adjacent strata; required `source`, `target`, and `mechanism_status` (unmodeled or absent); optional `drawn chain`.
- `port` - a typed interface borne by a continuant; required `bearer`, `label`, `direction` (one of `in`, `out`, `bidirectional`), and `accepts`; optional `realizable`.
- `conduit` - a directed, typed connection from port to port, transmissive or computational; required `from`, `to`, and `carries`; optional `label`, `transform`, and `realized_by`.
- `quality` - a property TYPE a thing can bear, such as `mass` or `membrane_potential`; required `label` and `datatype` (one of `quantity`, `categorical`, `boolean`); optional `unit` (a Unified Code for Units of Measure (UCUM) code, required when the datatype is `quantity` and absent otherwise) and optional `stratum`.

The seven token-tier content kinds (immutable, content-addressed, LOCAL BY DEFAULT):

- `token_individual` - a particular thing; required `instantiates` (the continuant it instantiates); optional `designator` (an opaque local designator) and `part_of` (token mereology).
- `token_occurrence` - a particular happening; required `instantiates` (the occurrent it

instantiates) and interval; optional participants, locus, and observer.

- `state_assertion` - a particular individual (the subject) bearing a particular value of a particular quality over an interval; required subject, quality, value, and interval; the value is exactly one of a quantity with a UCUM unit, a categorical string, or a boolean.
- `token_causal_claim` - a particular happening causing another; required causes and effects (each a `token_occurrence`); optional `covering_law`, `actual_delay` (a duration with a unit), and a but-for counterfactual.
- `attitude` - one mind's stance toward a content, what the holder's mind CONTAINS rather than what is true (the philosophers' term is a propositional attitude); required holder (a continuant or a `token_individual` - a modeled agent, never a signing key), `attitude_type` (one of believes, desires, intends, knows, expects, fears, a closed enumeration), and content (a reference by identity to any content object, including another attitude, which is how attributions nest); the content may be false, and an attitude carries no strength of its own.
- `predicted_occurrence` - an occurrence a source expects rather than reports; required `instantiates` (the occurrent type predicted), `interval` (exactly one temporal dimension: a wall-clock start with optional end, or an ordinal `start_tick` with optional `end_tick`, never both and never neither), and `predictor` (the predicting agent or construct - a modeled identity, never a key); optional strength, the predicted probability, identity-

bearing when present because it is part of the predicted content itself.

- `prediction_error` - the graded reckoning between a prediction and the world; required predicted (the predicted_occurrence it grades) and discrepancy (a signed number read as actual minus expected: positive exceeded, negative fell short, zero fulfilled); optional observed (the token_occurrence that answered the prediction - absent means nothing did).

The four provenance kinds (signed, add-only):

- `assertion` - source *S* asserts content object *C*; required about, source, evidence_type, confidence, and timestamp; optional strength (meaningful only when *C* is a causal_relation_object), evidence, and evidenced_by citing the particular token records that are the evidence.
- `enrichment` - the signed addition of one entry to one field of one object; about, field, entry, source, and timestamp; the fields are aliases, participants, subsumes, part_of, realized_in, occurrent_subsumes, and occurrent_part_of.
- `retraction` - a source's withdrawal of its own earlier assertion or enrichment; retracts, source, and timestamp.
- `succession` - a signed declaration that one key succeeds another, signed by the old key; predecessor, successor, and timestamp.

THE LOAD-BEARING DECISION

Content is separated from provenance, uniformly.



The Roadmap to AGI and ASI

Content objects are pure and immutable; every mutable datum – alias, link, claim, withdrawal, or key rotation – is a signed, add-only record with an author.

THE PRINCIPLES

There are seven, numbered P1 through P7, and this is the whole set.

P1, no free-floating tokens, and the arrow runs one way: every token instantiates a type, and a token naming no type is a schema failure; and no type record may reference a token, the sole exception being an assertion's evidenced_by, which is provenance, not content – so the shared type-tier commons is never polluted by any particular mind's token history.

P2, tokens local by default: a conformant store never publishes token-tier records to a shared commons unless its operator explicitly says so, per record or per policy; the rationale is safety, made normative and non-optional as safety rule R1.

P3, levels are scheme-relative: a stratum has meaning only within a named stratification scheme, and two strata are comparable only when they share that scheme, so rival stratifications coexist and a cross-scheme ordinal comparison is refused as a hard failure.

P4, content bears no provenance: no content kind carries strength, confidence, probability, source, or timestamp-of-assertion; the one recorded 4.0.0 exception is the predicted_occurrence's strength, which is the grade of the expectation itself –



The Roadmap to AGI and ASI

predicted content, not a provenance evaluation - relaxed for that single kind deliberately, narrowly, and by extension.

P5, absence made positive: setting skips to true converts a "no mechanism" gap into a positive finding rather than a defect.

P6, every addition is forced: no construct enters the standard unless a recorded, specific failure forced it, and constructs that are tempting but unforced are declined, with the declinations recorded.

P7, whole words, never abbreviations: every identifier scheme and type value is one whole English word; only the proper names of external standards are kept verbatim (ed25519, SHA-256, RFC 8785, RFC 3339, UCUM, UTC, JSON, JSON-LD, BFO, RO, PROV).

MODALITY

A `causal_relation_object`'s modality is one of five: necessary, sufficient, contributory, enabling (the cause makes the effect possible without bringing it about), and preventive (which opposes the other four).

TIME

A temporal window carries a unit that is one of a closed enumeration.

The eight wall-clock units are instant, seconds, minutes, hours, days, weeks, months, and years, and temporal admissibility converts each to seconds by fixed constants that every conformant



The Roadmap to AGI and ASI

implementation must use identically: instant is 0, seconds 1, minutes 60, hours 3,600, days 86,400, weeks 604,800, months 2,629,746 (the average Gregorian month), and years 31,556,952 (the average Gregorian year).

An implementation that uses any other constant is non-conformant.

The ninth unit, ticks, is the ordinal dimension - a dimensionless step ordered by integer comparison; it carries no conversion constant, is held as a dimension disjoint from wall-clock time, and converting ticks to seconds is refused as a category error.

PROVENANCE, SIGNATURES, AND EVIDENCE

A source identity is an Ed25519 public key (RFC 8032), written as ed25519 followed by sixty-four hexadecimal characters.

Every provenance record is signed over its canonical identity-bearing bytes, so verification needs only the record itself and re-submission is idempotent.

Evidence is graded on a ladder: intervention outranks observation, observation outranks simulation, then derivation, then human_hint, then imported - acting beats watching, and watching beats simulating.

An assertion may cite the particular token records that are its evidence through evidenced_by, so a law induced from ten interventions can be weighed against one induced from a thousand observations.



The Roadmap to AGI and ASI

An object's displayed links are materialized from its unretracted enrichment records, and the same entry from two sources shows as one entry with two contributors.

Trust is not forced: the store keeps every signed claim and lets each consumer choose a trust policy by source, by evidence type, by confidence, and by recency; contradictions coexist and are surfaced.

The recommended default aggregation, advised but never mandated, is a confidence-weighted mean over trusted sources, with intervention-grade evidence weighted above observation.

THE TWENTY-FIVE SEMANTIC RULES

A document is schema-valid if it matches its JSON Schema, and semantically valid if it also satisfies these rules.

1. Temporal windows: `minimum_delay` is at most `maximum_delay`.
2. Dangling reference: any identifier referenced but absent from the store surfaces `dangling_reference`.
3. Refinement: R refines P only if their causes and effects are equal, R repeats every field P specifies with the same value, and R adds at least one field P left unspecified.
4. Acyclicity: the mechanism, refines, subsumes, `part_of`, `bridge`, `occurrent_subsumes`, `occurrent_part_of`, and token `part_of` graphs are each acyclic.

5. Strength: strength is the source's estimate of the probability that the effects become possible given a do-intervention on the causes, in context.

6. Conflict: necessary, sufficient, contributory, and enabling are mutually compatible, and preventive opposes all four; two relations conflict when they share causes, effects, and compatible contexts, and their windows overlap, but their modalities oppose.

7. Hierarchy consistency, bridged reachability: after bridge-closing each cause and effect, every cause must reach some effect through the mechanism graph, and absent members make the check indeterminate rather than failed.

8. Open world: absence is not denial, and negative causal knowledge uses the preventive modality, never absence.

9. Merge: immutable objects, add-only records, set-union.

10. Retraction: valid only from the retracted record's own source or lineage.

11. Succession: chains define a lineage, one successor per key.

12. Enrichment validity: each enrichment field accepts only entries of its declared kind.

13. Stratal comparability: two strata are comparable only if they share a scheme, and comparing ordinals across schemes surfaces `scheme_mismatch`.



The Roadmap to AGI and ASI

14. Bridge well-formedness: coarse and every fine member are stratified, all fine share one stratum, the strata share a scheme, and coarse's ordinal is strictly greater than fine's.

15. Stratal classification: every fully stratified relation is exactly one of intra-stratal, adjacent-stratal, skipping, or mixed, recomputed on ingest.

16. Skip coherence: skips true requires an empty mechanism, skips true on a non-skipping relation is vacuous, and a skipping relation with an empty mechanism and skips absent is incomplete.

17. Conduit well-formedness: a transmissive conduit's carries must be accepted by both ports, and a computational conduit's target must accept its transform's effects.

18. Token instantiation: every token_individual names a continuant and every token_occurrence names an occurrent, so there are no free-floating tokens.

19. State type and unit coherence: a state value's shape must match its quality's datatype, and a quantity value's unit must equal the quality's unit.

20. Covering-law coherence: a token claim's tokens should instantiate its covering law's occurrents, and any actual delay must fall within the law's window.

21. Temporal coherence: every cause token's interval begins no later than every effect



The Roadmap to AGI and ASI

token's, and token-tier backward causation is not represented.

22. Cross-stratal seam well-formedness: both endpoints are stratified in one scheme with non-adjacent ordinals; a drawn chain forbids mechanism_status absent and must climb strictly through intervening strata; and the seam belongs to the coarsest stratum it touches.

23. Ordinal units are a separate dimension: a tick window and a wall-clock window never compare, never overlap, and never convert.

24. A prediction is not a report: a predicted_occurrence satisfies no rule that requires a token_occurrence; its interval carries exactly one temporal dimension, wall-clock or ordinal, never both (dimension_conflict) and never neither (missing_dimension); a prediction_error's observed occurrence, when present, must instantiate the same occurrent its prediction instantiates (pairing_mismatch); and discrepancy reads actual minus expected.

25. An attitude's content is quarantined: no reasoning, conflict, reachability, or classification rule treats an attitude's content as asserted; content that contradicts the actual record raises no conflict, because that mismatch is a false belief, first-class and shareable; the holder is a modeled agent, never a signing key; the content may reference any content object, including another attitude; and the attitude_type enumeration is closed.

Beyond the twenty-five, one store rule governs merge under decentralized replication: where a



The Roadmap to AGI and ASI

cycle can arise from concurrent writes, an enforcing tier rejects a cycle-creating enrichment, and every conformant view breaks an emergent cycle the same way - excluding the cycle-completing record with the latest timestamp, ties broken by lexicographic record identifier - and surfaces it as an inconsistent_hierarchy repair gap.

The determinism of that tie-break keeps every replica's view identical, which is what makes the store a conflict-free replicated data type.

THE NORMATIVE ALGORITHMS

Six algorithms make the executable rules decidable, and every binding implements them identically.

Algorithm A, bridge closure: the set of occurments identical to a given one across strata.

Algorithm B, hierarchy consistency: the bridged reachability that Rule 7 requires.

Algorithm C, stratal classification: intra, adjacent, skipping, or mixed, for Rule 15.

Algorithm D, skip coherence: the skip asymmetry of Rule 16.

Algorithm E, unit normalization: comparing wall-clock delays against a window, for Rules 19 and 20.

Algorithm F, cross-stratal seam well-formedness: the checks of Rule 22.



The Roadmap to AGI and ASI

THE GAP TAXONOMY

The store answers a gaps query with diagnostics of two kinds.

The HARD gaps are schema or semantic failures that MUST be rejected: `schema_mismatch`, `malformed_bridge`, `bridge_cycle`, `malformed_conduit`, `contradictory_skip`, `individual_cycle`, `value_type_mismatch`, `unit_mismatch`, `occurrent_hierarchy_cycle`, `occurrent_mereology_cycle`, `retrocausal_claim`, `malformed_seam`, `contradictory_seam`, `dimension_conflict`, `missing_dimension`, and `pairing_mismatch`.

The rest are stigmergic invitations, standing requests for the next contribution: `missing_field`, `dangling_reference`, `empty_mechanism`, `inconsistent_hierarchy`, `conflict`, `demand_supply`, `unstratified_occurrent`, `collided_realizable`, `vacuous_skip`, `stratum_mismatch`, `uncovered_causation` (the most valuable), `delay_outside_window`, `subsumption_crosses_strata`, `covering_law_mismatch`, `mixed_stratal_endpoints`, and `incomplete_mechanism`.

THE STORE

Every conformant implementation provides the complete abstract operation set: `canonicalize`, `identify`, `validate_schema`, `validate_semantics`, `admissible`, `conflicts`, `hierarchy_consistent` (bridged reachability), `bridge_closure`, `classify_cro`, `skip_gaps`, `delay_within_window`, `sign`, and `verify`; and, against a store, `put`, `put_record`, `get` (with materialized enrichments),



The Roadmap to AGI and ASI

assertions_about, enrichments_about,
retractions_of, lineage, resolve, query, and gaps.

The reference protocol is a Hypertext Transfer Protocol (HTTP) binding, the Tier A reference: a write to POST /objects stores a content object with its identifier recomputed; GET /objects/{id} returns the object with its materialized view; a write to POST /records stores a signed assertion, enrichment, retraction, or succession, and anything unsigned is quarantined; GET endpoints read /records/{id}, /assertions, /enrichments, /retractions, /successions, /resolve, /gaps, and /conflicts; and POST /query runs query-by-example.

The binding carries three required disciplines: every list endpoint is paginated by a limit (default one hundred, maximum one thousand) and an opaque cursor that returns a next_cursor; Tier A write access requires an account bearer token in addition to the per-record signature - the token governs resource use, the signature carries trust - with per-account and per-key rate limits; and two minimums are mandatory, a query-by-example minimum of equality and set-containment filters over one kind, and a resolve minimum in which exact canonical-label and alias matches must be returned and ranked ahead of any permitted fuzzy match.

The data model runs unchanged across three tiers: A, one central hosted store; B, federated nodes that push and pull like Git; and C, decentralized content-addressed replication over a Merkle directed acyclic graph, which is a conflict-free replicated data type by construction, with the deterministic cycle-breaking view as its only edge case.



The Roadmap to AGI and ASI

SAFETY: ABUSE RESISTANCE

The commons is for general, type-level knowledge, and its safety rests on the signature.

- Forgery is impossible without the private key, because every record is signed.
- Sybil attacks gain nothing, because keys are free to mint and trust never attaches to key-count; it attaches to consumer-chosen trust policies over specific sources, Tier A adds rate limits, and reputation weights history and corroboration, never volume.
- Data poisoning is attributable, because signed provenance pins poison to its key; retraction gives an honest source an exit while a poisoner's history stays pinned to its key, and conflicts surface rather than silently average.
- Vocabulary vandalism is contained, because every enrichment is signed - attributable at every tier, suppressible by source key, and removable by the author's own retraction - and grow-only views mean vandalism can never destroy good data.
- Unsigned or unverifiable material is never silently trusted: a record whose Ed25519 signature does not verify is rejected or placed in a clearly marked quarantine tier that default queries exclude.

Claims of health, safety, legal, or financial consequence are always served with their provenance and confidence shown, never as bare fact, and must be signed or be quarantined without



The Roadmap to AGI and ASI

exception; and every interface carries a standing notice that the commons records who claims what, with what evidence, and does not itself dispense medical, legal, or financial advice.

Takedown is honest and stated in advance, by tier: a Tier A operator hard-deletes with a public log; a Tier B origin deletes locally and issues a signed suppression notice that peers should honor; and a Tier C copy cannot be force-deleted, so takedown is a signed tombstone honored by every conformant default view, plus the public log.

SAFETY: THE UNFORGETTING STORE

Because Causalontology is content-addressed, add-only, and merged by set union, a record cannot be edited and cannot be redacted, retraction withdraws a claim but never erases a record, and federation replicates by union; therefore a personal identifier written into a published token_individual designator is permanent, global, and irrevocable – a mathematical consequence of content addressing that no future version of the specification can repair.

Six rules bind the token tier, and they are a blocking precondition of the tier itself, not advice:

- R1, tokens are local by default: a conformant store MUST NOT publish token-tier records to any federated peer or public endpoint unless its operator explicitly instructs it, per record or by explicit policy, and every implementation's default configuration MUST be non-publishing.



The Roadmap to AGI and ASI

- R2, no natural identifiers in published designators: a published designator **MUST NOT** be a name, medical record number, national identifier, email address, device serial, or street address, and **SHOULD** be a salted hash whose salt is held locally and never published.
- R3, the store must warn: an implementation **MUST** warn when a `token_individual` with a non-hash-shaped designator is submitted for publication, and **SHOULD** refuse by default and require an explicit override.
- R4, special categories: token records concerning health, biometrics, genetics, sexual life, political opinion, religious belief, trade union membership, or criminal history are special category data under multiple legal regimes, and an implementation serving such data **MUST** document its legal basis and **MUST NOT** publish such tokens to a public commons under any configuration.
- R5, the commons accumulates laws, not diaries: the intended flow observes tokens locally, induces type-level claims from them, and contributes the types, citing its tokens only by hash in `evidenced_by` if at all.
- R6, erasure is not satisfiable by retraction: implementers operating under the General Data Protection Regulation (GDPR) Article 17 or equivalent **MUST NOT** represent retraction as erasure, because the only way to satisfy an erasure request for published token data is to never have published it, which is why R1 exists.

SNAPSHOTS



The Roadmap to AGI and ASI

A snapshot is a signed, content-addressed, point-in-time dump of the commons that anyone can download, verify byte for byte with nothing but its own bytes, and use to stand up a mirror, giving durability before any live peer exists.

The rule is self-verifying and reproducible: every object proves itself by its own hash, every record by its own signature, and the whole is committed by one Merkle root and authenticated by one Ed25519 signature over the manifest.

A snapshot is a body and a manifest.

The body is newline-delimited canonical JSON, one content object or provenance record per line in its exact canonical bytes, all entries sorted by identifier, so the same store always emits byte-identical bytes and a large snapshot is read one line at a time; each line carries its own type, so the body is self-describing.

The manifest is a small signed JSON header carrying nine fields: `snapshot_format`, the format version; `spec_version`, the specification version; `created_at`, a Coordinated Universal Time timestamp and the only non-reproducible field; `content_objects` and `provenance_records`, the counts in the body; `includes_tokens`, false by default; `merkle_root`, one hash committing to every byte of the sorted body; `signed_by`, the publishing node's key; and `signature`, the Ed25519 signature over the manifest minus the signature field.

The `created_at` is inside the signed bytes but is excluded from the Merkle root, so the root is content-only and the same store yields the same



The Roadmap to AGI and ASI

root forever while each publication still carries an honest timestamp.

The Merkle root follows RFC 6962 domain separation: a leaf is the SHA-256 of a zero byte followed by the line, an interior node is the SHA-256 of a one byte followed by the left then the right child, an odd node is promoted unchanged, and an empty body has the root SHA-256 of the empty string.

A default snapshot is the shareable commons and MUST NOT include token-tier records - token_individual, token_occurrence, state_assertion, token_causal_claim, attitude, predicted_occurrence, prediction_error - nor any provenance record that references one, unless the operator explicitly opts in, recorded as includes_tokens true; this makes the local-by-default rule concrete, so a snapshot is exactly the laws-not-diaries boundary.

Export is deterministic and writes four files: the body, the signed manifest, a detached SHA-256 checksum, and a detached signature.

Verify checks a dump with nothing but its own bytes - the manifest signature, the Merkle root recomputed from the delivered body, every content object's identifier against the hash of its canonical bytes, every provenance record's Ed25519 signature, and the manifest's declared counts.

Import runs the full verification and then merges by set union; any failure raises before any write, so a tampered snapshot never partially loads, and importing the same valid snapshot twice is a no-op.



The Roadmap to AGI and ASI

CONFORMANCE

An implementation is Causalontology-conformant if and only if it passes every published conformance vector, V01 through V137, for the specification version it declares.

Beneath the vectors, thirty design decisions stand LOCKED - numbered, written down, and changeable only by a public, versioned process; the additions that grew the model to its present shape - the twenty-one-kind enumeration, the closed attitude_type list, the doxastic quarantine, the predicted strength, and the widened assertion reference among them - were each carried in the open as extensions to that locked set, under the enumeration rule that reserves a MAJOR version for exactly such a change.

VERSIONING

Specification versions follow Semantic Versioning: a new object kind, an extension of any closed enumeration, or any break to identity is a MAJOR version, and the vectors are re-frozen once per release.

The closed enumerations include the occurrent, continuant, and quality datatypes and categories, the modalities, the evidence types, the attitude types, the enrichment fields, and the temporal-window unit set, and extending any of them is MAJOR even where old data round-trips, because an old validator would reject the new value and identity computed over new vocabulary would not be comparable.



The Roadmap to AGI and ASI

Each release is frozen and permanently reachable by its git tag, so a consumer can pin an exact specification; the appendix Version History records what each release added, and every addition across the standard's life has been additive and identity-preserving, witnessed at the current specification by vector V136, which re-pins the frozen bytes of earlier records unchanged.

The current specification is 4.0.0.

GOVERNANCE

The locked decisions and the closed enumerations change only by a public, versioned process, in the spirit of a World Wide Web Consortium (W3C) or Internet Engineering Task Force (IETF) working group.

A change is a written proposal, reviewed in the open, carrying a conformance impact statement that names exactly which vectors change, and a decision recorded in the repository.

Recovery from a compromised key is a governance action, not a cryptographic one, and is stated honestly as the limit of cryptographic identity.

Every takedown, and every registry re-binding after a key compromise, is publicly logged.

LICENSING

The license is deliberately split so the commons can be reused without friction.



The Roadmap to AGI and ASI

The specification and all reference code are under the Apache License 2.0 – permissive and patent-protective.

The data in the commons is under Creative Commons Zero (CC0), the public-domain dedication, so a downstream user need track no attribution chain; attribution survives as provenance, a matter of record, not as a legal obligation.

Externally-licensed content that is incompatible with this is linked, never copied in.

A WORKED EXAMPLE

Two occurrents, one bare causal claim refined into a signed law, and a signed assertion – the whole loop, in JSON:

```
{ "type": "occurrent", "label": "press_button",  
  "category": "action" }
```

```
{ "type": "occurrent", "label": "light_on",  
  "category": "state_change" }
```

```
{ "type": "causal_relation_object", "causes":  
  ["occurrent:press_button"], "effects":  
  ["occurrent:light_on"], "temporal": {  
    "minimum_delay": 0, "maximum_delay": 1, "unit":  
    "seconds" }, "modality": "sufficient" }
```

```
{ "type": "assertion", "about":  
  "causal_relation_object:...", "evidence_type":  
  "intervention", "strength": 0.98, "confidence":  
  0.95, "source": "ed25519:...", "timestamp": "2026-  
07-13T00:00:00Z" }
```

Its meaning: pressing the button is a sufficient cause of the light coming on within a second, asserted on intervention-grade evidence by a signed source - a complete claim, its identity computable from its content, its truth carried by a signature nobody else can forge.

Two more examples show the full reach of the standard - one mind modelling another, and a mind measuring its own expectations against the world.

THEORY OF MIND - A FALSE, NESTED BELIEF

Here a mind is right, on the public record, about another mind being wrong.

The actual record says the lamp was off across the interval:

```
{ "type": "state_assertion", "subject":  
  "token_individual:lamp_seven", "quality":  
    "quality:power_state", "value": "off", "interval":  
      { "start": "2026-07-22T09:04:00Z", "end": "2026-  
07-22T09:11:00Z" } }
```

Agent B, an agent modelled inside the ontology, nonetheless believes the lamp was on - a belief whose content points, by identity, at a rival state_assertion reading "on" over that same interval:

```
{ "type": "attitude", "holder":  
  "continuant:agent_b", "attitude_type": "believes",  
  "content":  
    "state_assertion:...lamp_seven_power_state_on..."  
}
```

And Agent A believes that B believes it - the outer attitude's content is the inner attitude itself, named by identity, so a two-story thought becomes two honest records, the outer pointing at the inner:

```
{ "type": "attitude", "holder":  
  "continuant:agent_a", "attitude_type": "believes",  
  "content":  
    "attitude:...agent_b_believes_lamp_on..." }
```

Its meaning: A correctly models B's false belief that the lamp was on, while the record shows it off.

Semantics Rule 25 raises NO conflict - the mismatch between B's belief and the actual record is not an error in the commons; it IS a false belief, first-class and shareable, exactly as an attitude's quarantined content is meant to be.

Each holder is a modelled agent named as a continuant, forever distinct from whatever signing key (an ed25519 identity) asserts these records into the commons - the mind being described and the hand doing the signing are two different things.

THE PREDICTION LOOP - EXPECT, OBSERVE, COMPARE

Here a mind runs the oldest loop a learner runs: expect something, watch what comes, and grade the difference.

A predictor expects the light within a short window of its own internal steps - an ordinal tick window, not the wall clock - and states how strongly it expects it:

```
{ "type": "predicted_occurrence", "instantiates":  
  "occurrent:light_on", "interval": { "start_tick":  
    10, "end_tick": 12 }, "predictor":  
  "continuant:agent_a", "strength": 0.9 }
```

The light then actually comes on – a real happening at a real moment, recorded as a token occurrence:

```
{ "type": "token_occurrence", "instantiates":  
  "occurrent:light_on", "interval": { "start":  
    "2026-07-22T09:04:07Z" } }
```

And a prediction_error pairs the two and grades the forecast, its discrepancy read as actual minus expected:

```
{ "type": "prediction_error", "predicted":  
  "predicted_occurrence:...", "observed":  
  "token_occurrence:...", "discrepancy": 0 }
```

Its meaning: the predictor called it, the world obliged, and the signed discrepancy of zero records a prediction fulfilled – and because the observed occurrence instantiates the very occurrent the prediction named, the pairing is honest and could not have been graded against the wrong event.

Had nothing arrived, the observed field would simply be absent, and the discrepancy would carry a negative value – the surprise of nothing coming, a silence the record can grade as plainly as it grades an event.

APPENDIX: VERSION HISTORY



The Roadmap to AGI and ASI

This is the one place in the book that tells the standard as a story of versions; everywhere else it is told as one finished whole, which is what it now is.

The standard grew in four public releases, and each one is frozen forever behind its own git tag, so anyone can fetch exactly the version they want and check it byte for byte.

Version 1.0.0 was the founding move made real: eight kinds of object and thirty-eight conformance vectors.

It reified causation - it turned the causal arrow into a Causal Relation Object you can hold - it gave every object a content-addressed identity (a name computed from the object's own content), and it gave every claim a signed provenance (a tamper-proof mark of who vouched for it).

That was the seed: cause as a thing, identity as arithmetic, authorship as a signature.

Version 2.0.0 was the great widening, and it paid an honest price to get there.

To obey Principle P7 - the refusal of abbreviations, so that no reader ever has to decode a cipher to read a claim - every name in the standard was re-minted in whole words, a deliberate, once-only break made in the open.

It grew to seventeen kinds and one hundred and seven conformance vectors, and it added the machinery of levels and connection: strata (named rungs of description), bridges (the declaration that a coarse happening simply IS a set of finer



The Roadmap to AGI and ASI

ones), ports and conduits (typed interfaces and the typed connections that run between them), the whole token tier (the record of what actually happened, not merely what always happens), the enabling modality (a cause that opens the door without pushing), bridged reachability (Rule 7, which lets a chain of causes be checked across several scales at once), and honest skip semantics (a cause may declare that it leaps across non-adjacent levels with no mechanism claimed between).

Version 3.0.0 was three careful closures, each additive and identity-preserving.

It grew to eighteen kinds and one hundred and nineteen conformance vectors, and it let the fabric of structure clasp the living dynamics beneath it.

It added the ordinal tick unit, so a mind that counts its own steps can record time on the phenomenon's own terms rather than being forced into the wall clock.

It added the managed `cross_stratal_seam`, which turns a bare leap into a first-class record stating, in its `mechanism_status` field, whether the skipped mechanism is merely unmodeled (it exists, undrawn) or genuinely absent (there is none) – honest ignorance made explicit – checked by Algorithm F.

And it added the conduit's `realized_by` field, the signed pointer that binds a structural connection to the native law that actually does its computing.



The Roadmap to AGI and ASI

Version 4.0.0 is the current specification, and it is the reason the token tier now holds not only the was but the will-be and the believed.

It grew to twenty-one kinds and one hundred and thirty-seven conformance vectors, adding three token-tier kinds - the attitude (what a mind holds), the predicted_occurrence (what a mind expects), and the prediction_error (how the expectation met the world) - governed by two new semantic laws, Rule 24 (a prediction is not a report) and Rule 25 (an attitude's content is quarantined, never read as a fact).

But 4.0.0 was RELEASED BY A GATE, and the gate is the part worth telling slowly, because it is Principle P6 - no construct enters the ontology from the armchair - enforced with a paper trail.

Long before these kinds were written, a held change order forbade the additions outright until a real build should hit a real, recorded wall.

The build was konnectome, the glass-box synthetic mind, and it hit the wall twice in a single day.

Wall-1: the mind, modelling a companion, worked out that the companion held a false belief, and could hold that thought privately but could neither represent it nor share it, because no existing kind could say "this agent's mind contains that content" - and that wall released the attitude.

Wall-2: the mind's predictive loop found that a prediction minted as a token occurrence is byte-identical to a claim of fact, the very same name



The Roadmap to AGI and ASI

for "I expect the light" and "the light came on" - and that wall released the two prediction kinds.

Both walls were recorded on 2026-07-22; both additions passed the consumer-versus-ontology test - the standard's gate question of whether a need is one particular consumer's private business or a genuine gap that every mind would meet in the shared ontology; and both were admitted with the owner's written authorization.

Three further parts of that same change order - Section C, an imagination quarantine; Section D, quantitative sensitivity; and Section E, multi-language expression - remain HELD to this day, forbidden until their own real walls arrive.

That is the whole discipline, said once: every version is additive and identity-preserving - because content addressing hashes only the fields that are present, a record written under an earlier release keeps its exact name under a later one, witnessed at 4.0.0 by vector V136 (which re-pins two frozen earlier identifiers, unchanged) and by vector V137 (which refuses the abbreviated forms of the new names) - and the locked decisions at the standard's core are never quietly rewritten, only amended by extension, in the open, with every change dated and reviewed where anyone can watch.

Four releases, one identity arithmetic, and a single rule holding them together: the standard admits a kind only when the world forces its hand, and never forgets a thing it once said.

The current specification is 4.0.0.



The Roadmap to AGI and ASI

END OF MANIFESTO.



The Roadmap to AGI and ASI

BACK COVER MATTER

Dear Reader,

Thank you for reading:

Causalontology: AGI's Missing Link!

This is an evolving work. (As this is an evolving field!)

If you are aware of emerging alternative information that could further the field of Artificial Intelligence Engineering into the future, then please notify me so that the project may be improved in the next Edition!

<https://www.linkedin.com/in/d-r-dison>

I hope you enjoyed this book!

Positive reviews are the lifeblood of independent authors.

Praise is always welcome in the form of a review, as it helps other readers to find the work.

Constructive criticisms are welcome privately at my LinkedIn so that I can keep improving.

Copyright © 2026 by D. R. Dison (LinkedIn) All rights reserved.

No part of this book may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying,



The Roadmap to AGI and ASI

recording, or otherwise, without the prior written permission of the publisher, D. R. Dison, except as provided by U.S. Copyright Law or for the use of brief quotations in a book review. Acknowledgment: This work was produced with the aid of various technological tools, including Word Processing Tools (WPT), Desktop Publishing Tools (DPT), Image Processing Tools (IPT), and Artificial Intelligence Tools (AIT).

(International Standard Book Number) ISBN: 979-8-9950671-0-X

Published in the United States of America.

Printer's Key Control Number 10 9 8 7 6 5 4 3 2 1

